

# The Effect of Training Published Computer Vision Models with Unreal Engine Synthetic Images

An Introduction to the Synergy of Graphic Rendering Software and Machine Learning

John W. Belanger Smutny, Software Engineer, Northrop Grumman Corporation  
Bradley Department of Electrical and Computer Engineering, Virginia Tech

**Abstract**— All machine learning (ML) models require large amounts of varied training data to understand the real world. Gathering enough data can be a daunting task if you don't already have it or your task isn't funded well enough to gather it (manually or via web-scraping). The latter of which opens the door to trademark and privacy concerns. Is it possible to generate your own images for your use case in a controlled scalable environment? This project explores this possibility through five experiments to better understand the domain gap between real images and synthetic images created with the Unreal Engine (UE) photo-realistic rendering software. The used synthetic images are collected for two binary classification use cases: First, in a domestic setting (can the model identify if an image is of a cat or a dog). Second, to be used for quality assurance in industrial settings (can the model identify if a metal weld is defective).

The large VGG16 model showed slight improvements to model performance when synthetic images supplemented real images (real-to-synthetic ratio that is greater than or equal to 50:50). In addition, using synthetic padding to rebalance an imbalance classification dataset is possible, but should be done with caution based on the quality of synthetic images. The on-chip focused MobileNetV3-small architecture struggled to converge during learning for most experiments. Overall, the results from both models provide context to how different types of models react to synthetic images.

**Index Terms**— Computer Vision, Unreal Engine, Real Synthetic Ratio, Synthetic Utilization, Welding.

## 1.I. Introduction & Problem Statement

THIS paper investigates how synthetically generated images from the Unreal Engine can be used to improve CV model performance. This paper explores three questions; first, investigate how including synthetically generated images in a Computer Vision (CV) model's training set affects the end model performance (accuracy, loss, and AUCROC). Second, see if making more synthetic images available to be sampled changes model performance metrics. Third, can synthetic images rebalance an imbalanced classification dataset by padding the minority class. In addition, this project provides context on how practical it is to develop one's own synthetic images using Unreal Engine (UE), the potential business cases for it, and possible restrictions.

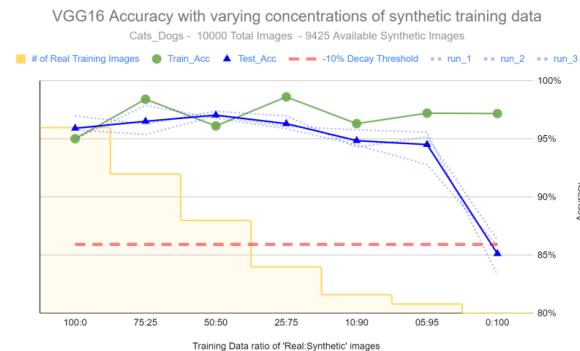


Figure 1.1: Example results from Experiment 1 showing how a transfer learning binary classifier model's accuracy changes as more of the training data is replaced with synthetic images.

To investigate these goals, sets of synthetic images are created for each case: 1) Cat vs Dog and 2) Metal Weld Defect Detection. Free to use .fbx files are used to create the Cat and Dog images, while the author applies their own changes to a purchased metal welding UE Material. Resulting in unique UE Material Instances for three weld defects: porosity, cracks, and incomplete welds. Two differently sized transfer learning models are used for these use cases: VGG16 and MobileNetV3-small. The variation in models used and use cases allows this research to investigate the domain gap between real and synthetic images, as well as how differently sized models are affected by them.

Models	VGG16, MobileNetV3	
Binary Classification use Cases	{Cat, Dog}, {Defective weld, Good weld}	
Experiments	How does the number of synthetic images in a model's training set affect binary classifier performance?	Can synthetic images be used to rebalance an imbalance classifier dataset?

Figure 1.2: General outline of methods used in this research.

## A. Research Outline

This paper covers all design decisions, methods, results, and interpretations on how training CV models on synthetic images can affect model key performance metrics for binary classification tasks. First, the paper covers two concepts relevant to the paper before describing the current literature related to training CV models off of synthetic images. Next, the paper describes the methodology used by the researcher, specifically describing the models used, architecture changes, image augmentations during training, dataset formation, and an outline of performed experiments. Followed by the graphical results, interpretations of said results and final conclusions.

Each section contains high level information meant to be an overview of what is done to understand the next section. At the end of each topic, the paper points the reader to a section in the appendix with more information on that topic. The appendix's "For More Information" section expands upon the high-level knowledge given in the body of the paper. This is so the reader can focus on what they care about the most.

## A. Key Terms and Concepts

The findings in this report rely on several background terminology and concepts. The most important terms being *CV model generalization*, *latent variation* in a set of images, and how metal MIG and TIG welding works. This paper aims to introduce two new concepts: *Real:Synthetic Ratio (RSR)* and *Synthetic Utilization Rate (SUR)*. RSR is the ratio of real images to synthetic images used in a model's training set. SUR is the percentage of synthetic data sampled to create a dataset versus what was available to be sampled.

Introduced in this Research	Real:Synthetic Ratio (RSR) Synthetic Utilization Rate (SUR)
Computer Vision	Model Generalization Latent Variation Transfer learning
UE Graphic Design	Material Texture map

Figure 1.3: Key terms to know to understand this paper's results.

For the definitions of each term and concept, please see the appendix the *A.I. Key Terms and Concepts* section "For More Information".

## B. Related Works

Using synthetic images to train ML applications is a relevant pursuit as applications become more diverse and as abundant real data is harder to find. In 2022, Keith Man and Javaan Chahl published a technical summary outlining various types of synthetic image creation, how synthetic data is generated, for what purposes, and provided an outline of how to measure the quality of Generative Adversarial Networks (GANs) images. The common measure of GAN quality is an FID Score (Frechet Inception Distance score). An FID score measures how far apart real and synthetic image feature spaces are from one another, where a score of 0.0 means the pictures are nearly identical. The survey documents that some of the factors to successfully using synthetic images:

1. How the experiment datasets are formed
2. How variation is there in the generated images
3. How to make the images as real as possible.  
Reducing the domain gap between real and synthetic is paramount [Man22].

The FID scores of synthetic images created for this report were not calculated but could be good for future research.

### i. Synthetic Data Substitution

Two papers specifically are used as the basis for this research. Papers published by an anonymous ECCV submission source [ECCV18] and Johan Thornström attempt to use images from the UE and (GANs) to train Convolutional Neural Networks (CNN) for classification. Both papers use the UE to render images with different variations for different binary classification use cases (airplane or bicycle and soldier or civilian respectively) [ECCV18] [Thornström19].

The ECCV paper's experiment trains the AlexNet CV model with different mixtures of real and synthetic images. Images of airplanes and bicycles against different UE full environments, backgrounds, lighting sources, and textures to create image variation. In addition, their experiment involved using GAN models (MrGAN or CycleGAN) to post-process the real and synthetically generated images to look "more real". The GANs were trained to improve the realism features in the image's feature vectors to mostly positive results. Their results indicate that mixing real and synthetic data at a ratio of 1:1 with 2D augmentations yields classifiers with the highest accuracy. In addition, using the GANs to further refine all training images improved test accuracy. The key point

being that adding any amount of synthetic data to a training set of real images improves model performance by introducing new feature variance [ECCV18]. It is believed that this submission is made by Nguyen, Tan, Chen, H, et al of Rice University as the paper was found on their personal blog, but the true author cannot be verified at this time.

Thornström's paper repeatedly trains a ResNet50 model with either a real image dataset or different variations of only synthetic images. Their paper "Domain Adaptation of Unreal Images for Image Classification" uses a python package, UnrealCV, to extract segmented images from Unreal Engine and feed them into a multi-staged GAN to produce unique images used to classify humans as soldiers or civilians. The model struggled to learn due to GAN images not producing enough image variation. After implementing many methods to stabilize the GANs, they ultimately suffered from domain divergence [Thornström19].

Additionally, the paper "*How Low Can you Go?*" published by Zoe Gastelum and Timothy Shead of Sandia National Laboratories, used the SideFx Houdini software to train CV models for industrial safety and compliance (a similar field to the second use case covered in this paper, weld defect detection). Like other papers covered in this section, Gastelum's experiments also varied the ratio of real and synthetic images in the training set, but the training set did not have a lot of images available to them. Only 160 real images were available to them for training, validation, and test sets. When trained on 1000 synthetic images and tested on the 160 real images available, their model classifiers struggled to surpass 60% accuracy and severed from overfitting after only a few epochs of training [Gastelum20].

## ii. Rebalancing Techniques

A paper by Gary Weiss, et al. provides the basis for the third goal of this paper; exploring how synthetic images can rebalance an imbalanced dataset. Weiss' research studies how different techniques can mitigate the common pitfalls of training classifiers on a dataset unequal number of samples for each class. They apply oversampling, undersampling, and applying class weighted Cost-Learning techniques to differently sized datasets. In the end, they conclude that for datasets with fewer than 10,000 samples, oversampling (the act of duplicating the minority class samples) provides the greatest benefit since Cost-Learning cannot accurately approximate the value of

classes with limited data. Meanwhile for datasets with greater than 10,000 samples, Cost-Learning provides the greatest benefit since it can appropriately approximate class value. The paper mentions synthetic padding, but it is not included in their tests. [Weiss07].

### iii. Welding 101

*This section of the paper on welding techniques and best practices is provided by professional welder (and volunteer instructor) Brandon Padayao. They provided their experience while instructing the Metal 3 course offered by the OpenWorks makerspace in Baltimore, Maryland.*

The second use case in the paper focuses on training a classifier to determine if a metal weld is defective or not. This requires an understanding of what welding is and how defects can occur. In general, metal weld defects occur when one of several factors are mismanaged. Including but not limited to; the operator's speed, set voltage, set filler feed rate, differences in metallic properties between metal plates and metal filler, and contamination of the molten pool during welding. All professional welders are certified by institutions and trained to be the primary inspector to identify surface level defects. After the fact, weld inspectors also perform various tests to examine the quality of an operator's weld.

Please see the appendix A.I. *Key Terms and Concepts* section For More Information on how metal weld defects are formed and current inspection practices.

400  
401  
402  
403  
404  
405  
406  
407  
408  
409  
410  
411  
412  
413  
414  
415  
416  
417  
418  
419  
420

## 2.II. Methodology

To better understand the effect of using synthetic images to train CV models, this research uses the Unreal Engine (published by Epic Games) to render images with custom photo-realistic backgrounds and models in a variety of situations. This is done by using the following software versions. For a full list of software packages, please see the *A.II. Installed Python Software Packages* section in the appendix For More Information.

Software Package	Version
Unreal Engine	4.27
TensorFlow	2.12
Keras API	2.12

Figure 2.1: Software versions used for this research.

Images rendered in the UE are then used in the training set of two reduced published transfer learning models for separate binary classification tasks; 1) determine if a picture is of a Cat or Dog and 2) Determine if a picture of a metal weld contains a defect. The Cat\_Dog and Welding use cases represent two separate cases that require different types of graphical models with different characteristics. Animals such as cats and dogs have finer features and irregular shapes that could be difficult to capture in a 3D model (hair, whiskers, legs, etc), but provide significant variation. Alternatively, modeling a weld defect occurs on a regular shape, but requires a different type of fine detail (dust, burn marks, holes, and other nuances) to create variance.

438  
439  
440  
441  
442  
443  
444  
445  
446  
447  
448  
449

### A. Model Architectures Used

This project depends on published models already trained on the ImageNet 1000 class dataset: A) VGG16 and B) MobileNetV3-small. The existing ImageNet 1000 weights for each model act as a starting point during training via transfer learning. Both models have a pre-defined architecture, existing research, and documented performance. Approximately 60% of each model's parameters are frozen during training to not tamper with the general feature extraction layers common to image classifiers. Only the last layers of the models were trainable to fit specifically to this research's use cases.

This paper utilizes transfer learning models to take advantage of work already done in other research and keep the paper focused on the effect of synthetic data. Using published models removes uncertainty of model

450  
451  
452  
453  
454  
455  
456  
457  
458  
459  
460  
461  
462  
463  
464  
465  
466  
467  
468  
469  
470  
471  
472  
473  
474  
475  
476  
477  
478  
479  
480  
481  
482  
483  
484  
485  
486  
487  
488  
489  
490  
491  
492  
493  
494  
495  
496  
497  
498  
499

architecture design. In addition, published state-of-the-art CV models have rigorous research as to why the model's architecture exhibits state-of-the-art results compared to any custom architectures. Moreover, using publicly available published models as a starting point is a common development practice. Focusing this research on how models will practically be used will allow any findings to be applied by more developers for their own projects.

i. Originally published Models

VGG16 and MobileNetV3-small are the two models examined in this research. Both represent different ways CV models can be applied. This research focuses on how synthetic data influences the largest (VGG16) and smallest (MobileNetV3-small) models to allow developers to make the most informed design decisions.

The VGG16 model is a model with many parameters that is best ran where there are minor limitations on power consumption and computer memory, such as hosted on an on-site or cloud server. The model has a common convolutional neural network (CNN) architecture involving 16 sets of *Convolution-Pooling-BatchNorm* layers [Simonyan14].

MobileNetV3-small is a variation of the MobileNetV3 model that has the fewest number of trainable parameters. The model is specifically designed for running on mobile devices & on-chip hardware, prioritizing lower power usage, a low memory footprint, and reduced latency from model input to output. MobileNetV3's architecture utilizes Squeeze-Excite layers that advance upon residual neural networks introduced by Dr Kaiming He's ResNet model [He]. After each convolutional layer, MobileNetV3 uses Squeeze-Excite residual network blocks to pool (squeeze the output) the Conv2D layer's output, passes them through fully connected (FC) layers (excite), before scaling the output for the next Conv2D layer. These FC layers apply weights to each of the previous Conv2D layer's filters, thus prioritizing which filters contain the most important features [Howard19].

From now on, any mention of MobileNetV3 will be in reference to MobileNetV3-small.

500  
501  
502  
503  
504  
505  
506  
507  
508  
509  
510  
511  
512  
513  
514  
515  
516  
517  
518  
519  
520  
521  
522  
ii. Modifications to the published Architectures

500  
501  
502  
503  
504  
505  
506  
507  
508  
509  
510  
511  
512  
513  
514  
515  
516  
517  
518  
519  
520  
521  
522  
Due to the differences in classification scope, both transfer learning model's training settings and output architectures are changed to better align with the desired task. Originally, both published models are trained to classify 1000 various classes of objects. Since this research only requires binary classification, the model's architecture and training sets are changed. These changes include changing the number of parameters in the output layers, Learning Rate, training Batch Size, dataset size, and how many parameters are trainable, as well as others. The base architecture layers are unchanged, with the top ~60% of trainable parameters being frozen. The output layers for each model are reduced as well, shrinking the total number of parameters for each model ~85% and ~47% respectively. For a complete list of model architecture values, please see the *A.III. Model Architecture changes from Published Models* section in the appendix For Your Information.

	VGG16	MobileNetV3
General Use Case	High Accuracy Adequately available computer resources On-demand power supply	Low power consumption Low prediction latency Low memory footprint
General Architecture	Conv-AvgPool-Dropout	Squeeze & Excite
Total Parameters in the published model for 1000 class classification	138,357,544	2,554,968
Total Parameters in the transfer learning models used in this research for binary classification	21,203,521 (15.3%)	1,365,617 (53.4%)
File size of weights (.h5)	80.9 MB	5.4 MB

541  
542  
543  
544  
545  
546  
547  
548  
549  
Figure 2.2: Magnitude of size reduction in trainable parameters from the original paper to the models used in this report  
[Simonyan14][Howard19]

## B. Model Training

540  
541  
542  
543  
544  
545  
546  
547  
548  
549  
Model hyper parameters such as Learning Rate and Batch Rate are chosen via trial and error when trained on the 100:0 RSR test case. The final values are chosen based on which combinations yielded consistent convergence to the feature space's minima. The learning gradients are calculated using binary cross entropy. Initial Learning and Early Stopping techniques are used for all experiments. The VGG16 and MobileNetV3 models trained for 5 epochs before being evaluated for early stopping. During

550  
551  
552  
553  
554  
555  
556  
557  
558  
559  
560  
561  
562  
563  
564  
565  
566  
567  
568  
569  
570  
571  
572  
573  
574  
575  
576  
577  
578  
579  
580  
581  
582  
583  
584  
585  
586  
587  
588  
589  
590  
591  
592  
593  
594  
595  
596  
597  
598  
599  
the Cat\_Dog experiment, model training ends if the model does not exhibit a -0.001 change in validation loss for seven consecutive epochs. During the Welding experiment, model training ends if the model does not exhibit a +0.001 change in validation AUC performance within seven epochs. As a note, during all experiments Synthetic images are only used during model training, model validation and testing sets only use real images.

$$AUC = \frac{\#TP + \#TN}{\#FP}$$

Question to the Classifier	"Is this image a picture of a Good weld?"
Defect weld	Negative (0)
Good weld	Positive (1)

580  
581  
582  
583  
584  
585  
586  
587  
588  
589  
590  
591  
592  
593  
594  
595  
596  
597  
598  
599  
Figure 2.3: Binary classifier design for the Welding use case to minimize the instances when the models predict an image is of a Good weld but is actually a Defect.

580  
581  
582  
583  
584  
585  
586  
587  
588  
589  
590  
591  
592  
593  
594  
595  
596  
597  
598  
599  
Class weighting for error gradient calculation in the Welding use case is based on the imbalance of the Welding dataset. Since there are more *Defect* welds than *Good* welds in the dataset, error from classifying *Good* welds is greater than classifying *Defects*. This means that the model is more sensitive to misclassifying a weld as Good. Thus, encouraging the model to be more accurate when it says an image is a *Good* weld. However, it leads models to skew towards defaulting every weld to be a *Defect*. Experiment 4 uses class weighting for its results.

## 600 601 602 603 C. Image Augmentation 604

605 During model training, the training set is expanded using a  
606 variety of 2D image augmentation techniques. Each of  
607 these techniques has a random chance of being applied  
608 and can stack on top of one another (an image could have  
609 all, none, or some of the augmentations). These techniques  
610 include standard augmentation such as horizontal &  
611 vertical flips, rotation, shifts,  
612 brightness change, shear, zoom, and  
613 channel shift. The model batch size and epoch  
614 steps are set so that the models train on all available  
615 training images every epoch. Meaning that every epoch,  
616 the entire training set is randomly augmented and fed into  
617 the model to calculate error gradients. Theoretically  
618 resulting in a more generalized model since every epoch is  
619 a slightly ‘brand new’ training set. For a list of applied 2D  
620 image augmentations, their available ranges, and  
621 reasoning for use. Please see the *A.IV. 2D Image*  
622 *Augmentations* section in the appendix For More  
623 Information.



624  
625  
626  
627  
628  
629  
630 *Figure 2.4: Example image augmentations for three separate*  
631 *epochs. The left image would be used for training in epoch 1, the*  
632 *middle in epoch 2, and the right in epoch 3.*

633 *Augmentations: (left) rotate right, shift up, shift right. (middle)*  
634 *zoom, shift up, shift right, brightness up. (right) vertical flip,*  
635 *rotate left.*

## 636 D. Dataset Formation 637

638 The datasets used for all experiments are a combination of  
639 real-world images from various public internet sources  
640 and synthetic screenshot images from Unreal Engine  
641 v4.27. Not every real image gathered is used in this  
642 paper’s experiments. The process for removing real data  
643 and the method of sampling is covered for each individual  
644 use case further in this paper. Used real images have  
645 varied original dimensions, all synthetic images are  
646 captured at the same dimensions based on the chosen  
647 camera in UE (782x440 for Cat\_Dog images or 853x480  
648 for Welding images). Regardless, all images are resized to  
649 224x224 to match input dimensions of the VGG16 and  
MobileNetV3 models.

## 650 651 652 ii. Collecting Synthetic Images via the UE Python API 653 654

655 UE contains a command line that can execute python  
656 scripts (if enabled). Executed python scripts can send UE  
657 commands and get information from the opened UE Editor  
658 that called the script. This leads to two notable computer  
659 limitations: installed python packages and threading.  
660

661 UE runs these scripts in the same environment as the  
662 computer program. Therefore, you must install all needed  
663 python packages to the same environment that you run the  
664 UE instance. Meaning that you will have package  
665 compatibility errors if you develop the scripts in a virtual  
666 environment but run UE from your desktop.  
667

668 Regarding threading, these scripts run on the same  
669 software threads as the UE Editor and use the python  
670 packages installed on the environment that UE is opened  
671 with [DowlingSoka]. Technical Art Engineer, Ryan  
672 DowlingSoka, refined a python class that can execute  
673 python code every ‘tick’ of the UE Editor process through  
674 a software callback. Every tick, the code in the callback  
675 executes. To gather all the needed synthetic image  
676 variations for the Cat\_Dog and Welding classes, a script  
677 based on this python class is created to run a predefined  
678 state-machine, that increments after every callback. Every  
679 ~30 ticks, the state-machine either A) sets the UE Editor  
680 environment (setting model positions, camera+lighting  
681 positions, texture properties, lighting intensity, etc) or B)  
682 takes a screenshot. At an execution rate of one callback  
683 action every ~30 ticks, the script can generate 80 images a  
684 minute if not interrupted. This rate is chosen to give each  
685 callback enough time to execute all necessary UE  
686 commands prior to the prompting of the next callback.  
687 (since both the script and UE Editor are on the same  
688 thread). To learn more about base python class and writing  
689 your own class, please see the *A.V. Collecting Synthetic*  
690 *Images via the UE Python API* section in the appendix For  
691 More Information.  
692

## 693 iii. Use Case: Cat\_Dog 694

695 The datasets used in Cat\_Dog experiments are sampled  
696 from the real Cats\_Dogs ImageNet dataset and four  
697 different collections of synthetic data using various free to  
use Cat & Dog .fbx models and UV textures. The binary  
698 classes are {Cat, Dog}. The real dataset contains 12500  
699 cats and 12500 dog images. The dataset is referenced  
700 directly from the Tensorflow datasets repository,  
701 distributed by Microsoft [Elson07]. All 25000 collected  
702 real images are available for sampling, none of the images  
703 are removed or edited.  
704

iv.1. Synthetic Data

Over the course of two collections (4 runs total), 47864 screenshots of 3D cat and dog models are collected against various backdrops in a variety of camera angles and lighting conditions. The used cat and dog .fbx 3D models are free to download from sites such as Turbosquid and Free3D. It is possible to purchase more advanced models with higher polygon, triangle counts and textures. The backdrop presets are different images and textures put behind the 3D model to represent background variation. Each backdrop preset is operated via a UE parameter that feeds into a software multiplexer in the UE material instance to decide which preset is shown. For all 3D models, authors, and credits, please see the appendix section *A.VI. Cat\_Dog – Collecting Synthetic Images* for more information.

	Collection 1		Collection 2	
	Run 1	Run 2	Run 3	Run 4
Collected Image Variations	6672 <sup>^</sup>	2753 <sup>^</sup>	27478	9881
# 3D Models	23	23	23	23
# of Camera Angles	28	10	72*	72*
# of Lighting Settings	1	3	4	1
# of Backdrop presets	4**	4**	6*	6*

*Figure 2.5: Summary of variation in screenshots collected for Cat\_Dog use case.*

<sup>^</sup> During collection, some images were lost due to nuances of running the python script through the UE Editor.  
\*Run 3 & 4 are the same except for the ‘unlit’ lighting setting.  
\*\*Backdrops for Run 1 & 2 are the same

Screenshots are collected via a python script that iterates through all possible combinations of these settings, as previously mentioned. During collection, the script moves a cat/dog model to a point in front of a backdrop at a specific lighting (a SunSky light modeled off earth’s sun), at hardcoded camera angles. Once all camera angles have been captured, the cat/dog model is moved out of the frame, then the next cat/dog model is moved to the point. The process repeats all set variations (cat and dog models, all backdrop presets, lighting intensities) in the UE test environment until a screenshot is taken of every variation. Table 2.5 summarizes the variation of this process.



*Figure 2.6: Example screenshots taken at different angles, models, and backdrops that are then possibly used to train the CV models prior to 2D image augmentation.*

iv. Use Case: Welding

The real images sampled for the welding experiments come from a variety of public sources. Several public repositories of real weld images are available online, but every dataset has various flaws. Flaws include multiple datasets using the same images, datasets including a mix of original and augmented images, or datasets that contain more than one weld (*Defect* or *Good*) in a single image. Dataset specifics are discussed in the next section. The synthetic data is generated from edits to a purchased UE material package from the Epic Games store. One of these materials in the package is applied to a flat mesh plate where various texture map bitmap files are edited to mimic three weld defects: porosity, crack, and incomplete welds. In general, all real and synthetic images are labeled as *{Defect, Good}*. The classifier does discern which defect is present, only that the weld has a defect.

v.1. Real Data

Through five sources, 2221 {624 Good, 1597 Defect} real images of various levels of quality form the real dataset. Out of the 5813 real images found from the five datasets, only 2221 are used. Most of the real images are removed to avoid repeat images, images with irrelevant classes, and images of multiple welds so close together that each weld could not be practically cropped. For more details on how each of the five datasets were used, please see the *A.VII. Welding – Dataset Formation* section in the appendix For More Information



#### iv.1. Synthetic Data

Over the course of two collections (4 runs total), 47864 screenshots of 3D cat and dog models are collected against various backdrops in a variety of camera angles and lighting conditions. The used cat and dog .fbx 3D models are free to download from sites such as Turbosquid and Free3D. It is possible to purchase more advanced models with higher polygon, triangle counts and textures. The backdrop presets are different images and textures put behind the 3D model to represent background variation. Each backdrop preset is operated via a UE parameter that feeds into a software multiplexer in the UE material instance to decide which preset is shown. For all 3D models, authors, and credits, please see the appendix section *A.VI. Cat\_Dog – Collecting Synthetic Images* for more information.

	Collection 1		Collection 2	
	Run 1	Run 2	Run 3	Run 4
Collected Image Variations	6672 <sup>^</sup>	2753 <sup>^</sup>	27478	9881
# 3D Models	23	23	23	23
# of Camera Angles	28	10	72*	72*
# of Lighting Settings	1	3	4	1
# of Backdrop presets	4**	4**	6*	6*

*Figure 2.5: Summary of variation in screenshots collected for Cat\_Dog use case.*

<sup>^</sup> During collection, some images were lost due to nuances of running the python script through the UE Editor.

*\*Run 3 & 4 are the same except for the 'unlit' lighting setting.  
\*\*Backdrops for Run 1 & 2 are the same*

Screenshots are collected via a python script that iterates through all possible combinations of these settings, as previously mentioned. During collection, the script moves a cat/dog model to a point in front of a backdrop at a specific lighting (a SunSky light modeled off earth's sun), at hardcoded camera angles. Once all camera angles have been captured, the cat/dog model is moved out of the frame, then the next cat/dog model is moved to the point. The process repeats all set variations (cat and dog models, all backdrop presets, lighting intensities) in the UE test environment until a screenshot is taken of every variation. To learn more about the nuances of this image collection and more details on the different backdrops, please see the *A.VI. Cat\_Dog - Collecting Synthetic Images* section in the appendix For More Information.



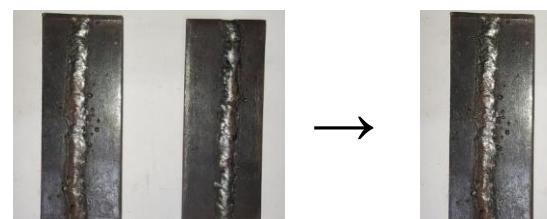
*Figure 2.6: Example screenshots taken at different angles, models, and backdrops that are then possibly used to train the CV models prior to 2D image augmentation.*

#### iv. Use Case: Welding

The real images sampled for the welding experiments come from a variety of public sources. Several public repositories of real weld images are available online, but every dataset has various flaws. Flaws include multiple datasets using the same images, datasets including a mix of original and augmented images, or datasets that contain more than one weld (*Defect* or *Good*) in a single image. Dataset specifics are discussed in the next section. The synthetic data is generated from edits to a purchased UE material package from the Epic Games store. One of these materials in the package is applied to a flat mesh plate where various texture map bitmap files are edited to mimic three weld defects: porosity, crack, and incomplete welds. In general, all real and synthetic images are labeled as  $\{\text{Defect}, \text{Good}\}$ . The classifier does discern which defect is present, only that the weld has a defect.

## v.1. Real Data

Through five sources, 2221 {624 Good, 1597 Defect} real images of various levels of quality form the real dataset. Out of the 5813 real images found from the five datasets, only 2221 are used. Most of the real images are removed to avoid repeat images, images with irrelevant classes, and images of multiple welds so close together that each weld could not be practically cropped. For more details on how each of the five datasets were used, please see the *A.VII. Welding – Dataset Formation* section in the appendix For More Information



*Figure 2.7: Example of an image from dataset 01 that shows two welds in one picture [welddefect]. For some images, the researcher copied the image and cropped both images into separate images. Resulting in two separate image data points.*

## v.2. Synthetic Data

Images for the Welding synthetic dataset are generated from taking screenshots of points of interest on used UE mesh metal plate actor. Synthetic images are generated from random camera angles, light offsets, metal color, and texture map presets. A ‘map preset’ is a combination of Metallic, Normal, Roughness, and Ambient Occlusion texture bitmap files specific to a certain welding variation. The sum of these four bitmaps creates the imperfections and details on the metal plate. Weld plate variations include 1) no defects, 2) imperfections but not enough to be defective, 3) porosity defects, 4) crack defects, and 5) incomplete weld defects.



Figure 2.8: (left) Image of the plate in which screenshots are collected (specifically shown is the purchased ‘Metal\_Weld’ material with the UE starter content ‘steel’ basecolor). (right) Side view of the metal plate used to collect images. Note: all texture map edits are to a 2D plane. Any perceived height is an optical illusion from the Normal and Ambient Occlusion maps.

Like the Cat\_Dog synthetic dataset; a total of 8160 synthetic images are collected to provide additional variation to the real dataset. Variation comes from the plate basecolor, texture map changes, camera angle and point light offset at several points. Texture map changes are applied to a base UE material purchased from the Epic Games Unreal Store create the weld imperfections (bitmap and targa files are 4096x4096 pixels) [4k Metal pack]. Six metal colors, 5 map presets and 68 points of interest are cycled through. The defect dimensions are randomly chosen from that dimension’s set range of values.

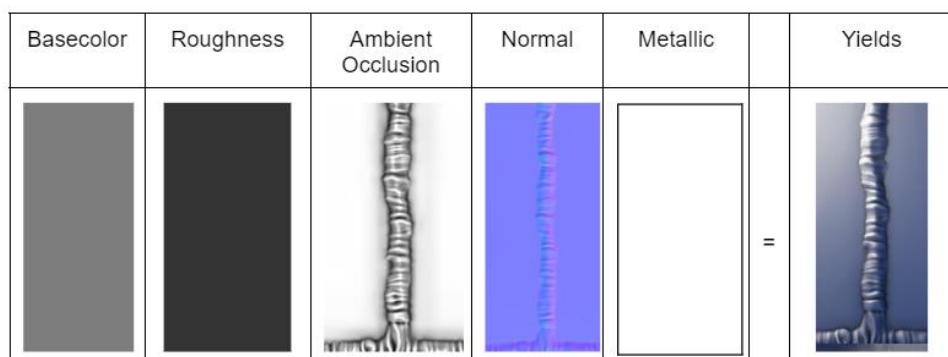


Figure 2.10: Illustration of how texture maps combine to create a 2D weld on the metal plate.

### v.3. Weld Imperfection Design

Three welding defects are modeled as a part of this research: Porosity, Crack, and Incomplete Welding. All three defects have different changes to the four main texture maps in a map preset. The following section details the researcher's interpretation of what visually defines each defect, what variables are used to create the UE Texture. In addition, non-defect versions of the three defects are defined for map preset 02. These non-defect versions are labeled as *Good* during model training and considered as imperfections that mimic a defect at a reduced scale or without certain characteristics.

A part of this design is the use of randomness to create a unique weld imperfection (defect or not) at each of the points of interest. Each of the welding defects uses random numbers to subtly change how a defect is designed. Every weld defect is a different size, length, and has different material properties by adding a level of noise to the resulting dimensions or texture map RGB colors. The randomness range for each dimension and RGB color pixel is set for each defect. This randomness helps add variation and make each weld modification unique. To learn more about what is specifically done to create each weld texture, please see the *A.VII. Welding – Dataset Formation* section in the appendix For More Information.

Defect	Visual Description	Real Weld Defect	Synthetic Weld Defect
Porosity	<p>Surface level Elliptical holes caused by trapped gas rising to the surface of the molten pool as a weld cools.</p> <p>Random Variables</p> <ul style="list-style-type: none"> <li>• xy-radius</li> <li>• xy-offset from weld center.</li> </ul> <p>Non-Defect = Smaller ellipse</p>		
Crack	<p>Surface level splitting or cutting of the metal, creating triangular separation compromising weld strength.</p> <p>Random Variables:</p> <ul style="list-style-type: none"> <li>• Length</li> <li>• Number of angled segments</li> <li>• Offset of each angle segment</li> </ul> <p>Non-Defect = a shorter length, lower offset per segment and fewer segments.</p>		
Incomplete Weld	<p>A missing area of filler on a welding line. Revealing the split between metal plates</p> <p>Random Variables:</p> <ul style="list-style-type: none"> <li>• Width of gap</li> <li>• Depth of gap towards the seam.</li> <li>• Seam size</li> </ul> <p>Non-Defect = Smaller gap width and not enough gap depth to reveal the plate seam</p>		

Figure 2.11: Summary table of each of the three weld defects modeled for map presets 02-05 in the synthetic Welding dataset.

### 3.III. Results

Five total experiments are performed to understand how synthetic data affects CV models in different ways. By design, different synthetic datasets in different experiments can reveal nuances to how synthetic images can be used to train CV models. The two types of experiments (RSR & Rebalancing) are documented below. Both experiments train the transfer learning models on different variations of training data, specifically how much and what synthetic data is used. The reported data results are the average of three training runs for each test within an experiment. Three runs are believed to give enough variation in training results to uncover the desired trends.

RSR experiments explore how different proportions of real to synthetic images in the training set effects model performance. The RSRs tested range from 100:0 (all real) to 0:100 (all synthetic). The sole Rebalancing experiment compares model performance when the training dataset is imbalanced to after it is balanced by four methods. This is an expansion of Dr Gary Weiss' 2007 experiment [Weiss07].

Experiment	RSR Experiments	Rebalancing Imbalanced Dataset Experiment
Goal	How do different RSRs affect CV Models	Can synthetic data be used to rebalance an imbalanced dataset
Cat_Dog	3x	Not performed
Welding	1x (Limited)	1x

Figure 3.1: Outline of the two types of experiments performed for Cat\_Dog (Experiments 1-3) and Welding (Experiments 4 & 5).

All five experiments and dataset variations use an 80/10/10 training/validation/test split to maximize training samples while attempting to provide enough testing samples (a constraint with small datasets). Furthermore, all images in the validation and test sets are only real images, no synthetic images are used in the validation or test sets. Synthetic images are only used in the training set.

#### A. Cat\_Dog Experiments

Three RSR experiments are performed with a balanced dataset of 10,000 total images. 1) model performance trained on different RSR proportions of real and synthetic images when with only 9425 synthetic images are available to be sampled, 2) model performance with different RSR values with 46784 synthetic images to sample, and 3) baseline performance of models with only real images but decreasing dataset sizes to compare against 1) & 2).

All tests are performed with balanced training/validation/test sets. The training set is a mix of real images and synthetic images from the Cat\_Dog synthetic dataset. Model accuracy and loss are the primary focus of the Cat\_Dog use case results.

Experiments 1) and 2) investigate how different levels of synthetic data availability can affect model performance. Different Synthetic Utilization Ratios (SURs) could affect the level of variance shown to a model with the chosen synthetic samples. The third test highlights the effect of padding real data with synthetic data. The displayed bar chart showing the differences between training datasets with or without synthetic data highlights if adding synthetic images improved model metrics. The definitions for all experiments & tests are outlined in the appendix. Please see the A.IX. *Results* section in the appendix For More Information.

RSR	Total Images	Number of Real Training Images	Number of Synthetic Training Images	Number of real VALIDATION images	Number of real TEST images
100:0	10000	8000	0	1000	1000
75:25	10000	6000	2000	1000	1000
50:50	10000	4000	4000	1000	1000
25:75	10000	2000	6000	1000	1000
10:90	10000	800	7200	1000	1000
05:95	10000	400	7600	1000	1000
0:100	10000	0	8000	1000	1000

Figure 3.2: Outline of tests for Experiment 1-3.

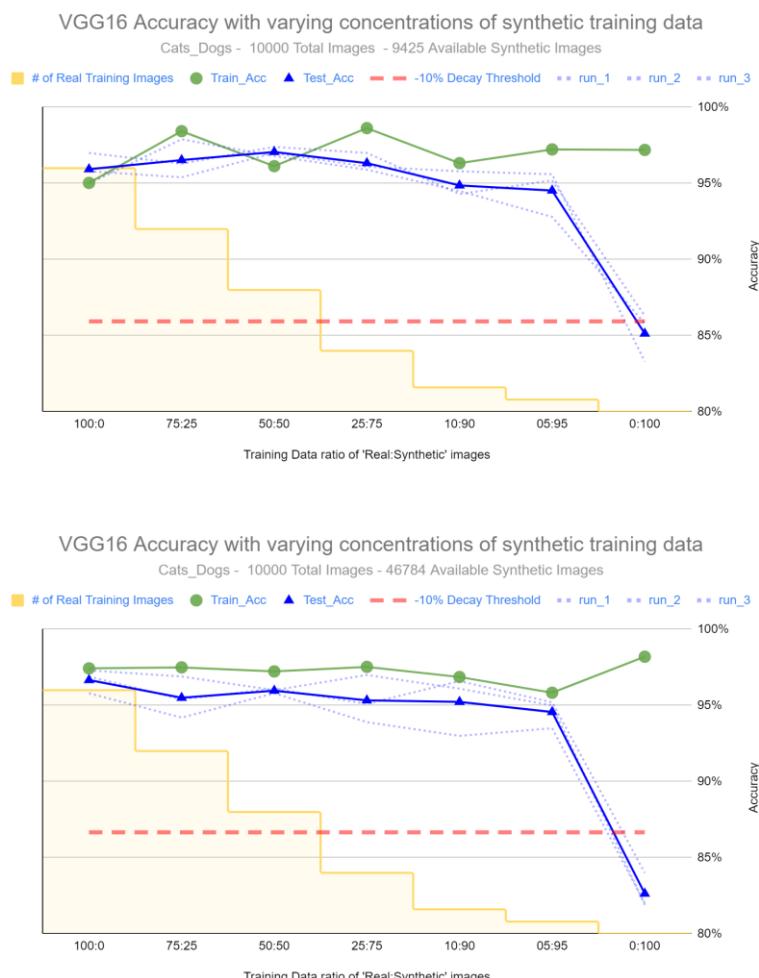
1100  
 1101  
 1102  
 1103 i. Experiment 1 – RSR with 9425 Synthetic Images  
 1104

1105 When the VGG16 model is trained on various RSRs with  
 1106 only 9425 synthetic images to sample, the test accuracy  
 1107 and test loss performance begins to degrade at RSR 25:75.  
 1108 The model's accuracy increases to its max at RSR 75:25 at  
 1109 96.5% (+0.6%) before decreasing for every subsequent  
 1110 RSR test. The lowest test accuracy occurs after a steep  
 1111 decrease in performance from RSR 05:95 to 0:100 with  
 1112 85.9% (-10.8%) binary accuracy. The model's test loss  
 1113 follows a similar pattern to the test accuracy. The  
 1114 performance is consistent before steadily increasing when  
 1115 RSR lowers past 25:75. Meanwhile, the model's training  
 1116 accuracy remains relatively flat throughout the experiment  
 1117 even as RSR lowers to 0:100. These results show that with  
 1118 a 1.4% decrease in test accuracy, a user can use 7600 less  
 1119 real training samples if RSR 05:95 is used.

1120  
 1121  
 1122  
 1123  
 1124  
 1125  
 1126  
 1127  
 1128  
 1129  
 1130  
 1131  
 1132  
 1133  
 1134  
 1135  
 1136  
 1137  
 1138  
 1139  
 1140  
 1141  
 1142  
 1143  
 1144  
 1145  
 1146  
 1147  
 1148  
 1149 Also, when RSR is at 75:25, the test accuracy is greater  
 1150 than the training accuracy indicating a positive level of  
 1151 generalization. However, the significant difference  
 1152 between training and test loss implies that the VGG16  
 1153 model is still overfit.

1154  
 1155  
 1156  
 1157  
 1158  
 1159  
 1160  
 1161  
 1162  
 1163  
 1164  
 1165  
 1166  
 1167  
 1168  
 1169  
 1170  
 1171  
 1172  
 1173  
 1174  
 1175  
 1176  
 1177  
 1178  
 1179  
 1180  
 1181  
 1182  
 1183  
 1184  
 1185  
 1186  
 1187  
 1188  
 1189  
 1190  
 1191  
 1192  
 1193  
 1194  
 1195  
 1196  
 1197  
 1198  
 1199 Regarding test accuracy variance, it showed a relatively  
 1158 positive parabolic shape. This means that the test accuracy  
 1159 variance peaks at RSR 0:100 and 100:0 respectively but is  
 1160 at a minimum during RSR 50:50 and 25:75. This should  
 1161 be inspected further and, in more experiments, to see if the  
 1162 trend holds. At worst, this is noise due to too few data  
 1163 points, at best this is a sign of the VGG1 model learning  
 1164 the latent features of the synthetic data and applying them  
 1165 to provide a more stable test accuracy performance.

1166  
 1167  
 1168 Please see table A.IX.R4 in the appendix section A.IX.  
 1169 Results For More Information.



1150  
 1151  
 1152  
 1153  
 1154  
 1155  
 1156  
 1157  
 1158  
 1159  
 1160  
 1161  
 1162  
 1163  
 1164  
 1165  
 1166  
 1167  
 1168  
 1169  
 1170  
 1171  
 1172  
 1173  
 1174  
 1175  
 1176  
 1177  
 1178  
 1179  
 1180  
 1181  
 1182  
 1183  
 1184  
 1185  
 1186  
 1187  
 1188  
 1189  
 1190  
 1191  
 1192  
 1193  
 1194  
 1195  
 1196  
 1197  
 1198  
 1199 Figure 3.3: (Top) accuracy and (bottom) loss graphs of VGG16's performance over different RSR settings with the Cat\_Dog use case.

The MobileNet model's performance is slightly different from VGG16 when sampling from 9425 synthetic images. The model's test accuracy decreases consistently until 05:95 where it steeply decreases to ~50%. The training accuracy remains flat after an initial decrease, diverging from test accuracy. The model's test loss appears slightly lower at RSR 75:25 before increasing for the remaining tests.

The combination of flat training accuracy while decreasing test accuracy and widening training & test loss gap signals that the MobileNetV3 model is very overfit to the training data. This is significant because at lower RSRs, the model is overfitting to the synthetic data and cannot apply the variations to the real-world test samples. In addition, unlike the VGG16 model, the MobileNet model's test accuracy degrades past 10% from training on only real images before the 0:100 RSR.

MobileNetV3-Small Accuracy with varying concentrations of synthetic training data  
Cats\_Dogs - 10000 Total Images - 9425 Available Synthetic Images



MobileNetV3-Small Loss with varying concentrations of synthetic training data  
Cats\_Dogs - 10000 Total Images - 9425 Available Synthetic Images

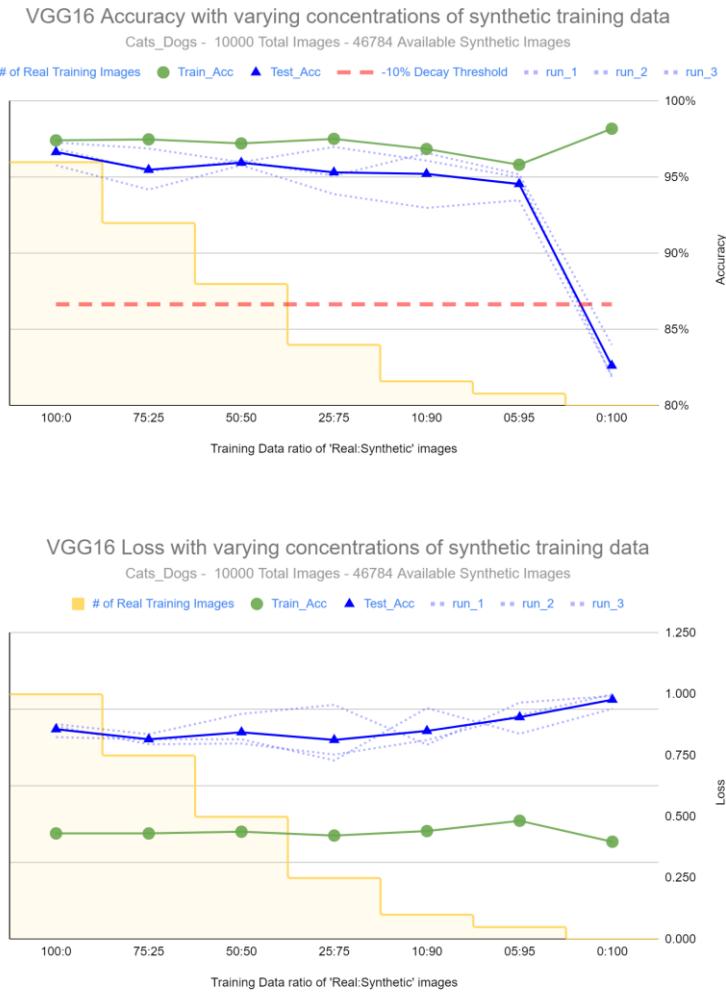


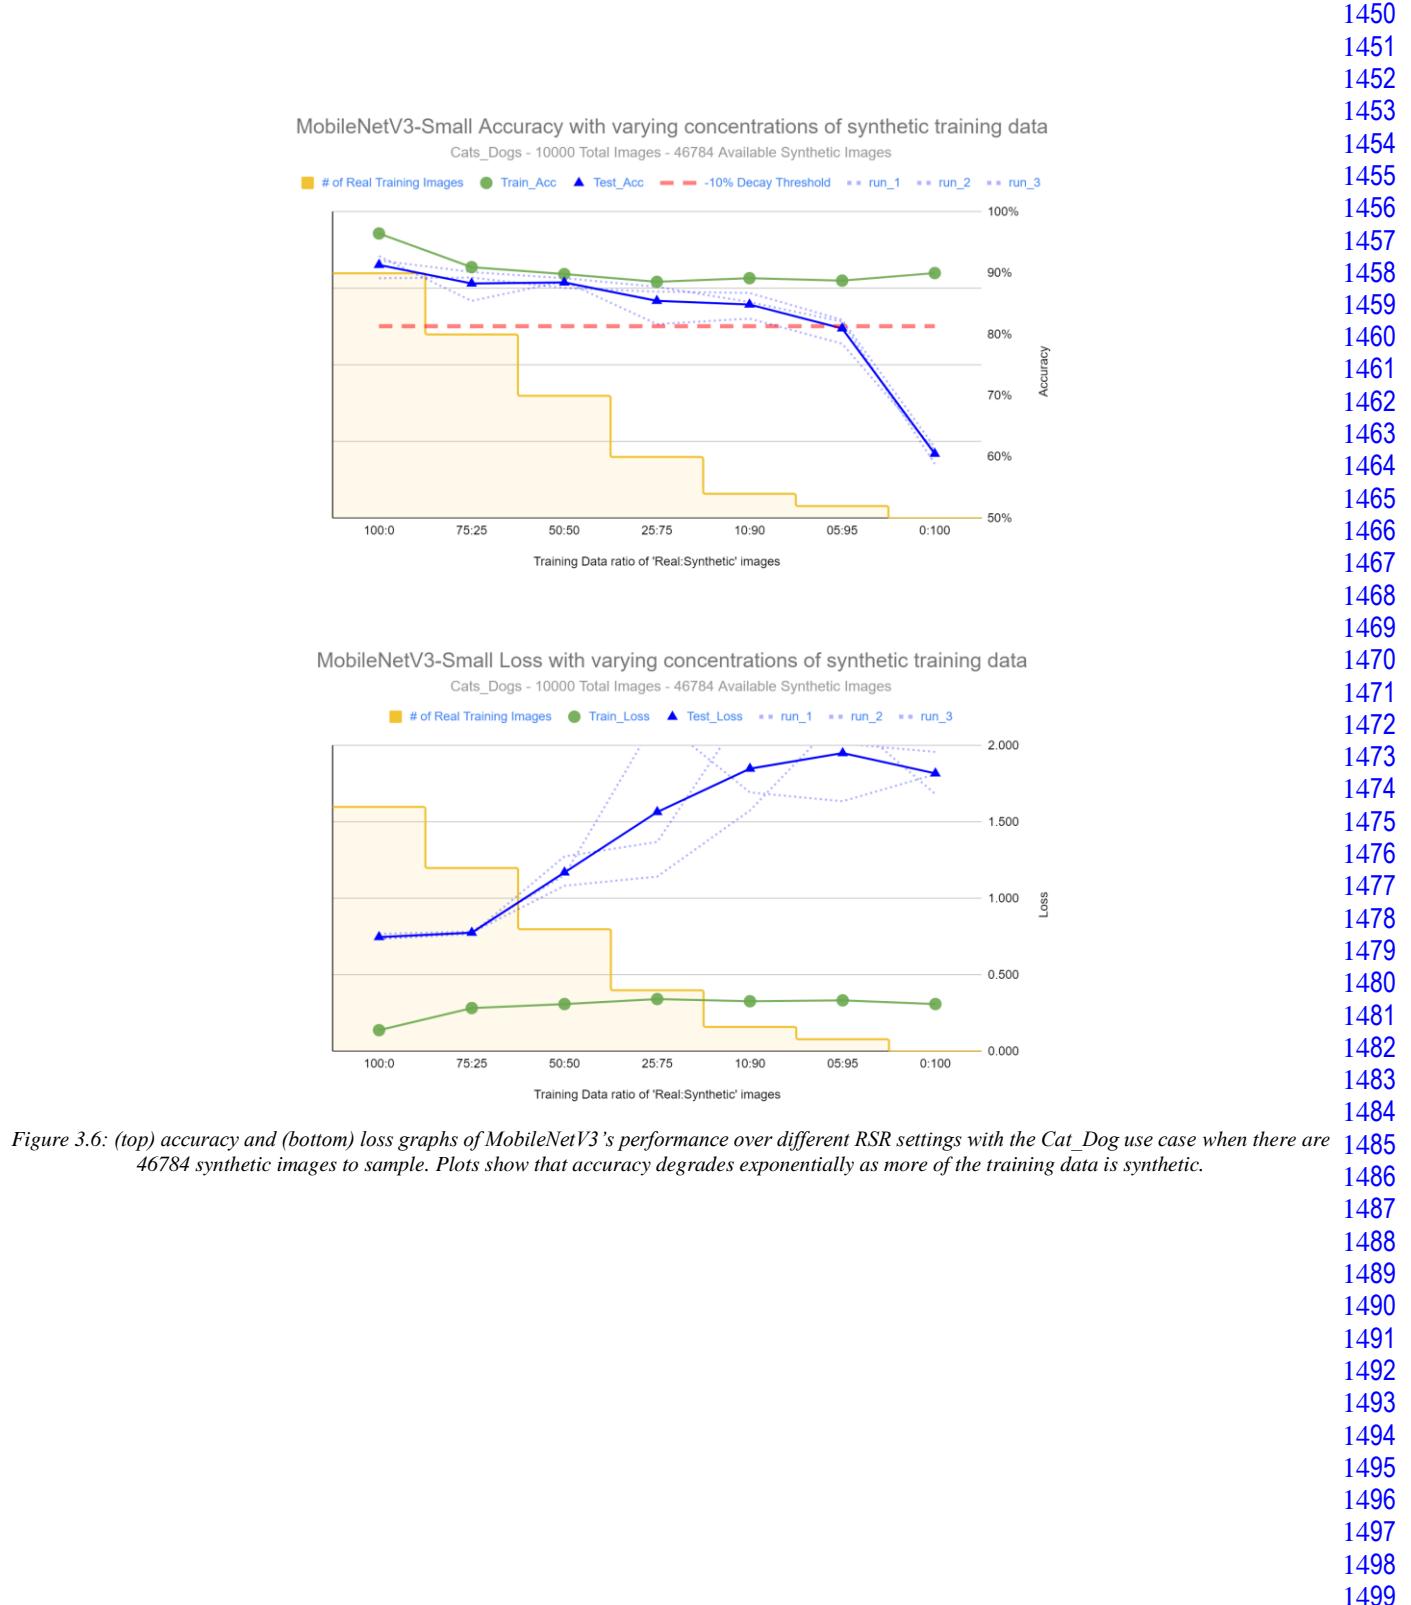
Figure 3.4: (top) accuracy and (bottom) loss graphs of MobileNetV3's performance over different RSR settings with the Cat\_Dog use case.

1300  
 1301  
 1302  
 1303 ii. Experiment 2 – RSR with 46784 Synthetic Images  
 1304  
 1305  
 1306  
 1307  
 1308  
 1309  
 1310  
 1311  
 1312  
 1313  
 1314

Compared to Experiment 1 above, increasing the number of synthetic images for sampling does not appear to significantly change model performance or observations expressed in Experiment 1 above. This implies that having image latent variation does not scale with the number of camera angles taken. Other changes like backdrops, models, and lighting require just as many variations. This is because the Cat\_Dog gathered dataset primarily used varying camera angles around a center location during the collection. The number of backdrops and lighting changes are limited.

1305 Some slight differences include the VGG16 model not being greater than training accuracy and its test accuracy dropping much more significantly from RSR 05:95 to 0:100. MobileNetV3 shows a slightly lower training accuracy to test accuracy gap, but the test Loss is significantly higher when there are 46784 available synthetic samples compared to only 9425.  
 1306  
 1307  
 1308  
 1309  
 1310  
 1311  
 1312  
 1313  
 1314  
 1315  
 1316  
 1317  
 1318  
 1319  
 1320  
 1321  
 1322  
 1323  
 1324  
 1325  
 1326  
 1327  
 1328  
 1329  
 1330  
 1331  
 1332  
 1333  
 1334  
 1335  
 1336  
 1337  
 1338  
 1339  
 1340  
 1341  
 1342  
 1343  
 1344  
 1345  
 1346  
 1347  
 1348 Figure 3.5: (top) accuracy and (bottom) loss graphs of VGG16’s performance over different RSR settings with the Cat\_Dog use case when there are 46784 synthetic images to sample. Shows that accuracy behaves approximately the same to when 9425 synthetic images are available. Implies that there is no performance gain by having more samples when those samples are collected in a ridged way.  
 1349

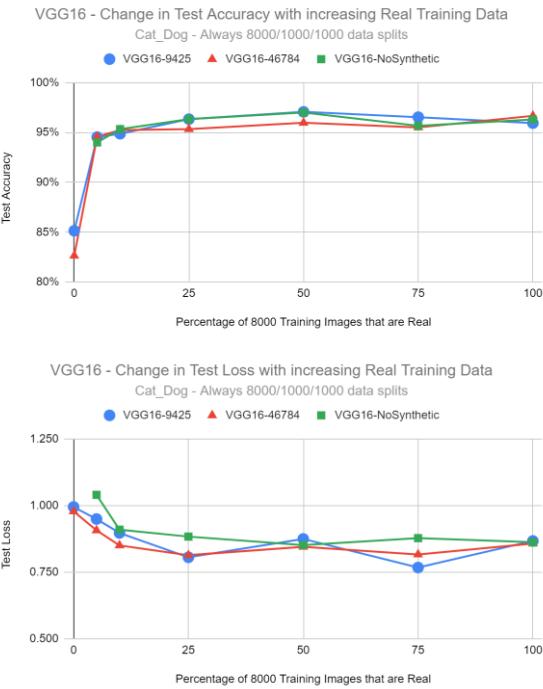




1500  
 1501  
 1502  
 1503 iii. Experiment 3 – RSR. Mixed vs Smaller Datasets  
 1504  
 1505  
 1506  
 1507  
 1508  
 1509

When comparing results of models trained solely off datasets with only real images versus datasets with a mix of real and synthetic images, the results below show different trends for test accuracy and test loss depending on the model.

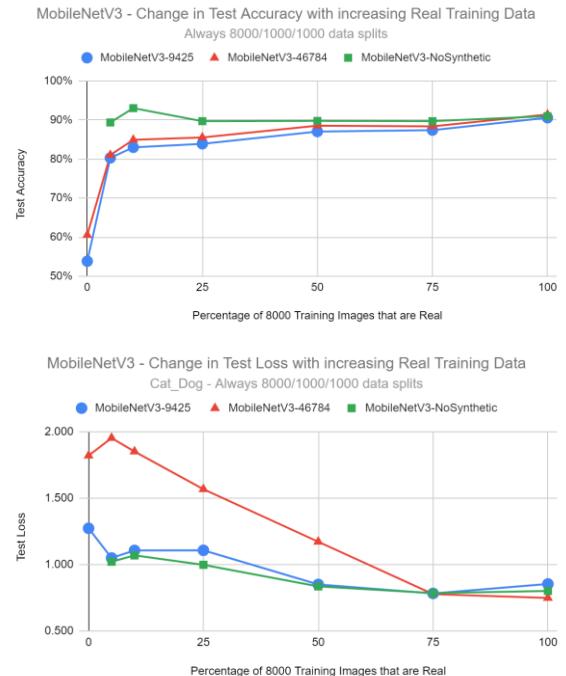
For the VGG16 model, the results do not show a meaningful change in test accuracy regardless of synthetic availability. However, having a mixed dataset of real and synthetic images lowered the test loss for nearly all RSR tests. For the 05:95 RSR test (5% in Figure 3.7), these trends highlight how synthetic image quality can be judged. In the 05:95 RSR test, the results show that padding 400 real images with 7600 synthetic images of cats and dogs yielded a test accuracy increase of only ~0.6% and a test loss decrease of approximately 0.1.



1510  
 1511  
 1512  
 1513  
 1514  
 1515  
 1516  
 1517  
 1518  
 1519  
 1520  
 1521  
 1522  
 1523  
 1524  
 1525  
 1526  
 1527  
 1528  
 1529  
 1530  
 1531  
 1532  
 1533  
 1534  
 1535  
 1536  
 1537  
 1538  
 1539  
 1540  
 1541  
 1542  
 1543 Figure 3.7: VGG16 comparison of three different training sets with their (top) test Accuracy and (bottom) test Loss. 'VGG16-NoSynthetic' uses a training set made of only real images but ranges from 400 images (5%) to 8000 (100%) images. The plots show that, compared to using a smaller training set (green), using additional synthetic images reliably lowers test Loss without a decrease in test Accuracy.  
 1544  
 1545  
 1546  
 1547  
 1548  
 1549

1550  
 1551  
 1552  
 1553  
 1554  
 1555  
 1556  
 1557  
 1558  
 1559  
 1560  
 1561  
 1562  
 1563  
 1564  
 1565  
 1566  
 1567  
 1568  
 1569  
 1570  
 1571  
 1572  
 1573  
 1574  
 1575  
 1576  
 1577  
 1578  
 1579  
 1580  
 1581  
 1582  
 1583  
 1584  
 1585  
 1586  
 1587  
 1588  
 1589  
 1590  
 1591  
 1592  
 1593  
 1594  
 1595  
 1596  
 1597  
 1598  
 1599

The MobileNetV3 model did not perform comparably when trained on only real images versus a mix of real and synthetic images. Training on only real images yielded the best test accuracy and test loss metrics for all RSR test cases. The results show that accuracy degraded somewhat exponentially as the percentage of synthetic images in the training set increased. In the big picture, this experiment describes that at RSR 05:95, datasets of only 400 real images yielded higher test accuracy and comparable loss to a dataset with those 400 real images plus an additional 7600 synthetic images.



1590  
 1591  
 1592  
 1593  
 1594  
 1595  
 1596  
 1597  
 1598  
 1599  
 Figure 3.8: MobileNetV3-small comparison of three different training sets with their (top) test Accuracy and (bottom) test Loss. 'MobileNetV3-NoSynthetic' uses a training set made of only real images but ranges from 400 images (5%) to 8000 (100%) images. The plots show that, compared to using a smaller training set (green), using additional synthetic images has a negligible effect on test Loss, but a negative effect on test Accuracy.

## A. Welding Experiments

Experiments 4 & 5 are performed with limited amounts of real weld data and the collected synthetic Weld dataset; 4) model performance trained on different RSR proportions of real and synthetic images while using the entire imbalance weld dataset, and 5) model performance changes based on different rebalancing methods. Model ROC AUC (receiver operating characteristics area under the curve), loss, and classifier confusion matrix percentages are the primary focus of the Welding use case results. As a reminder, the *Defect* class is considered Negative, and the *Good* weld class is considered Positive for these experiments. This is done so the models can use Test AUC as an early stopping metric during training.

Question to the Classifier	"Is this image a picture of a Good weld?"
Defect weld	Negative (0)
Good weld	Positive (1)

Figure 3.9: repeat of figure 2.3

Experiment 4 is like Experiments 1-3 for the Cat\_Dog use case. The major difference is that the entire real weld dataset (1400 Defect & 600 Good weld samples) is used for every test and the class imbalance from the real weld dataset is maintained for every test. Also, both use case experiments have the same goals; how the welding use case is affected by different RSRs and how that differs from the Cat\_Dog use case. However, there are a few key differences.

The first difference is that the RSR experiment is more limited due to the lack of available *Good* weld samples. Second, the training set is imbalanced at 7:3 Defect-to-*Good* class, but the test set is always balanced. Since the training set is imbalanced, the model class weights during training are inversely proportionally weighted, even if the test set is balanced. Therefore, the error gradient for *Defect* weld samples are weighted at 0.34 while the *Good* weld sample error gradients are weighted at 0.66. The imbalance class weighting encourages the model to not overlearn the majority class in the dataset (the potential consequence being that the model thinks all images are defective just because it trained on mostly *Defect* images). Finally, as the RSR decreases (skewing the training set towards mostly synthetic images), the number of real samples in the validation and test set increases. Adding Synthetic data allows for real samples to be moved to the validation and test sets to maintain the 80/10/10 training-validation-test data splits.

Data split	Balanced?	Type of Images
Training set	Imbalanced	Mix
Validation set	Balanced	Only Real
Test set	Balanced	Only Real

Figure 3.10: Summary of dataset splits for Experiment 4.

Experiment 5 expands upon research done by Weiss, et al investigating different methods of sampling on decision trees with imbalance classification data. Weiss concluded that when rebalancing 8 chosen datasets with three types of rebalancing methods (cost-sensitive learning, upsampling, and undersampling); cost-sensitive learning yielded the lowest cost for most datasets with over 10,000 samples, while oversampling yielded the lowest cost for datasets with less than 10,000 samples [Weiss07]. This second experiment expands Weiss' research to rebalance datasets with less than 10,000 samples with synthetic data. Specifically compare how padding the minority class with synthetic data compares against all their previous methods. In particular, the 'Padding with Synthetic' case balances the used dataset by adding 600 synthetic *Good* Weld samples to yield a dataset with 2400 samples (1200 samples per class). All 600 *Good* class synthetic images are screenshots from the Unreal Engine using map preset 01 (no changes to the UE material) & 02 (non-defect imperfections). This experiment focuses on metrics related to confusion matrix performance and model training (such as accuracy & loss).

Balancing Method	Total Number of Images	TP Cost	TN Cost	FP Cost	FN Cost
No Change	1800	0	0	0.5	0.5
Cost Learning	1800	0	0	0.667	0.334
Undersample	1200	0	0	0.5	0.5
Oversample	2400	0	0	0.5	0.5
Synthetic Padding	2400	0	0	0.5	0.5

Figure 3.11: Summary of rebalancing methods used in Experiment 5.

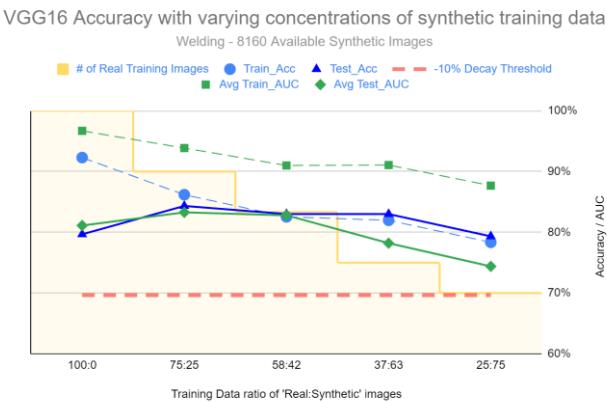
$$\text{Total Cost per Sample} = \frac{TP * \text{Cost}_{TP} + TN * \text{Cost}_{TN} + FP * \text{Cost}_{FP} + FN * \text{Cost}_{FN}}{\# \text{ of Test Samples}}$$

Figure 3.12: Equation used to calculate the cost of every model classification.

1700  
1701  
1702  
1703 i. Experiment 4 – RSR with 8160 Synthetic Samples

1704  
1705 The VGG16 model's test accuracy improved during RSR  
1706 75:25, 58:42, and 37:63 tests compared to 100:0. The  
1707 maximum test accuracy was 84.3% (+4.8%) compared to  
1708 the 100:0 test accuracy of 79.6%. Test Loss also decreased  
1709 for every RSR test compared to 100:0. The minimum test  
1710 loss was measured at RSR 75:25 at 0.987 (-0.100). The  
1711 training accuracy decreases as the RSR lowers towards  
1712 0:100, trending more in line with test accuracy. This is  
1713 similar performance to Experiments 1 & 2.

1714 Based on the training-test gaps of the loss and AUC  
1715 metrics, the VGG16 model is most likely still overfit, even  
1716 if introducing synthetic images reduces the gap slightly.  
1717 However, the VGG16 test AUC metric decreased as the  
1718 amount of synthetic data in the training set increased.  
1719

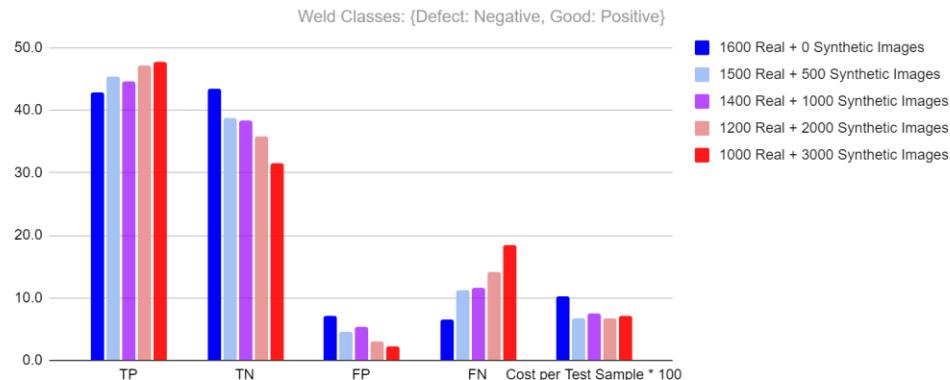


1725  
1726  
1727  
1728  
1729  
1730  
1731  
1732  
1733  
1734  
1735 Figure 3.13: Absolute training plots for VGG16 (left) Accuracy and (right) Loss. Shows max accuracy, max AUC, and min Loss when a training set is  
1736 75% real images and 25% synthetic. Accuracy and AUC decline for RSRs lower than 58:42, but Loss stays relatively unchanged when at least some  
1737 synthetic images are used in the training set.

1738  
1739 VGG16 Loss with varying concentrations of synthetic training data  
1740 Welding - 8160 Available Synthetic Images  
1741

Training Data ratio of 'Real:Synthetic'	% of Real Training Images	Train_Acc	Test_Acc	Loss
100:0	100%	~84.3%	~84.3%	~0.987
75:25	~75%	~83.5%	~83.5%	~0.987
58:42	~58%	~83.0%	~83.0%	~0.987
37:63	~37%	~82.5%	~82.5%	~0.987
25:75	~25%	~81.5%	~81.5%	~0.987

1742 VGG16 - Classifier Performance with Varying RSR



1744  
1745  
1746  
1747  
1748  
1749  
1750 Figure 3.14: Bar chart showing how VGG16 models trained on different sets of training data performed on confusion matrix focused metrics. {Defect  
1751 weld: Positive, God weld: Negative}. As the number of synthetic samples increased, the number of TN and FP classifications decreased. The number of  
1752 FNs increased. As such the Cost per Test Sample decreased as the number since a FP has a much higher cost associated with it.

MobileNet3's performance has much more worrying trends. The test accuracy metric shows similar trends to the performance in Experiments 1 & 2 with the Cat\_Dog synthetic dataset. Test accuracy does steadily decrease with more synthetic data, but the test loss appears to fluctuate. One of the more problematic trends is the MobileNetV3 test AUC metric hovering at approximately 50%. This signifies that the model has not decoupled what distinguishes a True Positive from a False Positive. These results and the bar chart further below send mixed signals as to how the classifier is learning from the imbalance training set and performing on the balanced test set.

One interesting observation to note is that the test accuracy for both VGG16 and MobileNetV3 are only a few percentage points different until RSR 50:50. This is less of a reflection on the quality of performance between the two models, but of how an imbalance dataset can skew a metric like accuracy and why other metrics are important to analyze. Figure 3.15 below shows that on average the MobileNetV3 model classified more and more samples as Positive (aka *Good*) as more synthetic data is introduced. This could be for many reasons; possibly that the MobileNetV3 model is more sensitive to the increased weight that error from *Good* samples created.

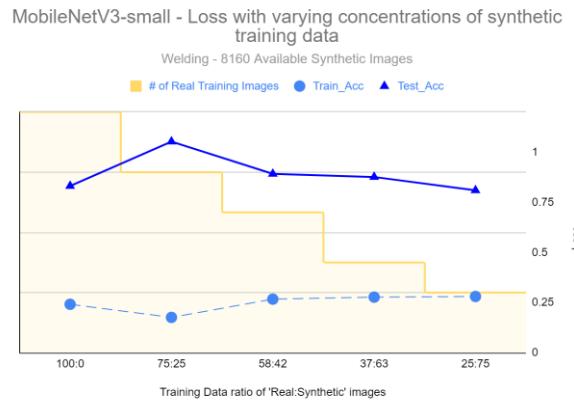
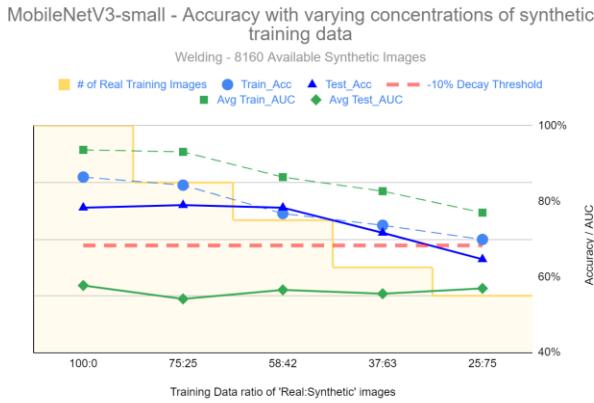


Figure 3.14: Absolute training plots for MobileNetV3-small (left) Accuracy and (right) Loss.

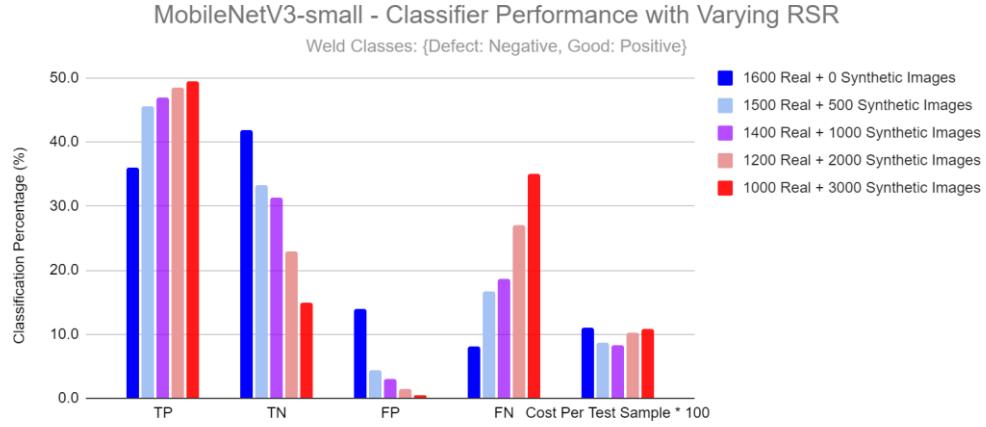


Figure 3.15: Bar chart showing how MobileNetV3 models trained on different sets of training data performed on confusion matrix focused metrics. {Defect weld: Positive, Good weld: Negative}. Highlights how more synthetic samples in the dataset skewed the MobileNetV3 models to classify most samples as Positives (aka. Good welds)

## ii. Experiment 5 - Rebalancing

The VGG16 model's results show that in terms of Cost per Test Sample, the oversampling method showed the greatest decline in cost (-18%) followed by synthetic padding (-12%). This is in line with Weiss' research) since the dataset used is below 10,000 samples. Compared to an imbalanced dataset, all methods decreased the Cost per Test Sample and increased test accuracy, but decreased test AUC and increased test Loss. A mix of desired and undesirable reactions. The synthetic padding results for loss are surprising. Dr Man's survey of the literature suggests that using synthetic images tends to reduce loss, not increase it [Man07].

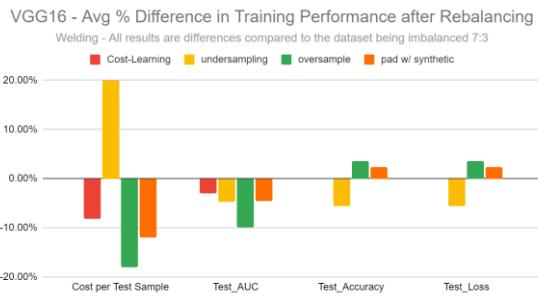


Figure 3.16: Bar chart showing the difference in performance metric between A) a VGG16 model that was trained on a sampled version of the original 2000 image real weld images, no synthetic (1200 Defect & 600 Good images are sampled), vs B) a VGG16 model that was trained a balanced dataset by one of the four methods shown.

In terms of Confusion Matrix classifications, all methods yielded a slight decrease in the cost per test sample and a mix of other reactions if each type is broken down individually. Overall, the oversample method showed a significant improvement of classifying defective welds (labeling TN instead of FN). Synthetic padding showed slight improvements over the imbalanced dataset, but not enough to draw any relevant conclusions. In addition, Cost-Learning closely replicated the results of the original dataset. More VGG16 results should be gathered for this experiment.

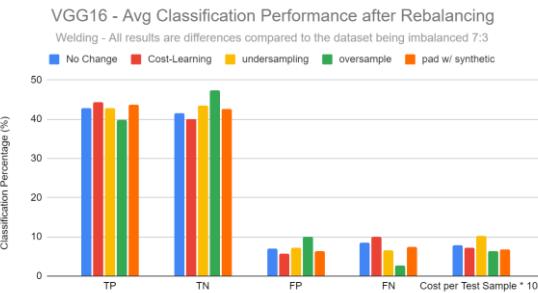


Figure 3.17: Confusion Matrix classifications for each rebalancing method. Only undersampling increased Cost per Test Sample.

For the MobileNetV3 model, the Cost-Learning method had the largest cost decrease, but the smallest change in loss. undersampling had the second largest decrease in

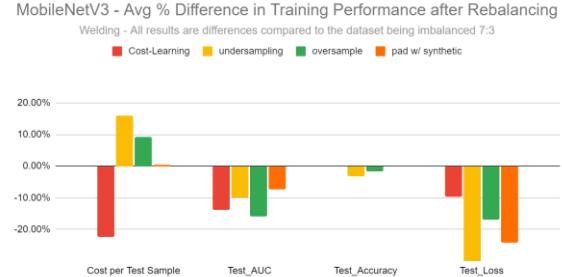


Figure 3.18: Bar chart showing the difference in performance metric between A) a MobileNetV3-small model that was trained on a sampled version of the original 2000 image real weld images and no synthetic (where 1200 Defect & 600 Good images were sampled), vs B) a MobileNetV3-small model that was trained a balanced dataset by one of the four methods shown.

The MobileNetV3 Confusion Matrix classification results show that Cost-Learning had smallest Cost per Test Sample value, followed by synthetic padding and oversampling. This can be misleading since Cost-Learning is the only method with imbalanced Cost classifications ( $\{0, 0, 0.667, 0.334\}$ ). Even though the method performed nearly identically to the unchanged imbalanced dataset. For MobileNetV3, these results do not provide a clear 'best rebalancing' method compared to using the dataset as is without class weights during training.



Figure 3.19: Confusion Matrix classification for each rebalancing method. Only oversampling decreased Cost per Test Sample.

Overall, it is not clear for the welding use case if any method provides a clear benefit for rebalancing the dataset. For VGG16, oversampling or synthetic padding may yield slightly better FP performance. For both models, undersampling negatively impacts performance. For this use case, removing samples for the sake of balance is not a good idea.

2000	2050
2001	2051
2002	2052
2003	2053
2004	2054
2005	2055
2006	2056
2007	2057
2008	2058
2009	2059
2010	2060
2011	2061
2012	2062
2013	2063
2014	2064
2015	2065
2016	i. Notes on MobileNetV3-small Training
2017	2066
2018	2067
2019	Unlike VGG16, the MobileNetV3-small model struggled
2020	2068
2021	to provide a clear connection of its own between synthetic
2022	2069
2023	images and model performance. MobileNetV3 reacted in
2024	2070
2025	similar ways to VGG16 like narrowing the training-test
2026	2071
2027	gap until after RSR 50:50, but the model was significantly
2028	2072
2029	more sensitive to the synthetic image datasets. It is
2030	2073
2031	possible that there were internal issues to how
2032	2074
2033	MobileNetV3 was being trained. During the training
2034	2075
2035	process, the best epoch of training was on average the first
2036	2076
2037	epoch to be judged for early stopping. In addition, while
2038	2077
2039	trying on the Cat_Dog synthetic dataset for RSR 25:75
2040	2078
2041	and lower, the test loss of the model could not converge.
2042	The validation loss diverged after 2-3 epochs like what
2043	2079
2044	happened in Gastelum's results, but with a training set
2045	2080
2046	with up to 2000 real images [Gastelum20]. However, in
2047	2081
2048	Experiment 4, VGG16's validation accuracy was slower to
2049	2082
	respond to every epoch's error gradient. As having more
	2083
	synthetic data created two separate feature space minimas
	2084
	(see Figure A.VII.3 in the appendix).
	2085
	2086
	For various accuracy and loss model training graphs
	2087
	showing converge and divergence, please see A.VIII.
	2088
	<i>Model Training Charts</i> section in the appendix For More
	2089
	Information.
	2090
	2091
	B. Synthetic Availability's Effect on Performance
	2092
	Based on Experiment 3's results and these two binary use
	2093
	cases (Cat_Dog & Weld Defect), there is no statistically
	2094
	relevant change in model performance if there is more
	2095
	synthetic data to be sampled. These results could be
	2096
	different if there was more varied synthetic data (such as
	2097
	being collected in full UE environments) or in different
	2098
	use cases. It is possible that having more synthetic samples
	2099
	can allow for more variation in the training set, but the
	way that the training set is sampled would also have to be
	more advanced. Furthermore, as the number of generated
	samples increases, each sample needs to be increasing

variation amongst the images, not repeating the feature space. The variations collected from the python state-machine could not be varied enough to expand the feature space. One interesting theory of expanding this need for variation to an infinitely large synthetic dataset is about becoming a simulation. If you require a sufficiently large synthetic dataset, conceptually the model might become the simulation that generated the images by being able to account for every little variation of ‘the real world’.

Please consider the following recommendations to improve collected synthetic data. First, utilize randomness in every facet of ‘setting up’ a screenshot (like the Welding collection), instead of collecting images on a set track with hardcoded values (like the Cat\_Dog collection). Second, in addition to sampling the synthetic collections proportionally (like what was done in this paper), a developer should proportionally sample the collection’s meta-data (model, lighting, etc). Proportionally sampling based on the collection’s meta-data can help prevent the oversampling of one 3D model over another to maintain image variation. Lastly, try incorporating more expansive UE environments during image collection. Only the UE test environment is used in this experiment, there are fully fleshed out landscapes, cities, and environments ready to insert the desired 3D models into.

### C. Efficacy of Rebalancing datasets with Synthetic data

Yes. Synthetic data can be used to rebalance an imbalance dataset if the synthetic data has enough variation. Synthetic Padding can be used to balance the minority class so more can be reallocated to the validation and test sets, if developers confirm the performance of a model does not change with synthetic padding. In terms of performance, Experiment 5 details how synthetic padding does not provide a clear advantage over other common sampling methods. It is possible that by mixing Cost-learning (class weighting) and synthetic padding that there would be a greater benefit. More research on more use cases will need to be done to make a clear determination.

Despite how rebalancing though synthetic padding performed, this is a reminder that undersampling without enough remaining data can harm a CV model more than it can benefit.

### D. Quantifying the Domain Gap between Real and Synthetic images

One way to quantify the domain gap between real and synthetic images is to use the model’s Validation Accuracy and Loss training graphs. As designed in Experiment 4, Figure 3.7 provides a way to visually represent the differences between a training set with and without synthetic data. The key is that the number of real images in the dataset needs to stay the same. With the number of real images changing from test to test (like in Experiments 1, 2, and 3), it is unclear if the change in performance is due to the added synthetic data or the loss of the real data.

### E. Commentary on created Image quality

The 3D models used in the Cat\_Dog and Welding use cases were fine representations of real life. They shared the shapes of those in real life and were quick to implement in the UE Editor. Dozens of models and weld defect variations can be created for little cost (\$12 for the base metal pack, all cats and dogs were free) and little startup time. After creating the python script, the Unreal Engine could output 80 image variations per minute. However, it is the author’s opinion that the generated images struggled from two critical flaws: resolution and feature variation. Resolution is the factor that determines how detailed the models or textures can be while zoomed in. Feature variation being how different the images are from one another.

The 3D Cat\_Dog model’s resolution is controlled by the 1) pixel resolution of their textures and 2) the polygon count that outlines the smoothness of their body. After reviewing the model’s descriptions, the authors report that most of the models are made from 10,000 Polygons. According to 3D Ace, 10,000 Polygons is considered a low-detailed character where the most complex environments can reach 60,000 polygons [3D Ace]. Therefore, the cat and dog models used in this report are of lower quality. In addition, the models lacked a ‘skeleton’ component, so it was impossible to adjust the limbs of said models to place them in different poses to increase image variation in the dataset.

The weld textures are constrained for a different type of resolution, pixel count. All of the welds provided in the Samur Art 4K Metal Panel Material Pack are described to be “4K”. It is true that the bitmap files that contain the groups of weld seams are 4096x4096 pixels, but the thickness of the weld seams in those 4096x4096 bitmaps are only ~35 pixels wide. Therefore, any created weld defects have less than 35 pixels to create dynamic ranges

2200		2250
2201		2251
2202		2252
2203	of color and light refraction. One alternative would be to diffuse the weld seam into a larger bitmap to get more pixels to work with. Regarding the defect design, the weld defect textures did provide significant feature differences since the size and shape of the defects are randomized, but the range of that randomization was limited. Also, more can be done to additional variation, such as using more than 6 noise texture maps to merge with the ‘base color’ textures.	2253
2204		2254
2205		2255
2206		2256
2207		2257
2208		2258
2209		2259
2210		2260
2211		2261
2212		2262
2213		2263
2214		2264
2215		2265
2216		2266
2217		2267
2218		2268
2219		2269
2220		2270
2221		2271
2222		2272
2223		2273
2224		2274
2225	One idea to improve image variation is to collect the images in more dense environments (as mentioned at the beginning of the Discussion). All of the created synthetic images for this research are taken in the UE project default test environment, a place without external shapes and objects. A blank canvas. There are maps pre-made in the UE marketplace such as a city block, a park, a spaceship, Italian riverways, etc. Having more complex objects in the background and foreground of the image could improve training results.	2275
2226		2276
2227		2277
2228		2278
2229		2279
2230		2280
2231	F. Practicality of generating Synthetic Images	2281
2232		2282
2233	UE created screenshots of cats, dogs and welds that are relevant representations of what can be done to gather a dataset quickly. With prior knowledge of a graphical engine and python, a user can generate a screenshot in a fraction of a second. They can perfectly control how the images are made, what environments to use, and what noise to add. If your application requires regular shapes and not a lot of detail (such as an airplane vs bicycle classifier like what the ECCV submission created) then generating your own synthetic data can save you time, money, and risk of errors should the data you do collect be incorrect, corrupted, or lost [ECCV18]. The ability to regenerate your entire dataset is a powerful tool to have in case the unexpected happens or if there needs to be a change in project scope. Alternatively, any use cases that requires complex shapes or finer details (like fur), generating your own data might not be practical unless there is a skilled graphic designer creating the models and is able to overestimate the required model variations to train a successful model.	2283
2234		2284
2235		2285
2236		2286
2237		2287
2238		2288
2239		2289
2240		2290
2241		2291
2242		2292
2243		2293
2244		2294
2245		2295
2246		2296
2247		2297
2248		2298
2249		2299

## 2300 5.V. Contributions

2301 This work contributes to research related to using  
2302 synthetic images in model training. First, having low  
2303 quality 3D models (<10,000 polygons) can replicate parts  
2304 of the feature space of complex shapes (animals) and  
2305 minor alterations (welding defects). It is possible to train  
2306 CV models on only low-quality 3D model images and  
2307 produce significant classifier accuracy (as shown in Figure  
2308 4.1 when training with 75% real samples). Second, CV  
2309 models with millions of parameters can train on small  
2310 amounts of real images and yield significant classifier  
2311 accuracy. This supports an alternative for needing  
2312 thousands of images to train an adequate classifier. Lastly,  
2313 if large amounts of real images are available, this research  
2314 also shows that low-quality 3D synthetic images can  
2315 reduce model loss and possibly increase model accuracy.  
2316

## 2317 6.VI. Future Research

2318 Future research can expand upon several areas. First, how  
2319 to reduce the domain gap between real and synthetic  
2320 images, such as the relationship between classifier  
2321 accuracy and factors such as model resolution and using  
2322 UE pre-built environments. Second, these tests can be  
2323 repeated to explore how synthetic images can introduce  
2324 object variations unknown to your model (reducing the  
2325 Curse of Dimensionality). Finally, work can be done to  
2326 improve the image pipeline software to create images with  
2327 bounding boxes that can train more advanced CV models.  
2328

2329 Specific to this research, further information can be  
2330 gathered by 1) repeating the RSR and Rebalancing  
2331 experiments for both use cases more than three times, 2)  
2332 repeating each experiment with a simplified CNN (made  
2333 from three decreasing kernel Convolution layers & two  
2334 Fully-Connected layers for instance) to better understand  
2335 the level that the published models are overfit, under  
2336 parameterized or the synthetic images poison the model,  
2337 3) repeat the experiments with less trainable parameters as  
2338 proposed by other literature (that suggest transfer learning  
2339 models to be ~10%-30% trainable) [Pan10], and 4) repeat  
2340 the Welding experiments with even more variation  
2341 provided by other weld textures in the 4K Metal Materials  
2342 pack or in UE environments other than the test  
2343 environment. Regardless of use case, experiments like  
2344 these should be repeated to better understand the current  
2345 domain gap between simulation and the real world.  
2346

## 2347 7.VII. Conclusion

2348 Variation. The key to a successful CV classifier is having  
2349 a dataset with a diverse feature space to learn from.  
2350 Through the documented RSR experiments, this research  
2351 can conclude that there are benefits to supplementing a  
2352 CV training set with synthetic Unreal Engine images. For  
2353 particular RSRs, replacing training set real images with  
2354 synthetic images can decrease your CV model's loss and  
2355 narrow the training-test accuracy gap. This also allows  
2356 developers to move valuable real samples from the  
2357 training set to the validation & test sets. However,  
2358 choosing a model architecture requires experimentation to  
2359 balance model size with the application's hardware  
2360 limitations. The 'server based' VGG16 model benefitted  
2361 from the injection of synthetic training images, while the  
2362 on-chip focused MobileNetV3-small model was  
2363 hampered. Investing in processes to produce synthetic data  
2364 can reduce the work, cost, and risk of collecting real  
2365 images manually.  
2366

## 2367 Source Code

2368 Please see the following Git repository for all source code  
2369 related to this report. The repository includes 1) the python  
2370 base class to implement a UE python callback, 2) proof of  
2371 concept python scripts to edit a UE environment for any  
2372 reader's reference, 3) python scripts to train the VGG16  
2373 and MobileNetV3-small model and 4) any other assisting  
2374 scripts.  
2375

2376 Link:  
2377

2378 [2379 https://github.com/smutnyjw/training\\_cv\\_models\\_on\\_ue\\_i](https://github.com/smutnyjw/training_cv_models_on_ue_i)  
2380 images\_jws  
2381

## APPENDIX

### A.I. Key Terms and Concepts

#### Paper Specific Terms

Real:Synthetic Ratio	(RSR) The ratio of real images to synthetic images used in a model's training set. This is heavily used to describe all results in this report.
Synthetic Utilization Rate	(SUR) The percentage of synthetic data included in a dataset versus what was available to be sampled. Used primarily in Experiment 3 of this report

#### Computer Vision Terms

Model Generalization	How well a Machine Learning model performs on samples it has not trained on. Characterized by the difference in a model's performance metrics during training and during testing.
Latent Variation	The amount of variance in features and/or the amount of variety in the images that make up a dataset. Examples include differences in camera perspective, backgrounds, foregrounds, color, shapes, positions, etc.
Transfer Learning	During model training, only apply loss gradients to specific layers of a model.

#### Graphic Design in Unreal Engine Terms

UE Editor	The central place in which a 3D simulated environment is viewed. This is where you can control the actors and their parameters. Python scripts can be executed from a UE Editor's command line if the Python API is enabled. As explained in the Generating Synthetic Images from the Unreal Engine 4.27, an executed python script through the python API will run on the same thread as the editor. Therefore, you must provide a frame buffer to your actions
Actor Mesh	The polygons that make up an actor's (object's) shape.
Material	The surface characteristics of an object (or mesh) in a 3D environment. A resulting output of Textures, Layering, UV Mapping, shading, and other techniques are what we see as the 'skin' of an object.
Map	A collection of pixels represented as a vector of one to four dimensions. [Red, Green, Blue, Alpha]
Texture map	A characteristic of an object that can be simulated over the course of all pixels in a map. Most common: <ul style="list-style-type: none"> <li>• Basecolor: The RGB color</li> <li>• Metallic: How similar a surface is to metal.</li> <li>• Roughness: Amount of light reflection on a surface</li> <li>• Ambient Occlusion: Definition of how light should reflect.</li> <li>• Normal: Definition of when light should reflect. Simulate height.</li> </ul>
Texture Parameter	Common ways to describe a particular texture, such as its metallic, roughness, etc qualities. This affects the shininess and light reflections of the object. Normally represented as a float between 0 and 1.
Texture mask	A way to selectively mix a texture with another in a material
Rendering	The act of converting a 3D simulation into a 2D image of a particular resolution.
UV Mapping	Flattening a 3D model into a 2D image. 'U' and 'V' reference the resulting 2D image's axis, not the x,y,z axis of the 3D model.
SunSky Lighting	A light ray generator to act like the sun
Individual Light	A light source of light rays that acts like a light bulb or spotlight
Camera Location	The geographic coordinates of a CineCameraActor
Camera Tracking	Boolean setting to force a camera to always follow an actor's center point

2500		2550
2501		2551
2502		2552
2503	Welding 101	2553
2504		2554
2505	Terms	2555
2506	A Metal Weld	2556
2507	Using heat to break the molecular bonds of different metals and then reforming bonds to another material (filler or another piece of metal). Non-ionizing gas is used to create an environment for welding to occur and prevent weld contamination.	2557
2508	Weld Filler	2558
2509	Another alloy is added between two metal plates that you wish to weld. The filler bonds to the plates when melted and results in the weld when cooled.	2559
2510		2560
2511	MIG Welding	2561
2512	Metal Inert Gas. Welding is where a filling alloy is extruded through a welding gun into a gas mixture to facilitate an arc. Heating the metal plates and filler material. Bonding them when cooled [Openworks].	2562
2513		2563
2514	TIG Welding	2564
2515	Tungsten Inert Gas. Welding where an electrode is at the end of a welding gun. Gas is emitted from the gun to facilitate an arc between the electrode and desired metal. Adding filler material via a separate filler rod is optional [Openworks].	2565
2516		2566
2517	<u>How weld defects form?</u>	2567
2518		2568
2519	The quality of the weld is determined by human technique (speed), heat of the metals, coverage of shielding gas, and metallic properties of the destination metal & filler. Defects in a welding material can occur when there is a mismanagement in any of these. The biggest concern with a weld is balancing a welder's output voltage (temperature) and feed rate (how much filler material to output) to ensure that the weld has sufficient penetration (the filler material bonds to the entire depth of the weld [Padayao]).	2569
2520		2570
2521		2571
2522		2572
2523		2573
2524		2574
2525	Shielding gas such as Argon or Argon Carbon Dioxide mix is emitted from the arcing tool that heats up the metal. This gas is meant to protect the molten pool of metal that eventually creates the weld when cooled. Defects can occur when particles or gasses mix get past the shielding gas and mix with the molten pool before it cools. Potentially causing splatter, gas bubbles to form and cause porosity or internal holes [Openworks].	2575
2526		2576
2527		2577
2528		2578
2529	Significant differences in metallic properties of the base and the filler material such as heat diffusion and thermal conductivity can make forming a good weld difficult. Thermal conductivity describes how quickly a metal heats up while heat diffusion describes how quickly heat spreads throughout a metal. If a base and filler materials have different properties it is possible that one can get too hot and melt through material while the other is still heating up. Or inversely, one is ready while the other is too cold for the weld to adhere too. The wrong choice of filler for your base material can cause holes in your material or incomplete welds [Padayao].	2579
2530		2580
2531		2581
2532		2582
2533		2583
2534		2584
2535		2585
2536	Finally, operator speed, material output rate, and chosen heat output control the molten pool's penetration into the weld and completeness of the weld as just discussed with metallic properties. Think of managing the molten pool like using a hot glue gun. Move too fast or output not enough material and your weld will not penetrate through the entire weld (resulting in a weak weld and potential cracking). Moving too slow and outputting too much material results in splatter, contamination, and excess [Padayao].	2586
2537		2587
2538		2588
2539		2589
2540		2590
2541		2591
2542	<u>Current Weld Defect detection practices</u>	2592
2543		2593
2544	According to professional welder Brandon Padayao, there are several checks to uncover the integrity of welds involving different parties. The most common defect detector is the welder. All professional welders are certified by the American Welder Society to do their job safely and effectively. Becoming certified signals the skills to visually identify at the point of creation if a weld is sufficient or not. The second check is other professional inspectors that perform various tests to ensure the weld meets specifications, such as being liquid tight. Tests involve flashlight inspection, applying dyes to find holes, applying water to find leaks, x-rays to check weld penetration, or even cutting the weld open to visually inspect the operator's performance (the welder then re-welds the cut. The last check comes from post-work visual inspections [Padayao]	2594
2545		2595
2546		2596
2547		2597
2548		2598
2549		2599

2600  
 2601  
 2602  
**2603 A.II. Installed Python Software Packages**  
 2604  
 2605  
 2606

2607	Package	Version	2608	Package	Version	2609	Package	Version
	<i>absl-py</i>	2.1.0		<i>jsonpointer</i>	2.1		<i>pyOpenSSL</i>	23.2.0
	<i>archspec</i>	0.2.1		<i>keras</i>	2.10.0		<i>PySocks</i>	1.7.1
	<i>astunparse</i>	1.6.3		<i>Keras-Preprocessing</i>	1.1.2		<i>requests</i>	2.31.0
	<i>boltons</i>	23.0.0		<i>libclang</i>	16.0.6		<i>requests-oauthlib</i>	1.3.1
	<i>Brotli</i>	1.0.9		<i>libmambapy</i>	1.5.6		<i>rich</i>	13.7.1
	<i>cachetools</i>	5.3.2		<i>Markdown</i>	3.5.2		<i>rsa</i>	4.9
	<i>certifi</i>	2024.2.2		<i>markdown-it-py</i>	3.0.0		<i>ruamel.yaml</i>	0.17.21
	<i>cffi</i>	1.16.0		<i>MarkupSafe</i>	2.1.5		<i>ruamel.yaml.clib</i>	0.2.6
	<i>charset-normalizer</i>	2.0.4		<i>mdurl</i>	0.1.2		<i>setuptools</i>	68.2.2
	<i>colorama</i>	0.4.6		<i>menuinst</i>	2.0.2		<i>six</i>	1.16.0
	<i>conda-package-handling</i>	2.2.0		<i>namex</i>	0.0.7		<i>tensorboard</i>	2.10.1
	<i>conda_package_streaming</i>	0.9.0		<i>numpy</i>	1.26.4		<i>tensorboard-data-server</i>	0.6.1
	<i>cryptography</i>	41.0.7		<i>oauthlib</i>	3.2.2		<i>tensorboard-plugin-wit</i>	1.8.1
	<i>distro</i>	1.8.0		<i>opt-einsum</i>	3.3.0		<i>tensorflow</i>	2.10.1
	<i>dm-tree</i>	0.1.8		<i>packaging</i>	23.1		<i>tensorflow-estimator</i>	2.10.0
	<i>flatbuffers</i>	23.5.26		<i>pip</i>	23.3.1		<i>tensorflow-io-gcs-filesystem</i>	0.31.0
	<i>gast</i>	0.4.0		<i>platformdirs</i>	3.10.0		<i>termcolor</i>	2.4.0
	<i>google-auth</i>	2.27.0		<i>pluggy</i>	1.0.0		<i>tqdm</i>	4.65.0
	<i>google-auth-oauthlib</i>	0.4.6		<i>protobuf</i>	3.19.6		<i>truststore</i>	0.8.0
	<i>google-pasta</i>	0.2.0		<i>pyasn1</i>	0.5.1		<i>typing_extensions</i>	4.9.0
	<i>grpcio</i>	1.60.1		<i>pyasn1-modules</i>	0.3.0		<i>urllib3</i>	1.26.18
	<i>h5py</i>	3.10.0		<i>pycosat</i>	0.6.6		<i>Werkzeug</i>	3.0.1
	<i>idna</i>	3.4		<i>pycparser</i>	2.21		<i>wheel</i>	0.41.2
	<i>jsonpatch</i>	1.32		<i>Pygments</i>	2.17.2		<i>win-inet-pton</i>	1.1.0
							<i>wrapt</i>	1.16.0
							<i>zstandard</i>	0.19.0

2631  
 2632  
 2633  
 2634  
 2635  
 2636  
 2637  
 2638  
 2639  
 2640  
 2641  
 2642  
 2643  
 2644  
 2645  
 2646  
 2647  
 2648  
 2649  
 2650  
 2651  
 2652  
 2653  
 2654  
 2655  
 2656  
 2657  
 2658  
 2659  
 2660  
 2661  
 2662  
 2663  
 2664  
 2665  
 2666  
 2667  
 2668  
 2669  
 2670  
 2671  
 2672  
 2673  
 2674  
 2675  
 2676  
 2677  
 2678  
 2679  
 2680  
 2681  
 2682  
 2683  
 2684  
 2685  
 2686  
 2687  
 2688  
 2689  
 2690  
 2691  
 2692  
 2693  
 2694  
 2695  
 2696  
 2697  
 2698  
 2699

2700  
 2701  
 2702  
 2703 A.III. Model Architecture changes from Published Models  
 2704  
 2705 Output Layers  
 2706  
 2707 Specifically for the model's output layers, minimal changes are made to the output layers of both VGG16 and  
 2708 MobileNetV3 models. This is done to maintain the architecture of the original models and the author's thought process  
 2709 behind their perspective research. However, changes are made to reduce the size and complexity of the models so that they  
 2710 are better aligned with the change from the 1000 class classification task to only a binary classification task. The reduced  
 2711 output layers result in an approximate 84.7% and 46.6% reduction in model size for the VGG16 & MobileNetV3 models  
 2712 respectively.  
 2713

	VGG16	MobileNetV3
Total Parameters in the published model	138,357,544	2,554,968
Total Parameters in the transfer learning model	21,203,521 (15.3%)	1,365,617 (53.4%)

2714  
 2715 *Figure A.III.1: Summary of model size reduction from the 1000 class case to binary classification [Simonyan14][Howard19].*  
 2716  
 2717  
 2718  
 2719

2720 In conjunction with reducing model size, only a fraction of the model's parameters are trainable (thus transfer learning). In  
 2721 this research only 41.7% and 39.4% of parameters are affected by learning gradients for layers in the output and non-output  
 2722 layers. The primary cause for this reduction is the reduction in output layer size for both models. The VGG16 output layers  
 2723 are a pair of Fully-Connected (FC) Dense layers which are reduced from 4096 nodes each to 256 nodes. The  
 2724 MobileNetV3's pair of Convolutional output layers are also reduced from 1024 filters each to 512 filters and 256 filters  
 2725 respectively. Both reductions are chosen to yield a total trainable parameter percentage of approximately 40% for each  
 2726 model. 40% is the target percentage of trainable parameters due to the disproportionate size of the model output layers even  
 2727 after the previously stated reductions above and a desire to retrain at least one convolutional layer in each respective model.  
 2728 More review should be done to determine what is the best approach to setting the number of trainable parameters.  
 2729  
 2730

	VGG16	MobileNetV3
Total Parameters in the transfer learning model	21,203,521	1,365,617
Total Trainable Parameters	8,848,641 (41.7%)	538,433 (39.4%)
<u>Further breakdown</u>		
Number of parameters in the non-output layers	14,714,688 (69.4%)	939,120 (68.8%)
Percentage of parameters in the non-output layers that are trainable	26.7%	20.8%
Number of parameters in the output layer	6,488,833 (30.6%)	426,497 (31.2%)
Percentage of Parameters in the output layers that are trainable	100%	100%
Total Conv Layers	13	28
Trainable Conv Layers	1 trainable conv layers (3 trainable layers total)	11 trainable conv layers (24 trainable layers total)

2731  
 2732 *Figure A.III.2: Summary of the parameters in the VGG16 and MobileNetV3 transfer learning models used in this report. This documents the location of  
 2733 parameters for the reduced models to binary classification.*  
 2734  
 2735  
 2736  
 2737  
 2738  
 2739  
 2740  
 2741  
 2742  
 2743  
 2744  
 2745  
 2746  
 2747  
 2748  
 2749

The tables below present more specific architectural changes between the published paper and the models used in the paper. Details include output layer size, amount of layer regularization, additional learning techniques, etc. Any row in **BOLD** represents a setting that is different from the published model's settings.

VGG16		
Settings	Original Paper	Transfer Learning Model
Optimizer	Mini-batch SGD	Batch SGD
Optimizer settings	Momentum=0.9	beta_1=0.9 (default) beta_2=0.999 (default) epsilon=1e-7 (default)
<b>Initial Learning Rate</b>	<b>0.01</b>	<b>0.0001</b>
<b>Batch Size</b>	<b>256</b>	<b>16</b>
<b>Learning Rate Decay</b>	<b>0.1 when Val_accuracy stopped improving</b>	<b>0.95 every epoch</b>
<b>Dataset Used</b>	<b>ILSVRC 2012-2014 (1.45 million images)</b>	Various (Cats_Dogs & Welding)
<b>Dataset Split</b>	<b>1.45M Total (1.3M/50k/100k)</b>	Various (up to 10k total) (80%/10%/10%)
Input Image Size	(224, 224, 3)	(224,224,3)
<u>Output Layer Information</u>		
<b>Number of Nodes per FC Dense Layer</b>	<b>4096</b>	<b>256</b>
Dropout Rate in Output Layers	None	None
<b>Last Layer Activation Function</b>	<b>Softmax, 1000</b>	<b>Sigmoid, 1</b>
Kernel Regularization	L2=5e-4	L2=5e-4
Kernel Weight Initialization	Glorot & Bengio (2010) o Normal Distribution	GlorotNormal (see documentation)
<u>Early Stopping</u>		
<b>Metric</b>	<b>Val_Accuracy</b>	<b>Val_Loss (Cat_Dog) Val_AUC (Welding)</b>
<b>Minimum Delta</b>	<b>0.0</b>	<b>0.0005</b>
<b>Patience (# epochs)</b>	<b>unknown</b>	<b>7</b>
<u>Initial Training</u>		
<b>Number of Epochs before evaluation</b>	<b>Unknown</b>	<b>5</b>

Figure A.III.3: Summary of the training differences between the originally published VGG16 model and what is used in this report [Simonyan14].

2900  
2901  
2902  
2903 Deployed VGG16 Model Architecture  
2904  
2905 The following output layer information is taken from the Keras API's implementation of the VGG16 model. Only the  
2906 number of parameters in each Fully-Connected layer is different between the implementation below and the published  
2907 VGG16 model. The published model uses 4096 nodes in the *fc1* and *fc2* Dense layers.  
2908  
2909  
2910  
2911  
2912  
2913  
2914  
2915  
2916  
2917  
2918  
2919  
2920  
2921  
2922  
2923  
2924  
2925  
2926  
2927  
2928  
2929  
2930  
2931  
2932  
2933  
2934  
2935  
2936  
2937  
2938  
2939  
2940  
2941  
2942  
2943  
2944  
2945  
2946  
2947  
2948  
2949

Model: "sequential"				
Layer (type)	Output Shape	Param #	Trainable	
vgg16 (Functional)	(None, 7, 7, 512)	14714688	Y	
flatten (Flatten)	(None, 25088)	0	Y	
fc1 (Dense)	(None, 256)	6422784	Y	
fc2 (Dense)	(None, 256)	65792	Y	
Prediction (Dense)	(None, 1)	257	Y	
=====				
Total params: 21,203,521				
Trainable params: 8,848,641				
Non-trainable params: 12,354,880				
=====				

Figure A.III.4: Model architecture of the VGG16 transfer learning model. 'vgg16' is the keras pre-built VGG16 model [Simonyan14].

The tables below present more specific architectural changes between the published paper and the models used in the paper. Details include output layer size, amount of layer regularization, additional learning techniques, etc. Any row in **BOLD** represents a setting that is different from the published model's settings.

MobileNetV3-Small		
Settings	Original Paper	Transfer Learning Model
Optimizer	Tensorflow RMSProp	Tensorflow RMSProp
Optimizer settings	Momentum=0.9	Momentum=0.9 (default)
<b>Initial Learning Rate</b>	<b>0.1</b>	<b>0.0001</b>
<b>Learning Rate Decay</b>	<b>0.99 every 3 epochs</b>	<b>0.95 every epoch</b>
<b>Batch Size</b>	<b>4096</b>	<b>100</b>
Input Image Size	(224, 224, 3)	(224,224,3)
<b>Dataset Used</b>	<b>ImageNet</b>	<b>Various (Cats_Dogs &amp; Welding)</b>
<b>Dataset Split</b>	<b>Unknown</b>	<b>Various (up to 10k total) (80%/10%/10%)</b>
<u>Output Layer Information</u>		
Dropout rate in Output Layers	0.8	0.8
<b>Last Layer Activation Function</b>	<b>Softmax, 1000</b>	<b>Sigmoid, 1</b>
<b>Number of Filters per Convolutional Layer in the Output Layers</b>	<b>{1024, 1024}</b>	<b>{512, 256}</b>
<b>Kernel Regularization</b>	<b>L2=1e-5</b>	<b>L2=1e-4</b>
<b>Kernel Weight Initialization</b>	<b>Unknown</b>	<b>GlorotNormal (see documentation)</b>
<u>Early Stopping</u>		
<b>Metric</b>	<b>Unknown</b>	<b>Val_Accuracy (Cat_Dog) Val_AUC (Welding)</b>
<b>Minimum Delta</b>	<b>Unknown</b>	<b>0.0005</b>
<b>Patience (# epochs)</b>	<b>Unknown</b>	<b>7</b>
<u>Initial Training</u>		
<b>Number of Epochs before evaluation</b>	<b>Unknown</b>	<b>5</b>

Figure A.III.5: Summary of the training differences between the originally published MobileNetV3-small model and what is used in this report [Howard19]

3100  
 3101  
 3102  
 3103 Deployed MobileNetV3-small Architecture  
 3104  
 3105 The output layers for the MobileNetV3-small architecture were recreated based on the published paper's description. The  
 3106 main architecture changes are 1) the use of AvgPooling instead of GlobalPooling, 2) the number of filters for convolutional  
 3107 layers *Conv\_2* and *logits*. Instead of 1024 filters for each Conv2D layer, 512 and 256 filters are used. Howard, et al. uses a  
 3108 new activation layer called hard-swish in order to complete its set of convolutional layers. Hard-Swish uses a residual  
 3109 network split, starting from the *Conv\_2* layer and ending at the *multiply\_18* layer. One of the splits is not processed further,  
 3110 but the other goes through several steps summarized by the equation below [Howard19].  
 3111  
 3112

x	The output of the pool_7x7 layer. The split starting the residual network
y	The merging of the residual network. Input into the multiply_18 layer

$$y = x * \frac{\text{Relu}(x+3)}{6}$$

Layer (type)	Output Shape	Param #	Connected to	Trainable
input_2 (InputLayer)	[None, 224, 224, 3]	0	[]	Y
MobilenetV3small (Functional)	(None, 7, 7, 576)	939120	['input_2[0][0]']	Y
pool_7x7 (AveragePooling2D)	(None, 1, 1, 576)	0	['MobilenetV3small[0][0]']	Y
Conv_2 (Conv2D)	(None, 1, 1, 512)	294912	['pool_7x7[0][0]']	Y
h_swish-tf_TFOpLambda_add_3 (Lambd	(None, 1, 1, 512)	0	['Conv_2[0][0]']	Y
h_swish-re_lu (ReLU)	(None, 1, 1, 512)	0	['h_swish-tf_TFOpLambda_add_3[0][0]']	Y
h_swish-tf_TFOpLambda_multiply_6th (Lambda)	(None, 1, 1, 512)	0	['h_swish-re_lu[0][0]']	Y
multiply_18 (Multiply)	(None, 1, 1, 512)	0	['Conv_2[0][0]', 'h_swish-tf_TFOpLambda_multiply_6th[0][0]']	Y
dropout (Dropout)	(None, 1, 1, 512)	0	['multiply_18[0][0]']	Y
logits (Conv2D)	(None, 1, 1, 256)	131328	['dropout[0][0]']	Y
flatten (Flatten)	(None, 256)	0	['logits[0][0]']	Y
dense (Dense)	(None, 1)	257	['flatten[0][0]']	Y
Total params:	1,365,617			

Figure A.III.4: Model architecture of the VGG16 transfer learning model. 'vgg16' is the keras pre-built VGG16 model.

3200  
 3201  
 3202  
 3203 A.IV. 2D Image Augmentations  
 3204  
 3205

3206 The following list of augmentations are available to be applied to every image used in a training epoch. All augmentations  
 3207 have a uniform random change of being applied and a uniform random sampling of the value range. These are all  
 3208 implemented through the Tensorflow ImageDataGenerator class.  
 3209

Augmentation	Value	Reason
Rotation range	+/-15 deg	Desire to give enough variation especially when combined with horizontal and vertical flipping. 15 seemed about right. No elaborate testing
width_shift_range	+/-10%	Since the camera is always centering the object, I wanted at least a +/- 10% shift from a (224, 224, 3) image. See the <i>Fill_mode</i> row for information on the resulting fill algorithm.
height_shift_range	+/-10%	Since the camera is always centering the object, I wanted at least a +/- 10% shift from a (224, 224, 3) image. See the <i>Fill_mode</i> row for information on the resulting fill algorithm.
Fill_mode	nearest	Between nearest and constant. Chosen “nearest” so there is continuity in the image. Other options result in hard dividing lines if combined with augmentations such as rotate or shear.
brightness_range	[0.1, 1.0]	The pictures are already bright, so more brightness was not needed. However, the range was put to near 0 to prevent any BLACK images but still try to lower the brightness. Upon inspection, it doesn't look like this setting did much.
shear_range	+/-10 deg	Desire to give enough variation. 10 seemed about right. No elaborate testing
zoom_range	0.15	Through trial and error, 15% zoom was enough to give valid variation. Anymore and there were too many close ups, and the models were too small when zoomed out.
channel_shift	30	Gave enough rgb alterations ('halo' effect) to believe that 30 is enough. Too much channel_shift led to wild color contrast that took away from the image
horizontal_flip	True	General augmentation.
vertical_flip	True	General augmentation

3236 *Figure A.IV.1: List of 2D Augmentations and values that could be applied to an image in the training set every epoch. Augmentations stack so none, all,  
 3237 or some may be applied.*

3238  
 3239  
 3240  
 3241  
 3242  
 3243  
 3244  
 3245  
 3246  
 3247  
 3248  
 3249 3250  
 3251  
 3252  
 3253  
 3254  
 3255  
 3256  
 3257  
 3258  
 3259  
 3260  
 3261  
 3262  
 3263  
 3264  
 3265  
 3266  
 3267  
 3268  
 3269  
 3270  
 3271  
 3272  
 3273  
 3274  
 3275  
 3276  
 3277  
 3278  
 3279  
 3280  
 3281  
 3282  
 3283  
 3284  
 3285  
 3286  
 3287  
 3288  
 3289  
 3290  
 3291  
 3292  
 3293  
 3294  
 3295  
 3296  
 3297  
 3298  
 3299

3300	3350
3301	3351
3302	3352
3303	A.V. Collecting Synthetic Images via the UE Python API 3353
3304	3354
3305	Synthetic images are collected by taking screenshots of the current UE Editor being viewed from a CineCameraActor 3355
3306	and in the Editor's Cinematic mode. It is possible to write a python script that imports the unreal python package and 3356
3307	then the .py file be executed in a currently active UE Environment through its terminal. The python script sends commands 3357
3308	to the UE Editor and can manipulate numerous parts of the Editor such as an object's geospatial position, lighting intensity, 3358
3309	displayed materials on a mesh, and texture parameters on an object's material instance. For this research, all Editor 3359
3310	variations are done through a single python script instead of experimenting with UE Blueprint functions or other associated 3360
3311	timetables due to the researcher's introductory knowledge of the UE. 3361
3312	3362
3313	The developed python script requires that the UE Editor make precise changes in a structured sequence (like a state- 3363
3314	machine). The script takes a defined sequence and iterates through all model/background/lighting variations required to set 3364
3315	the Editor in unique ways for a unique screenshot. This is possible by creating a python class to facilitate a user created 3365
3316	software callback that is tied to the UE's internal frame counter. This means that every 'tick' of the UE software process, 3366
3317	the python callback is executed, and UE commands are processed before the next process 'tick'. The python scripts created 3367
3318	for this research is based on the class structure outlined by the Technical Art Engineer, Ryan DowlingSoka, to control the 3368
3319	UE through asynchronous loops to accomplish this [DowlingSoka]. 3369
3320	3370
3321	3371
3322	class MyClass(object): 3372
3323	def __init__(self) -> None: 3373
3324	self.frame_count = 0 3374
3325	self.max_count = 1000 3375
3326	def start(self) -> None: 3376
3327	self.slate_post_tick_handle = unreal.register_slate_post_tick_callback(self.tick) 3377
3328	self.frame_count = 0 3378
3329	def tick(self, delta_time: float) -> None: 3379
3330	print(self.frame_count) 3380
3331	self.frame_count += 1 3381
3332	if self.frame_count >= self.max_count: 3382
3333	unreal.unregister_slate_post_tick_callback(self.slate_post_tick_handle) 3383
3334	test = MyClass() 3384
3335	test.start() 3385
3336	3386
3337	3387
3338	3388
3339	3389
3340	3390
3341	3391
3342	3392
3343	Figure A.V.1: Base python class DowlingSoka used from UE documentation to create a 'tick' callback. 3393
3344	3394

Operationally, the developed python scripts that gather all the synthetic Cat\_Dog and Weld screenshots operate like a state-machine iterated via a counter. Every number of UE Editor frame 'ticks' increments the python scripts sequence counter. Based on this counter, the material parameters (like what background is showing or what defect is present), lighting position, camera position, and focused model changes. The script ends when the sequence is completed, the frame count exceeds the user defined limit, or if the script takes too many actions. These error checks are put in place to prevent the python executable from running an infinite loop. The author has not identified how to stop a python script executed from the UE Editor command line once it starts other than terminating the entire UE program.

3400  
 3401  
 3402  
 3403 Based on the author's research on online forums, Unreal Engine help threads, and Ryan DowlingSoka's documentation, it  
 3404 is the author's belief that there is one large limitation of controlling the UE Editor via tick callback. This limitation is that  
 3405 the python script is executed on the same thread as the actual UE Editor. The UE Editor does not appear to create a  
 3406 separate thread to execute the python script. Therefore, it is the developer's responsibility to not execute too many  
 3407 commands in a single UE frame tick. If a script callback has too many actions, then the callback won't be completed by the  
 3408 time the next tick executes the callback. Thus, the script drops the remaining commands. This can be mitigated by  
 3409 changing when the code in the `::tick()` class method is executed via the modulus operator (%). Creating a  
 3410 `frame_buffer` will slow down the rate at which callbacks can be executed but it guarantees that a single callback does  
 3411 all the actions that the user wants it to.  
 3412  
 3413

```

3414     class MyClass(object):
3415         def __init__(self, frames_buffer: int) -> None:
3416             self.frame_count = 0
3417             self.max_count = 1000
3418             self.frame_buffer = frames_buffer
3419
3420         def start(self) -> None:
3421             self.slate_post_tick_handle = unreal.register_slate_post_tick_callback(self.tick)
3422             self.frame_count = 0
3423
3424         def tick(self, delta_time: float) -> None:
3425             print(self.frame_count)
3426             self.frame_count += 1
3427             if self.frame_count >= self.max_count:
3428                 unreal.unregister_slate_post_tick_callback(
3429                     self.slate_post_tick_handle)
3430             elif self.frame_count % self.frame_buffer == 0:
3431                 print("Do things for one callback")
3432
3433
3434 frames_bw_ticks = 10
3435 test = MyClass(frames_bw_ticks)
3436 test.start()
3437
3438
3439
3440
3441
3442
3443
3444
3445
3446
3447
3448
3449
  
```

*Figure A.V.1: Appended python class from DowlingSoka to use a class parameterer 'frames\_bw\_ticks' to specify how many ticks are needed to fully execute each callbacak. The user should customize this value, so it is 1) large enough to fully complete the expected actions and 2) be small enough to generate images as fast as possible.*

3500  
3501  
3502  
3503 A.VI. Cat\_Dog - Collecting Synthetic Images  
3504  
3505  
3506  
3507  
3508  
3509

The generated synthetic Cat\_Dog dataset contains 46784 images of 3D models & accompanying textures of cats and dogs with a variety of camera angles, backgrounds, and lighting. A custom python script connects to the active UE project via the UE Python API and sends commands to the UE Editor to change the Editor's camera and environment to collect a unique screenshot. The following attributes are changed throughout the collection: camera angle (spherical position centered around a model), the 3D model in focus, the sky lighting, the background in the screenshot.



Figure A.VI.1: Examples of the UE test environment backdrop during Cat\_Dog synthetic image collection.

3510  
3511  
3512  
3513  
3514  
3515  
3516  
3517  
3518  
3519  
3520  
3521  
3522  
3523  
3524  
3525  
3526  
3527  
3528  
3529  
3530  
3531  
3532  
3533  
3534  
3535  
3536  
3537  
3538  
3539  
3540  
3541  
3542  
3543  
3544  
3545  
3546  
3547  
3548  
3549  
3550  
3551  
3552  
3553  
3554  
3555  
3556  
3557  
3558  
3559  
3560  
3561  
3562  
3563  
3564  
3565  
3566  
3567  
3568  
3569  
3570  
3571  
3572  
3573  
3574  
3575  
3576  
3577  
3578  
3579  
3580  
3581  
3582  
3583  
3584  
3585  
3586  
3587  
3588  
3589  
3590  
3591  
3592  
3593  
3594  
3595  
3596  
3597  
3598  
3599

The 3D models are free to download from websites such as Turbosquid and Free3D. Each model has different levels of quality, polygon count, texture variation, and anatomic position. Purchased models can have animations, skeletal positions, and more texture color palettes that could increase variation. In total 13 unique cat models & variations and 10 unique dog models, are set in front of 10 separate backgrounds to generate all synthetic data. Please see the References section for all 3D models used for authors and credits.

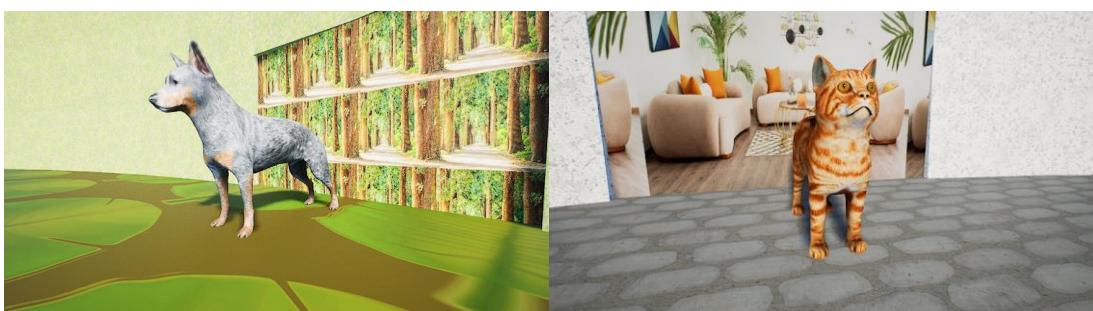


Figure A.VI.2: Examples of a cat and dog screenshot during the Cat\_Dog synthetic image collection.

3600  
3601  
3602  
3603 The used backgrounds consist of three of three main components, each with their own texture; Center Scene (directly  
3604 behind the model), Setting (additional surrounding background affiliated with the Center Scene), and floor (underneath the  
3605 model). During each collection, the python script flips switches in the loaded UE materials for all three of these mesh  
3606 objects to change which texture is shown. Across the two collections, ten unique combinations are used. Please see the  
3607 References section for image credit for all images used.  
3608  
3609  
3610  
3611  
3612  
3613  
3614  
3615  
3616  
3617  
3618  
3619



3620 Figure A.VI.3: Description of the backdrop used during Cat\_Dog synthetic image collection. A unique Center\_scene, setting, and floor are used for each  
3621 variation during image collection.

#### 3624 Collection Method

3625 46784 images were collected over the course of two collections. The first set of images (9424) contained 4 different scenes  
3626 while the second set of images (37360). The table on the following page outlines the iterated through settings to yield the  
3627 total amount of images. Due to how the python script sends commands over the UE Python API and how the UE operates  
3628 on the computer's threads, it is possible to lose images during collection. It is the researcher's belief that this is due to the  
3629 UE operating on a set number of computer threads but not being a multithreaded software process. Thus, if the user clicks  
3630 from the UE during collection, commands sent by the python script can be missed or skipped, losing that screenshot.  
3631  
3632

3633 Screenshots are collected by moving a cat/dog model to a center point in front of the set background, in set lighting, at  
3634 hardcoded camera angles. Once all camera angles have been captured, the cat/dog model is moved out of the center point  
3635 out of the frame and the next cat/dog model is moved to the center. The process repeats for all cat and dog models, then the  
3636 background preset rotates to the next. Once all background presets are completed then the lighting in the test environment  
3637 changes. The lighting element in these data collections is the 'sun'. The lighting applied to the model and environment is  
3638 simulated off earth's sun. The process of iterating through cat/dog models, background presets, and sun time variations  
3639 continues until a screenshot is taken of every variation.  
3640

3641 Please see the tables below for what specific camera angles, backdrops and light intensities were used to collect the  
3642 synthetic data.  
3643  
3644  
3645  
3646  
3647  
3648  
3649

		Collection 1	
		Run 1	Run 2
3707	Total Images	6672	2753
3709	Total Variations	7728	2760
3711	Image Loss (no screenshot taken)	13.7%* (6672 from expected 7728)	0.03% (2753 from expected 2760)
3713	<u>Variations</u>		
3714	Models	23	23
3716	Angles	Radi = [175, 300] Alpha = [-180, -135, -105, -90, -75, -45, 0] Phi = [55, 75]	Radi = [220] Alpha = [-120, -110, -92, -70, -60] Phi = [60, 80]
3718	(sets of angles that are unique to each run. Not shared)		
3719	Total Lighting settings	150	[90, 140, 190]
3720	Total available Presents	4	4

Figure A.VI.4: Description of the first Cat\_Dog synthetic image collection. Note that in Run 1, 13.7% of samples were lost. This was due to the user clicking off of the UE window during python callback execution. Moving off of the UE window caused the process to slow down enough that the python script needed more frame\_buffer (more ticks) to complete.

		Collection 2	
		Run 3	Run 4
3729	Total Images	27478	9881
3731	Total Variations	29808	9936
3733	Image Loss (no screenshot taken)	0.078% (27478 from expected 29808)	0.006% (9881 from expected 9936)
3735	<u>Variations</u>		
3736	Models	23	23
3738	Angles	Radi = [200, 240, 290] Alpha = [-140, -105, -95, -85, -75, -55] Phi = [45, 55, 65, 80]	Radi = [200, 240, 290] Alpha = [-140, -105, -95, -85, -75, -55] Phi = [45, 55, 65, 80]
3739	(sets of angles that are unique to each run. Not shared)		
3741	Total Lighting settings	[75, 90, 140, 170]	[unlit]
3743	Total available Presents	6	6

Figure A.VI.5: Description of the second Cat\_Dog synthetic image collection.

3800  
3801  
3802  
3803 Cats - 10 Models + 3 model repeat with a different texture  
3804  
3805 Please see the References section for 3D model credit.  
3806  
3807

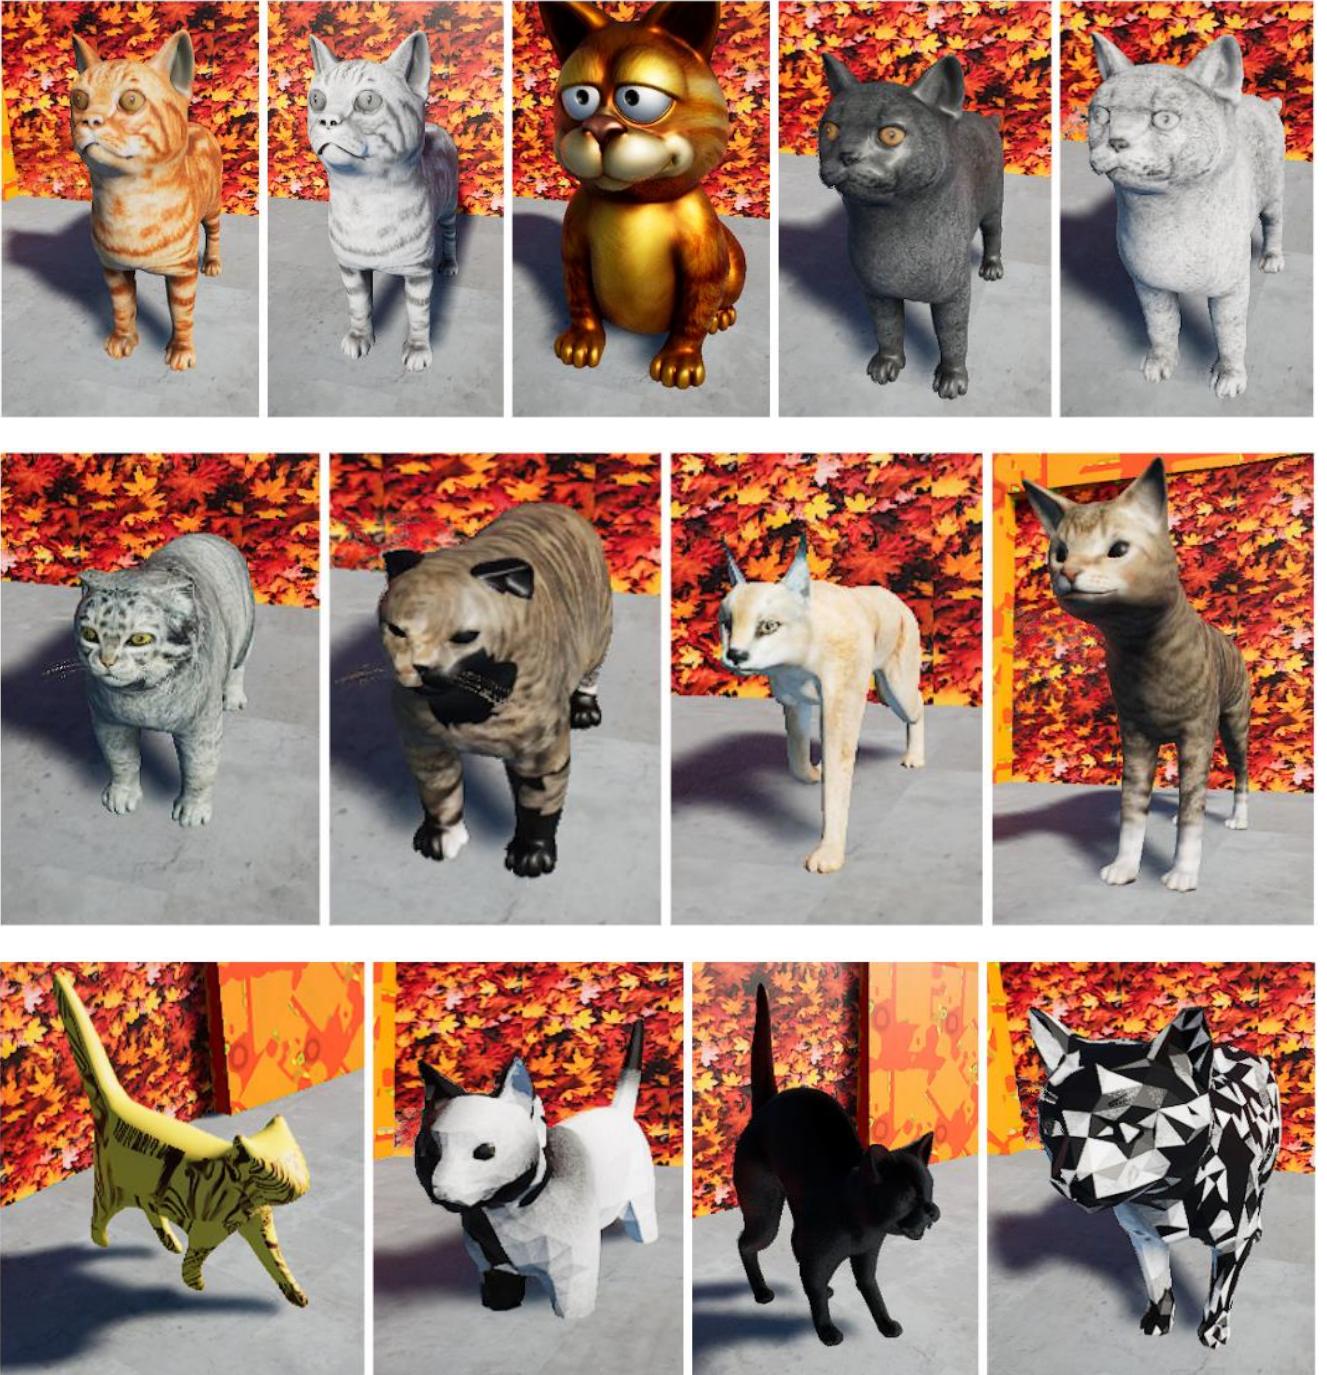


Figure A.VI.6: List of 3D model cats used in the Cat\_Dog synthetic image collection.

3850  
3851  
3852  
3853  
3854  
3855  
3856  
3857  
3858  
3859  
3860  
3861  
3862  
3863  
3864  
3865  
3866  
3867  
3868  
3869  
3870  
3871  
3872  
3873  
3874  
3875  
3876  
3877  
3878  
3879  
3880  
3881  
3882  
3883  
3884  
3885  
3886  
3887  
3888  
3889  
3890  
3891  
3892  
3893  
3894  
3895  
3896  
3897  
3898  
3899

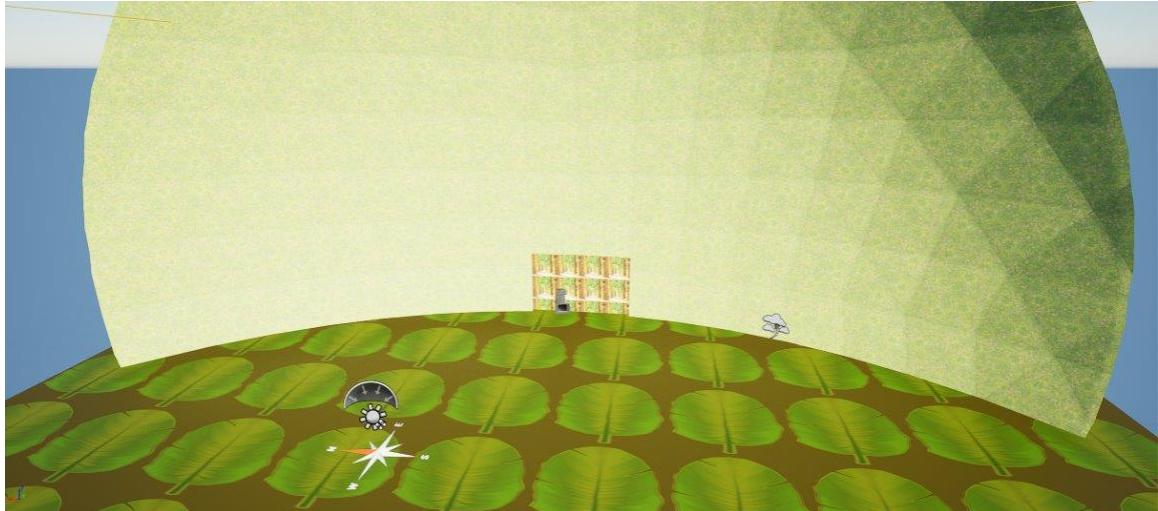
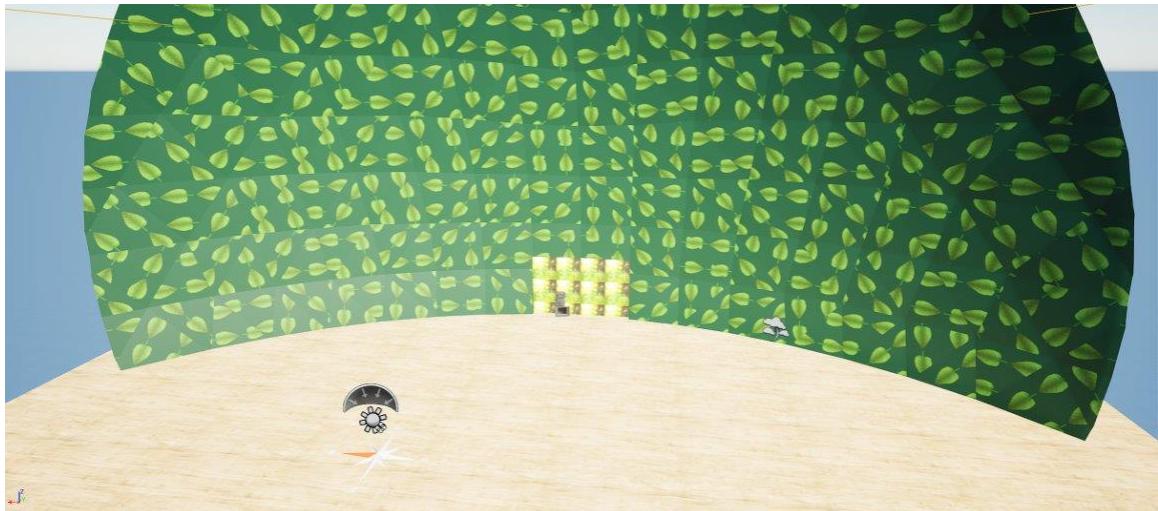
3900  
3901  
3902  
3903 Dogs - 9 Models + 1 model repeat with a different texture.  
3904  
3905 Please see the References section for 3D model credit.  
3906  
3907



3908  
3909  
3910  
3911  
3912  
3913  
3914  
3915  
3916  
3917  
3918  
3919  
3920  
3921  
3922  
3923  
3924  
3925  
3926  
3927  
3928  
3929  
3930  
3931  
3932  
3933  
3934  
3935  
3936  
3937  
3938  
3939  
3940  
3941  
3942  
3943  
3944  
3945  
3946  
3947  
3948  
3949  
3950  
3951  
3952  
3953  
3954  
3955  
3956  
3957  
3958  
3959  
3960  
3961  
3962  
3963  
3964  
3965  
3966  
3967  
3968  
3969  
3970  
3971  
3972  
3973  
3974  
3975  
3976  
3977  
3978  
3979  
3980  
3981  
3982  
3983  
3984  
3985  
3986  
3987  
3988  
3989  
3990  
3991  
3992  
3993  
3994  
3995  
3996  
3997  
3998  
3999

Figure A.VI.7: List of 3D model dogs used in the Cat\_Dog synthetic image collection.

4000  
4001  
4002  
4003 Pictures of Cat\_Dog Backgrounds  
4004  
4005 Set 01 (4 total)  
4006



4007  
4008  
4009  
4010  
4011  
4012  
4013  
4014  
4015  
4016  
4017  
4018  
4019  
4020  
4021  
4022  
4023  
4024  
4025  
4026  
4027  
4028  
4029  
4030  
4031  
4032  
4033  
4034  
4035  
4036  
4037  
4038  
4039  
4040  
4041  
4042 *Figure A.VI.8: Backdrops for the first Cat\_Dog synthetic image collection. (Top) Flowering garden, (bottom) woodland path.*  
4043  
4044  
4045  
4046  
4047  
4048  
4049  
4050  
4051  
4052  
4053  
4054  
4055  
4056  
4057  
4058  
4059  
4060  
4061  
4062  
4063  
4064  
4065  
4066  
4067  
4068  
4069  
4070  
4071  
4072  
4073  
4074  
4075  
4076  
4077  
4078  
4079  
4080  
4081  
4082  
4083  
4084  
4085  
4086  
4087  
4088  
4089  
4090  
4091  
4092  
4093  
4094  
4095  
4096  
4097  
4098  
4099

4100  
4101  
4102  
4103  
4104  
4105  
4106  
4107  
4108  
4109  
4110  
4111  
4112  
4113  
4114  
4115  
4116  
4117  
4118  
4119  
4120  
4121  
4122  
4123  
4124  
4125  
4126  
4127  
4128  
4129  
4130  
4131  
4132  
4133  
4134  
4135  
4136  
4137  
4138  
4139  
4140  
4141  
4142  
4143  
4144  
4145  
4146  
4147  
4148  
4149

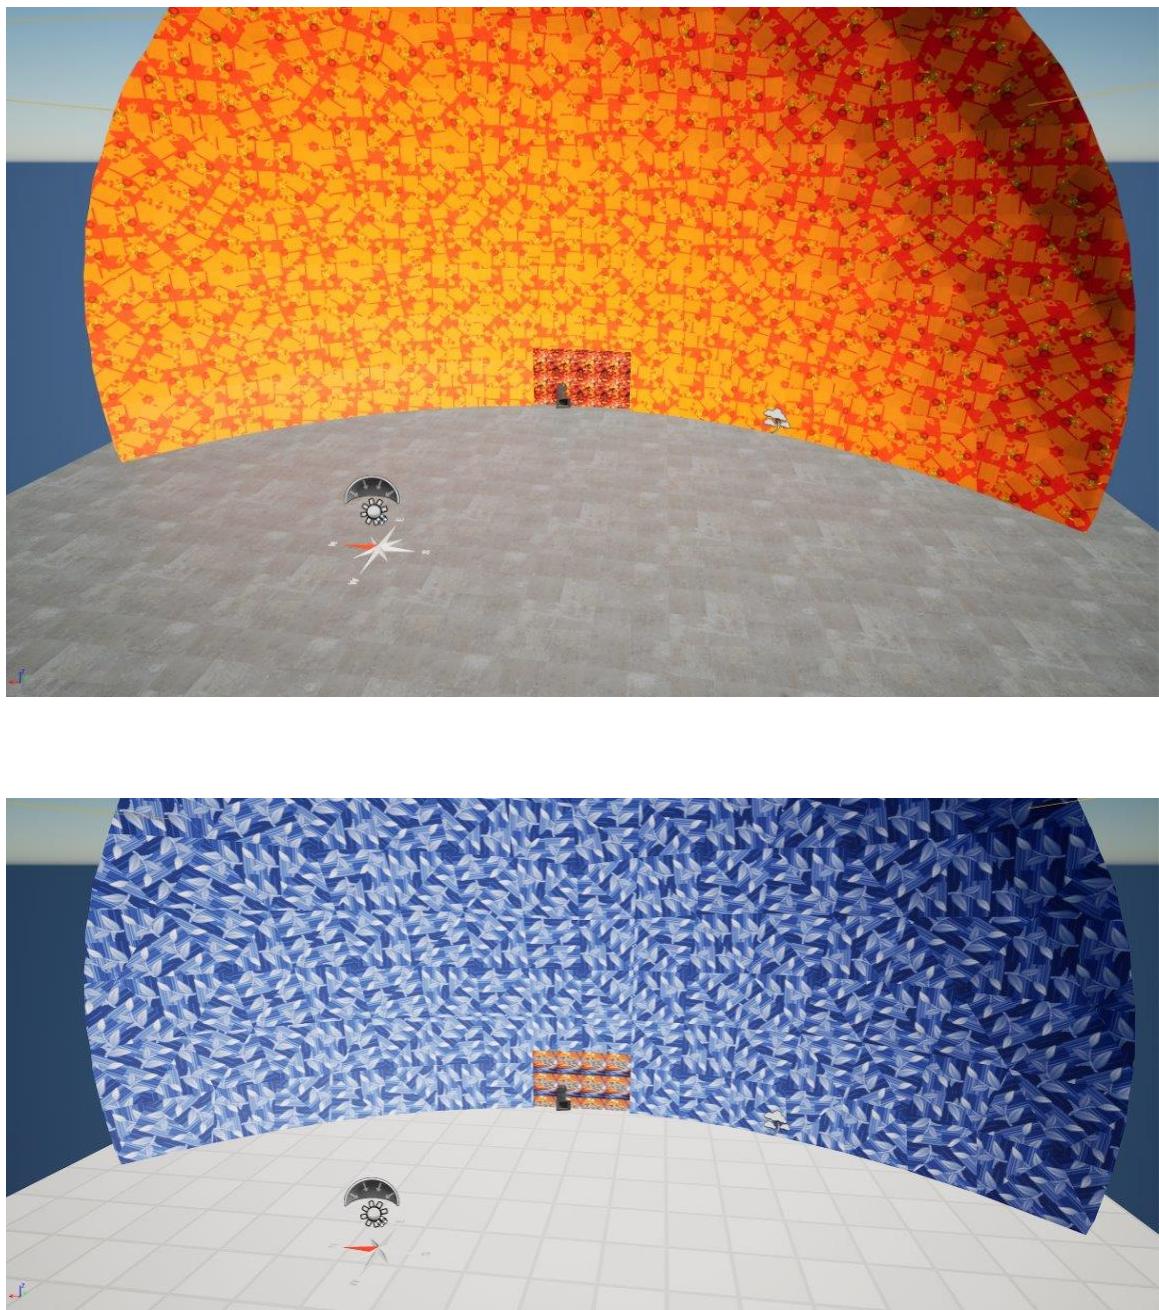


Figure A.VI.9: Backdrops for the first Cat\_Dog synthetic image collection. (Top) Fall leaves, (bottom) Futuristic interior.

4150  
4151  
4152  
4153  
4154  
4155  
4156  
4157  
4158  
4159  
4160  
4161  
4162  
4163  
4164  
4165  
4166  
4167  
4168  
4169  
4170  
4171  
4172  
4173  
4174  
4175  
4176  
4177  
4178  
4179  
4180  
4181  
4182  
4183  
4184  
4185  
4186  
4187  
4188  
4189  
4190  
4191  
4192  
4193  
4194  
4195  
4196  
4197  
4198  
4199

4200  
4201  
4202  
4203  
4204  
4205  
4206  
4207  
4208  
4209  
4210  
4211  
4212  
4213  
4214  
4215  
4216  
4217  
4218  
4219  
4220

Set 02 (6 total)



[PARKS & OPEN SPACE STAFF]



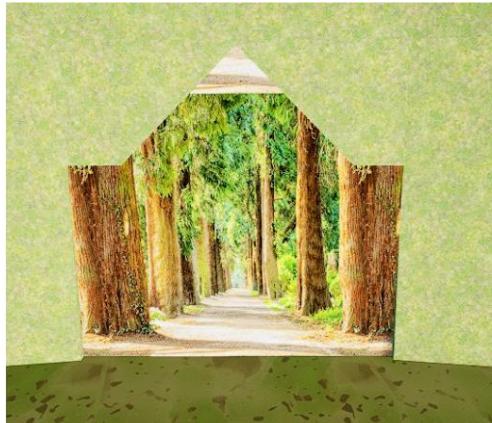
[KOSHKINA]



[PRITCHARD]



[WIESTOCK]



[HERRMANN]



[SmartFurniture]

4250  
4251  
4252  
4253  
4254  
4255  
4256  
4257  
4258  
4259  
4260  
4261  
4262  
4263  
4264  
4265  
4266  
4267  
4268  
4269  
4270

4271  
4272  
4273  
4274  
4275  
4276  
4277  
4278  
4279  
4280  
4281  
4282  
4283  
4284  
4285  
4286  
4287  
4288  
4289  
4290  
4291  
4292  
4293  
4294  
4295  
4296  
4297  
4298  
4299

Figure A.VI.10: Backdrops for the second Cat\_Dog synthetic image collection.

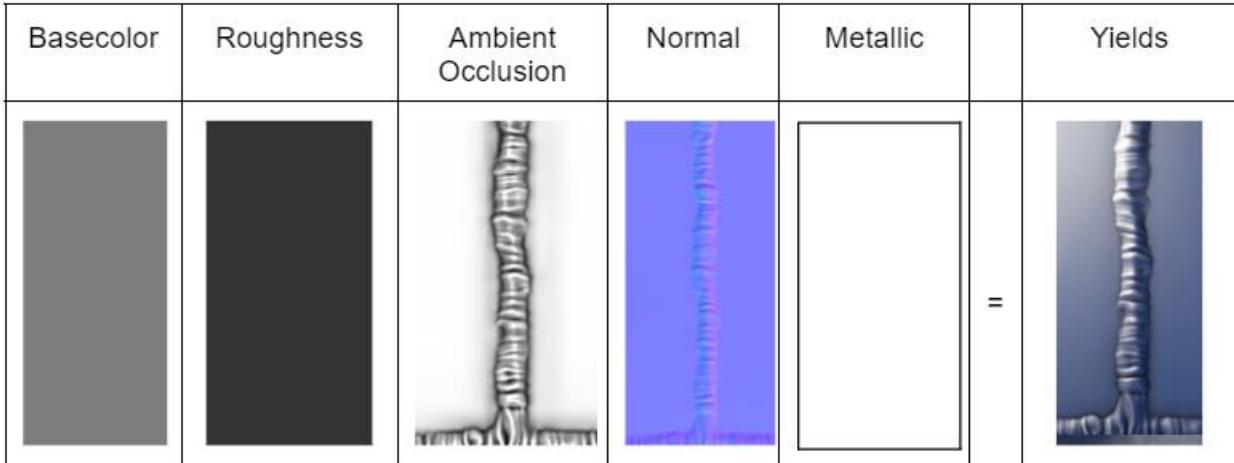
4300  
 4301  
 4302  
 4303 A.VII. Welding - Real Dataset collection  
 4304  
 4305  
 4306

4307 Below is a table showing how many images are found in each dataset, how many images are used in this experiment and  
 4308 reasons for why images were not used.

4309 Id	4310 Source	4311 Number of Used Images	4312 Number of Images in the Original dataset	4313 % Used	4314 Reasons for Removals and Edits
4315 01	4316 [weldefect]	4317 104	4318 408	4319 25.5%	4320 Repeat images with augmentations. Images with multiple welds that could not be cropped.
4321 02	4322 [Celebal Workspace]	4323 510	4324 1027	4325 49.7%	4326 Repeat images with augmentations.
4327 03	4328 [15DEM20F]	4329 173	4330 1184	4331 14.6%	4332 Repeated images from other datasets
4333 04	4334 [afiys]	4335 1279	4336 3038	4337 42.1%	4338 Repeat images with augmentations. Repeated images from other datasets
4339 05	4340 [Weld 1]	4341 156	4342 156	4343 100	4344 Cropped images to improve quality.

4345 *Figure A.VII.1: Descriptions of the found real welding datasets that form the 2221 real images.*

4346 A map preset refers to the combination of various graphical maps to yield a resulting image. In practice, a map is an image  
 4347 represented by [256, 256, 256] RGB values for each pixel in a bitmap file. A material being the result of these different  
 4348 images being used interpreted in different ways by the graphics engine. In this use case, each map preset is the combination  
 4349 of the Metallic, Roughness, Ambient Occlusion and Normal texture maps. By intentionally editing each of these maps,  
 4350 weld defects can be simulated in the UE. Please see the appendix for more detailed information on what effect each texture  
 4351 map has on the resulting material.



4364 *Figure A.VII.2: Demonstration of how different texture maps yield a UE Material.*

4400  
 4401  
 4402  
 4403  
 4404  
 4405  
 4406  
 4407  
 4408  
 4409  
 4410  
 4411  
 4412  
 4413  
 4414  
 4415  
 4416  
 4417  
 4418  
 4419  
 4420  
 4421  
 4422  
 4423  
 4424  
 4425  
 4426  
 4427  
 4428  
 4429  
 4430  
 4431  
 4432  
 4433  
 4434  
 4435  
 4436  
 4437  
 4438  
 4439  
 4440  
 4441  
 4442  
 4443  
 4444  
 4445  
 4446  
 4447  
 4448  
 4449

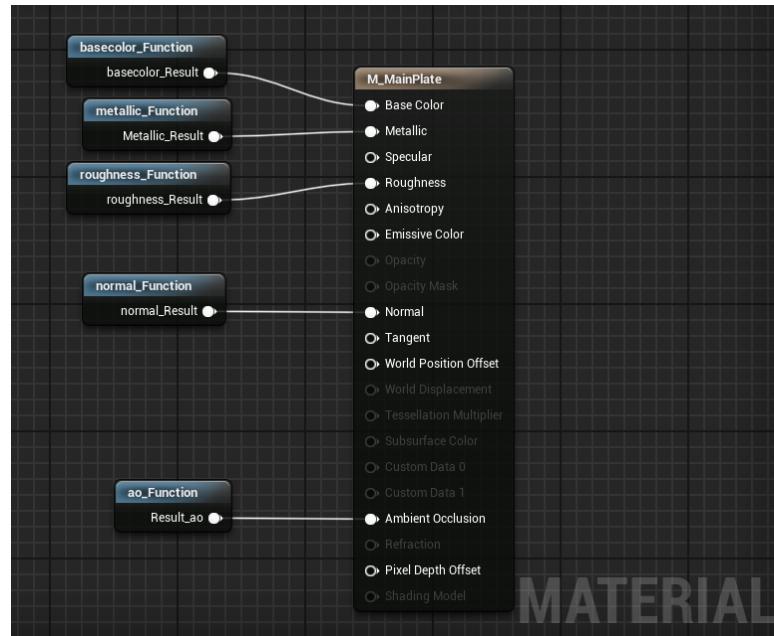


Figure A.VII.3: UE Material node view on how to create a map preset. This is repeated for all map presets.

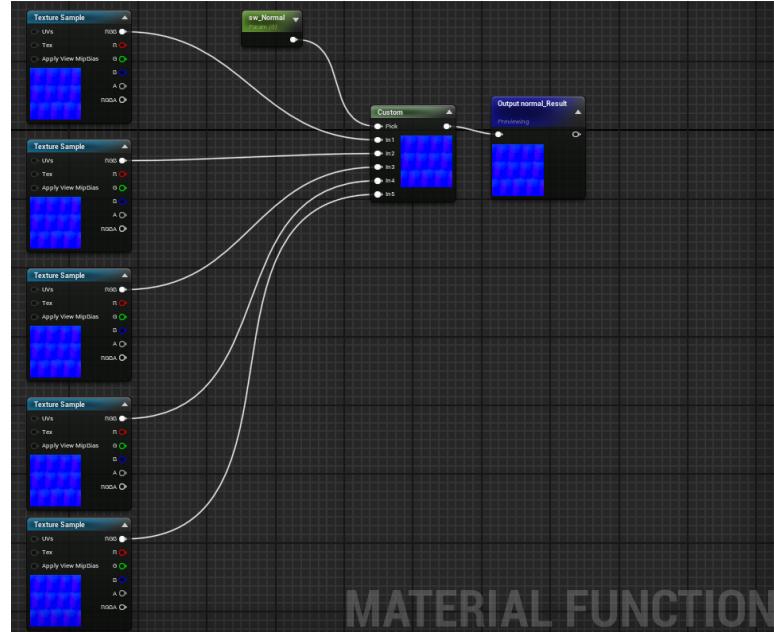


Figure A.VII.4: Example of a material function used in Figure A.VII.3 to create a map preset. The switch decides which texture map to use.

4450  
 4451  
 4452  
 4453  
 4454  
 4455  
 4456  
 4457  
 4458  
 4459  
 4460  
 4461  
 4462  
 4463  
 4464  
 4465  
 4466  
 4467  
 4468  
 4469  
 4470  
 4471  
 4472  
 4473  
 4474  
 4475  
 4476  
 4477  
 4478  
 4479  
 4480  
 4481  
 4482  
 4483  
 4484  
 4485  
 4486  
 4487  
 4488  
 4489  
 4490  
 4491  
 4492  
 4493  
 4494  
 4495  
 4496  
 4497  
 4498  
 4499

Collection Method

To collect all 8160 synthetic images, a camera and point light focused on points spread out on a flat mesh object. The purchased ‘Metal\_Weld’ material is applied to a flat rectangular plate in the UE test environment [4k metal pack]. Screenshots are taken along a set of points of interest where the weld seam is present in the original base material. The purchased material base texture size is defined from texture map bitmap files that are 4096x4096 pixels. No points of interest are defined beyond the 4096x4096 bitmap or within ~15 pixels of the image border (resulting in a 4066x4066 pixel area in which screenshots are focused). Starting at these defined borders, points of interest are recorded every 400 pixels along each horizontal and vertical weld seam. Resulting in 68 specific points in which defects are applied and the camera focuses on.

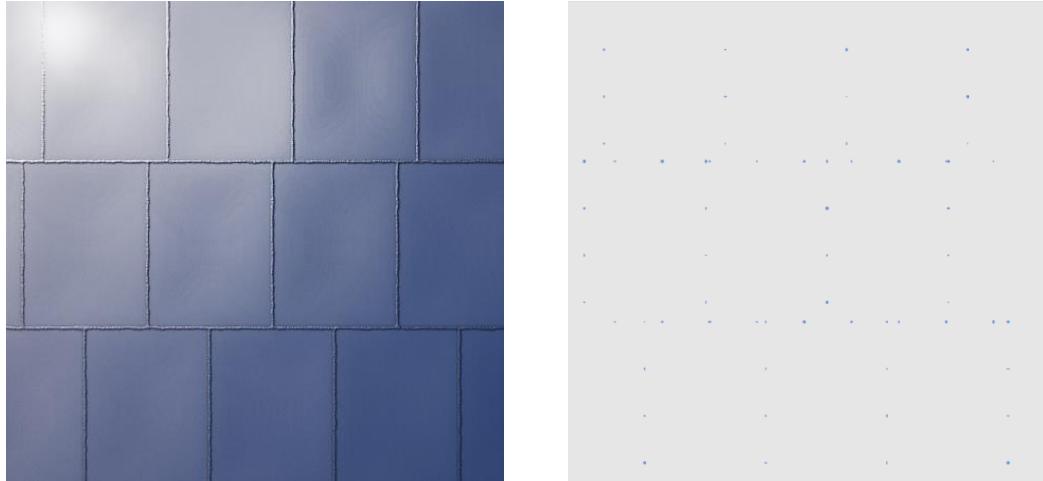


Figure A.VII.5: (left) top view of the base welding material. (right) dot locations of all 68 points of interest in which defects are created and screenshots taken.

Once the points of interest are defined, a custom python script sends commands to the UE Editor over the UE Python API each process frame tick. The process for creating the python scripts is documented earlier in this Dataset Formation section. The python script moves the UE Editor CineCamera object to each point of interest, where one of many actions in the collection sequence occurs; a screenshot is taken, the camera & light angle changes, the displayed metal plate changes color or the map preset changes. For every point of interest, four screenshots are taken from a camera and point light source that randomly is shifted in the x, y, and z directions. This allows for the displayed plate’s point of interest to have different camera angles and light behavior. This creates a unique screenshot for every plate variation. Across the entire collection process, screenshots are taken for each of the defined variations in plate colors, plate map sets, plate points of interest and point of interest camera angles.

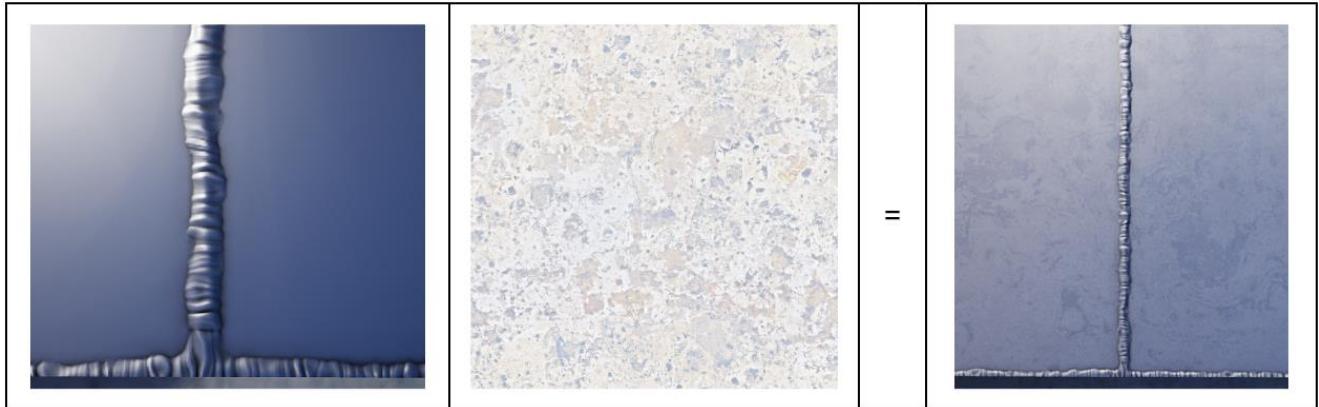
Plate Qualities	Variations
Plate base color	Aluminum, Steel, Copper, Gold, rusted metal, and a base metal
Plate map presets [Metallic, Normal, Ambient Occlusion, & Roughness]	1) no defects, 2) edits but not enough to be defective, 3) porosity defects, 4) crack defects, and 5) incomplete weld defects
Plate weld seam points of interest	68 (spaced 400 pixels apart)
Number of screenshots per plate point of interest (random Camera & Point Light offsets)	4

Figure A.VII.6: Overview of the variations used in the Weld Synthetic image collection.

Synthetic Weld Variation

Six base colors and five map presets are the basis for generating variation in the training set's synthetic images. Screenshots are taken when the weld plate is displaying each of these base colors and map presets. Collectively, the base color and map preset create unique defects at each of the 68 assigned points of interest (as shown above in the gray dotted square image in the previous section).

The various base colors highlight the different metals that could be practically (or impractically) welded with FILLER, Aluminum, Steel, Copper, Gold, rusted metal, and a base metal color. In addition, macro variations from the UE starter content are mixed with some colors to provide more variations. Additionally, the five map presets provide finer detailed changes to the welding material. Each map preset changes how the welded area appears at each point of interest. All defects have some element of randomness, so every point for every map preset has a variation of the defect, further increasing variation.



*Figure A.VII.7: Example of how using a LERP (linear interpolation) node to combine base colors can yield image variation.*

Defect Specific Texture Modifications

On the following pages are descriptions on how each weld defect model is created. Descriptions include which texture maps are edited and how to reach the shown effect.

4700  
 4701  
 4702  
**4703 Porosity**  
 4704  
 4705

4750

4751

4752

4753

4754

In the real weld dataset, a porosity defect is defined by elliptically shaped holes in the surface of a weld. The created weld UE defect textures use edits to the Ambient Occlusion and Normal maps to outline a feeling of depth to an elliptical shape.

4755

4756

<p>Ellipse logic to fill</p> <p>For pixel point <math>(x', y')</math> for ellipse with center <math>(x, y)</math> and radius <math>(r_x, r_y)</math></p> $d = \left( \frac{x'-x}{r_x} \right)^2 + \left( \frac{y'-y}{r_y} \right)^2$ <p>If <math>d \leq 1</math> then fill with color</p>	
---	--

Figure A.VII.8: Overview of how a pixel is determined to be inside of the designed ellipse to be in the porosity defect.

To add a difference in height to the hole, the RGB of the ellipse fill area is set as a radial gradient where the color in the Normal map is determined by the distance a pixel in the hole is to the ellipse center. light reflects and scatters as the pixel intensity changes around the ellipse center mimicking a hole. At every applied plate point of interest, a single ellipse is created with a random x-radius, y-radius, and xy-offset from the point of interest coordinate. Unfortunately, the UE defect does not replicate the trend in real images to create multiple porosity imperfections in a cluster.

4768

4769

4770

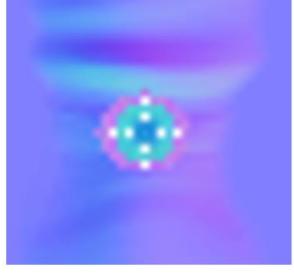
<p>Interpolation from Blue to Pink</p> <p>For pixel point <math>(x', y')</math> for ellipse center <math>(x, y)</math></p> $d = \sqrt{(x' - x)^2 + (y' - y)^2}$ <p>If <math>d &gt; 4</math> then color = PINK          Elif <math>2 &lt; d \leq 4</math> then color = LIGHT BLUE          Elif <math>d \leq 2</math> then color = BLUE</p> <p>Key</p> <ul style="list-style-type: none"> <li>- PINK = [242, 129, 189]</li> <li>- LIGHT BLUE = [220, 203, 83]</li> <li>- BLUE = [220, 158, 44]</li> </ul>	
--	---

Figure A.VII.9: Overview of how the Normal map's RGB intensity is selected for the Porosity Defect

Ellipse Characteristic	Range of the Defective Texture	Range of the Non-Defective Texture
Center x	[-3, 3]	[-3, 3]
Center y	[-3, 3]	[-3, 3]
Radius x	[2, 6]	[1, 2]
Radius y	[2, 6]	[1, 2]
Texture Map Pixel Color Noise	[-10, 10]	[-10, 10]

Figure A.VII.10: Outline of the design variables to create a Porosity defect and non-defect imperfection.

Crack

A crack defect is characterized by a continuous jagged line along a weld. The UE model defect attempts to replicate this by splitting a line, several pixels thick, into multiple points then moving each of those points off the original line in the positive or negative direction. The line is multiple pixels thick to create the illusion of depth. Each split line segment starts and end points are moved positively or negatively (depending on the direction of the weld seam) to create the jagged shape. All pixels inside of the line thickness are filled with the same value. The Ambient Occlusion and Normal texture maps create the shape of the crack while the Metallic and Roughness maps define how reflective the inside of the crack is.



Figure A.VII.11: Example of the crack defect. (left) 6 segments, (right) 8 segments

Crack Characteristic	Range of the Defective Texture	Range of the Non-Defective Texture
Length	[30, 80] pixels	[30, 80] Pixels
Number of Segments	[3, 10]	[3, 10]
Segment Offset	[4, 10]	[1, 1]
Line Thickness	[2, 4]	[1, 1]
Texture Map Pixel Color Noise	[-15, 15]	[-15, 15]

Figure A.VII.12: Outline of the design variables to create a crack defect and non-defect imperfection

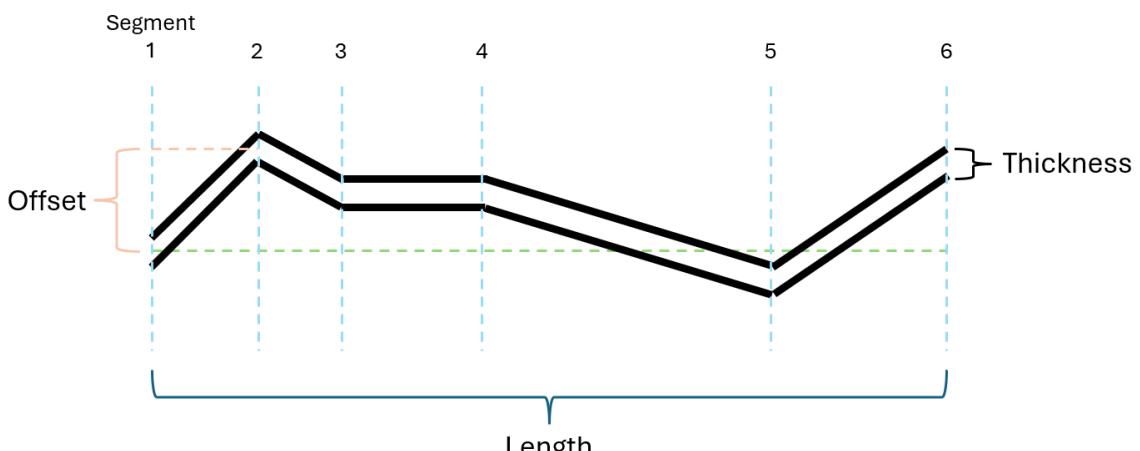


Figure A.VII.13: Visual representation of how a weld crack is defined by pixel. The shown crack is made up of a line with 6 segments with slight thickness variation and have positive offsets for segment point 2, & 6. Segment points 3 & 5 have negative offsets. Segment point 4 has an offset of zero.

Incomplete Weld

Incomplete welds are instances where the filler material is not completely continuous along a weld line. If enough filler is missing, the weld seam that you are attempting to weld could show. The UE models attempt to replicate this by inserting a rectangular area perpendicular to the weld seam. If the cut area extends past the weld midpoint, then a ‘seam’ is created in the Normal, Metallic, and Roughness texture maps. The seam is 1 pixel wide with 3 additional pixels on either side that is slightly more metallic than the center seam. These two areas represent the edge of the plate and the free space between the plates. The pixels in the rectangular region resample the pixel intensity of the pixel from the direction of the cut. As seen in the picture below, the cut is beyond the weld midpoint. The seam is showing and there is slight discoloration in the rectangular area that was cut out. This is the mismatching of pixels between what was resampled and what is around the overwritten pixel.



Figure A.VII.14: Example of the incomplete weld defect with the seam showing. (left) Partial incomplete weld, (right) missing weld.

Incomplete Weld Characteristic	Range of the Defective Texture	Range of the Non-Defective Texture
Gap Width	[3, 30]	[0, 30]
Gap Depth	[~40%, 100%]	[0%, 30%]
Texture Map Pixel Color Noise	[-10, 10]	[-10, 10]

Figure A.VII.15: Outline of the design variables to create an incomplete weld defect and non-defect imperfection



### A.VIII. Model Training Charts

The following are examples of the Accuracy and Loss training charts for the various experiments.

Experiment 1 – Cat\_Dog - 9425 Synthetic Images available

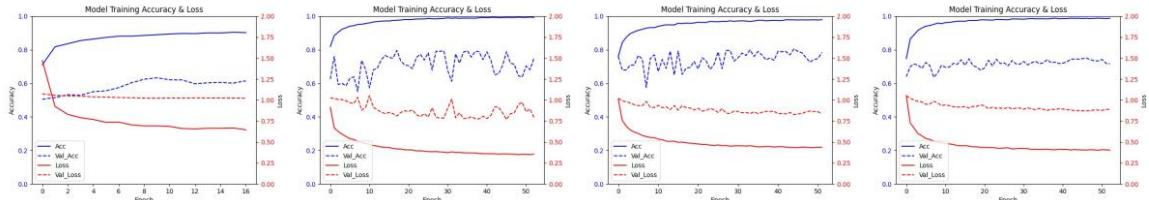


Figure A.VII.1: Accuracy and Loss training graphs for VGG16, RSR (left) 0:100, (middle left) 75:25, (middle right) 50:50, (right) 05:95

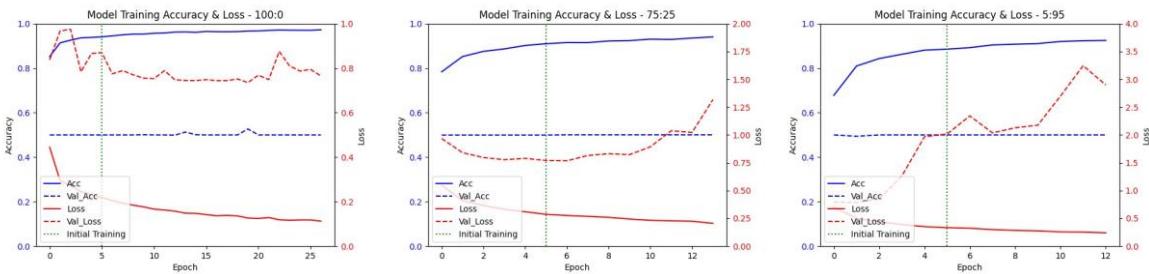


Figure A.VII.2: Accuracy and Loss training graphs for MobileNet, RSR (left) 0:100, (middle) 75:25, (right) 05:95. As discussed in the Discussion section of the paper, the MobileNet model struggled with training in several aspects. Although the training and test gap for loss was about the same as VGG16, the MobileNet model's absolute minima was so sensitive that the model diverged continuously during training. This is why the five epochs of initial training were put in place (to give the model more time to find a local minima). Experimenting with different learning rate and batch size combinations did not resolve this. In addition, the Kerasv2.12 implementation of the MobileNet model did not respond with any change to Validation Accuracy (blue dashed line). This is most irregular and, during this project, was present in MobileNetV3 variations (small, large, large-minimal).

Experiment 4 – Welding - 8160 Synthetic Images available

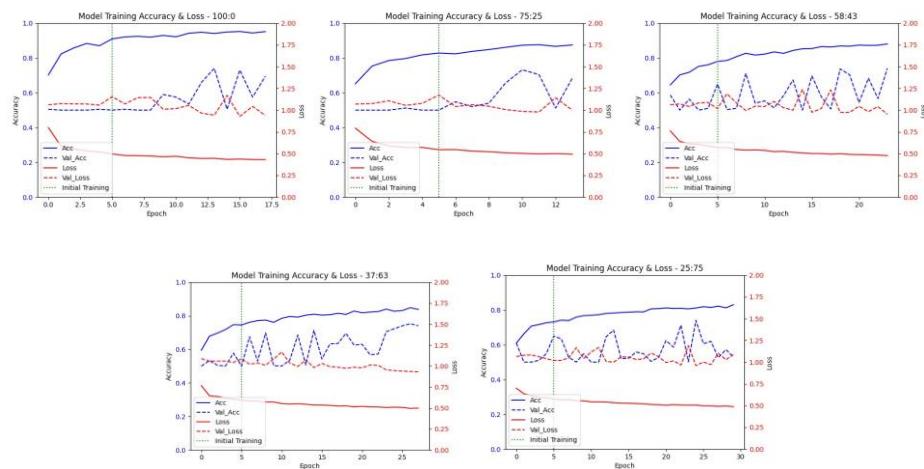


Figure A.VII.3: Accuracy and Loss training graphs for VGG16. All show a struggle to reduce Validation Loss, especially as the RSR decreases.

5200  
 5201  
 5202  
 5203  
 5204  
 5205  
 5206  
 5207  
 5208  
 5209  
 5210  
 5211  
 5212  
 5213  
 5214  
 5215  
 5216  
 5217  
 5218  
 5219  
 5220  
 5221  
 5222  
 5223  
 5224  
 5225  
 5226  
 5227  
 5228  
 5229  
 5230  
 5231  
 5232  
 5233  
 5234  
 5235  
 5236  
 5237  
 5238  
 5239  
 5240  
 5241  
 5242  
 5243  
 5244  
 5245  
 5246  
 5247  
 5248  
 5249

## A.IX. Results

Definition of Experiments 1 & 2

The tables below detail how many images of each type are used for Cat\_Dog Experiments 1 & 2 analyzing how using synthetic images and synthetic data availability affect model performance.

Real:Synth Ratio	Total Images	Required Number of Real Images	Data Split	Number of real <b>TRAINING</b> images	Number of synthetic <b>TRAINING</b> images	Number of real <b>VALIDATION</b> images	Number of real <b>TEST</b> images
100:0	10000	10000	80/10/10	8000	0	1000	1000
75:25	10000	8000	80/10/10	6000	2000	1000	1000
50:50	10000	6000	80/10/10	4000	4000	1000	1000
25:75	10000	4000	80/10/10	2000	6000	1000	1000
10:90	10000	3000	80/10/10	800	7200	1000	1000
05:95	10000	2800	80/10/10	400	7600	1000	1000
0:100	10000	2000	80/10/10	0	8000	1000	1000

Figure A.IX.1: Definition of Experiment 1. Mix of real and synthetic data in the training set. Always 1000 real images in the validation and test sets.

Real:Synth Ratio	Total Images	Number of <b>Real</b> Training Images	Number of <b>Synthetic</b> Training Images	Synthetic Utilization Rate (9425 Synthetic Images Available)	Synthetic Utilization Rate (46784 Synthetic Images Available)
100:0	10000	8000	0	0%	0%
75:25	10000	6000	2000	21.2%	4.2%
50:50	10000	4000	4000	42.4%	8.4%
25:75	10000	2000	6000	63.6%	12.6%
10:90	10000	800	7200	76.4%	15.0%
05:95	10000	400	7600	80.6%	15.9%
0:100	10000	0	8000	84.8%	16.8%

Figure A.IX.2: Definition of Experiment 1 expanded to Experiment 2. Details of the SUV for each test in Figure A.IX.1.

5300  
 5301  
 5302  
 5303 Raw Data - Experiments 1 & 2  
 5304  
 5305  
 5306 Experiment 1 - 9425 Synthetic Images Available  
 5307  
 5308  
 5309

5350

5351

5352

5353

5354

5355

5356

5357

5358

5359

5360

5361

5362

5363

5364

5365

5366

5367

5368

5369

5370

5371

5372

5373

5374

VGG16							
RSR	Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.950	0.00178	0.959	0.00010	-	-
75:25	8000	0.984	0.00009	0.965	0.00016	0.0060	58.2%
50:50	6000	0.961	0.00061	0.971	0.00001	0.0113	-90.8%
25:75	4000	0.986	0.00009	0.963	0.00003	0.0040	-66.1%
10:90	3000	0.963	0.00038	0.949	0.00007	-0.0107	-34.5%
05:95	2800	0.972	0.00035	0.945	0.00023	-0.0140	126.3%
0:100	2000	0.972	0.00027	0.851	0.00026	-0.1080	154.9%

Figure A.IX.3: Raw data

VGG16							
RSR	Required Number of Real Images	Avg Train_Loss	Variance	Avg Test_Loss	Variance	Loss change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.504	0.01584	0.867	0.01621	-	-
75:25	8000	0.387	0.00109	0.767	0.00079	-0.0997	-95.2%
50:50	6000	0.478	0.00481	0.875	0.00397	0.0077	-75.5%
25:75	4000	0.386	0.00175	0.806	0.00120	-0.0613	-92.6%
10:90	3000	0.472	0.00449	0.897	0.00406	0.0300	-75.0%
05:95	2800	0.454	0.00375	0.949	0.00858	0.0823	-47.1%
0:100	2000	0.445	0.00368	0.995	0.00089	0.1276	-94.5%

Figure A.IX.4: Raw data

5337  
 5338  
 5339  
 5340  
 5341  
 5342  
 5343  
 5344  
 5345  
 5346  
 5347  
 5348  
 5349

5387

5388

5389

5390

5391

5392

5393

5394

5395

5396

5397

5398

5399

MobileNetV3-small							
RSR	Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.959	0.00002	0.905	0.00000	-	-
75:25	8000	0.915	0.00007	0.874	0.00003	-0.032	1200.00%
50:50	6000	0.912	0.00044	0.870	0.00022	-0.035	621.98%
25:75	4000	0.895	0.00000	0.839	0.00065	-0.067	197.87%
10:90	3000	0.903	0.00006	0.830	0.00023	-0.076	-64.69%
05:95	2800	0.901	0.00000	0.803	0.00051	-0.103	119.39%
0:100	2000	0.903	0.00004	0.538	0.00063	-0.367	25.13%

Figure A.IX.5: Raw data

MobileNetV3-small							
RSR	Required Number of Real Images	Avg Train_Loss	Variance	Avg Test_Loss	Variance	Loss change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.160	0.00040	0.854	0.01116	-	-
75:25	8000	0.273	0.00065	0.782	0.00005	-0.0713	-99.6%
50:50	6000	0.277	0.00313	0.850	0.01196	-0.0037	7.2%
25:75	4000	0.327	0.00002	1.107	0.09706	0.2530	769.9%
10:90	3000	0.308	0.00043	1.106	0.06910	0.2523	519.3%
05:95	2800	0.308	0.00007	1.049	0.01413	0.1953	26.6%
0:100	2000	0.301	0.00008	1.272	0.11135	0.4187	898.0%

Figure A.IX.6: Raw data

Experiment 2 - 46784 Synthetic Images Available

VGG16							
RSR	Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.974	0.00011	0.967	0.00006	-	-
75:25	8000	0.975	0.00005	0.955	0.00018	-0.0117	203.3%
50:50	6000	0.972	0.00003	0.960	0.00000	-0.0070	-96.1%
25:75	4000	0.975	0.00045	0.953	0.00024	-0.0133	305.0%
10:90	3000	0.969	0.00071	0.952	0.00038	-0.0143	530.4%
05:95	2800	0.958	0.00078	0.946	0.00009	-0.0210	43.1%
0:100	2000	0.982	0.00007	0.826	0.00014	-0.1403	132.6%

Figure A.IX.7: Raw data

VGG16							
RSR	Required Number of Real Images	Avg Train_Loss	Variance	Avg Test_Loss	Variance	Loss change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.432	0.00190	0.857	0.00080	-	-
75:25	8000	0.432	0.00023	0.816	0.00042	-0.0410	-47.7%
50:50	6000	0.439	0.00056	0.845	0.00429	-0.0123	433.5%
25:75	4000	0.423	0.00656	0.814	0.01555	-0.0437	1833.6%
10:90	3000	0.442	0.00761	0.850	0.00667	-0.0070	729.3%
05:95	2800	0.484	0.00586	0.907	0.00408	0.0493	407.8%
0:100	2000	0.398	0.00075	0.978	0.00104	0.1207	29.2%

Figure A.IX.8: Raw data

MobileNetV3-small							
RSR	Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.965	0.00005	0.914	0.00036	-	-
75:25	8000	0.910	0.00002	0.883	0.00062	-0.030	70.81%
50:50	6000	0.899	0.00009	0.885	0.00007	-0.029	-89.23%
25:75	4000	0.886	0.00004	0.855	0.00110	-0.059	1540.30%
10:90	3000	0.892	0.00004	0.849	0.00045	-0.065	-58.78%
05:95	2800	0.888	0.00001	0.810	0.00047	-0.104	3.97%
0:100	2000	0.900	0.00006	0.605	0.00021	-0.308	-56.19%

Figure A.IX.9: Raw data

MobileNetV3-small							
RSR	Required Number of Real Images	Avg Train_Loss	Variance	Avg Test_Loss	Variance	Loss change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	10000	0.139	0.00073	0.748	0.00032	-	-
75:25	8000	0.284	0.00009	0.777	0.00006	0.0287	-80.3%
50:50	6000	0.310	0.00033	1.171	0.00944	0.4230	2811.0%
25:75	4000	0.343	0.00009	1.567	0.29963	0.8183	92284.7%
10:90	3000	0.328	0.00013	1.850	0.14130	1.1020	43465.5%
05:95	2800	0.334	0.00000	1.952	0.08354	1.2037	25657.1%
0:100	2000	0.310	0.00026	1.819	0.01907	1.0710	5778.3%

Figure A.IX.10: Raw data

5700  
 5701  
 5702  
 5703  
 Definition of Experiment 3

5704  
 5705 Cat\_Dog Experiment 3 is a baseline measurement of how the models perform if no synthetic images are used, but the dataset of real  
 5706 images decreases.

Total Images	Data Split	Percentage of TRAINING images are Real	Number of synthetic TRAINING images	Number of real TRAINING images	Number of real VALIDATION images	Number of real TEST images
7500	80/10/10	100	0	6000	750	750
5000	80/10/10	100	0	4000	500	500
2500	80/10/10	100	0	2000	250	250
1000	80/10/10	100	0	800	100	100
500	80/10/10	100	0	400	50	50

5718 Figure A.IX.11: Definition of Experiment 3 comparing mixed datasets with smaller datasets but only real images. The mixed dataset test results are  
 5719 collected from Experiments 1 & 2

5720  
 5721 Raw Data - Experiments 3

VGG16								
Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Avg Train_Loss	Variance	Avg Test_Loss	Variance
10000	0.967	0.00005	0.957	0.00023	0.462	0.00065	0.878	0.00293
8000	0.978	0.00014	0.970	0.00010	0.432	0.00172	0.852	0.00235
6000	0.984	0.00002	0.963	0.00043	0.421	0.00011	0.883	0.01554
4000	0.982	0.00001	0.953	0.00013	0.438	0.00035	0.909	0.00130
3000	0.933	0.00162	0.940	0.00480	0.575	0.01213	1.040	0.00430

5734 Figure A.IX.12: Raw data

MobileNetV3-small								
Required Number of Real Images	Avg Train_Acc	Variance	Avg Test_Acc	Variance	Avg Train_Loss	Variance	Avg Test_Loss	Variance
10000	0.954	0.00006	0.897	0.00023	0.184	0.00065	0.784	0.00307
8000	0.950	0.00000	0.897	0.00056	0.196	0.00008	0.835	0.01024
6000	0.926	0.00008	0.897	0.00103	0.272	0.00038	0.998	0.03269
4000	0.928	0.00006	0.930	0.00010	0.263	0.00068	1.069	0.04207
3000	0.937	0.00126	0.893	0.00053	0.266	0.00771	1.021	0.09085

5735 Figure A.IX.13: Raw data

5736  
 5737  
 5738  
 5739  
 5740  
 5741  
 5742  
 5743  
 5744  
 5745  
 5746  
 5747  
 5748  
 5749

5750  
 5751  
 5752  
 5753  
 5754  
 5755  
 5756  
 5757  
 5758  
 5759  
 5760  
 5761  
 5762  
 5763  
 5764  
 5765  
 5766  
 5767  
 5768  
 5769  
 5770  
 5771  
 5772  
 5773  
 5774  
 5775  
 5776  
 5777  
 5778  
 5779  
 5780  
 5781  
 5782  
 5783  
 5784  
 5785  
 5786  
 5787  
 5788  
 5789  
 5790  
 5791  
 5792  
 5793  
 5794  
 5795  
 5796  
 5797  
 5798  
 5799

Welding

Definition of Experiment 4 - Change in RSR values

For every test in Experiment 4, all available 2000 real images are used. As such, to maintain the 80/10/10 split, real images are transferred into the Validation and Test sets. The RSR is only reduced to 25:75 due to the results of 0:100 tests in Experiments 1-3 and that it is desired to always have real images for this test. It is possible to have a test of 0:100 where all real images are in the validation and test set with synthetic images still imbalanced at 7:3.

Real:Synth Ratio	Total Number of Images	# of GOOD	# of DEFECT	Data Split	# of real TRAINING	# of real VALIDATION	# of real TESTING
100:0	2000	600	1400	80/10/10	1600	200	200
75:25	2500	750	1750	80/10/10	1500	250	250
58:42	3000	900	2100	80/10/10	1400	300	300
37:63	4000	1200	2800	80/10/10	1200	400	400
25:75	5000	1500	3500	80/10/10	1000	500	500

Figure A.IX.14: Definition of Experiment 4. As the RSR decreases the number of real images in the validation and test sets increase.

Real:Rynth Ratio	Total Number of Images	Train Images: REAL	# of Training Real GOOD	# of Training Real Defect	Train Images: Synthetic	# of Added Synthetic GOOD	# of Added Synthetic DEFECT
100:0	2000	1600	400	1200	0	0	0
75:25	2500	1500	350	1150	500	250	250
58:42	3000	1400	300	1100	1000	420	580
37:63	4000	1200	200	1000	2000	760	1240
25:75	5000	1000	100	900	3000	900	2100

Figure A.IX.15: Definition of Experiment 4 continued. Shows the class breakdown for each class in the imbalanced dataset.

Raw Data

VGG16						
Real:Rynth Ratio	Total Number of Images	Avg Train_Acc	Avg Test_Acc	Variance Test_Acc	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.9229	0.7967	0.00823	-	-
75:25	2500	0.8618	0.8433	0.00223	0.0467	-72.9%
58:42	3000	0.8254	0.8300	0.00000	0.0333	-100.0%
37:63	4000	0.8197	0.8300	0.00030	0.0333	-96.4%
25:75	5000	0.7835	0.7933	0.00003	-0.0033	-99.6%

Figure A.IX.16: Raw data

VGG16						
Real:Rynth Ratio	Total Number of Images	Avg Train_Loss	Avg Test_Loss	Variance Test_Loss	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.4895	1.0880	0.00462	-	-
75:25	2500	0.5062	0.9873	0.00311	-0.1008	-32.6%
58:42	3000	0.5320	1.0186	0.00051	-0.0695	-89.0%
37:63	4000	0.5198	0.9896	0.00193	-0.0985	-58.1%
25:75	5000	0.5275	1.0217	0.00201	-0.0664	-56.4%

Figure A.IX.17: Raw data

VGG16						
Real:Rynth Ratio	Total Number of Images	Avg Train_AUC	Avg Test_AUC	Variance Test_AUC	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.9669	0.8112	0.00026	-	-
75:25	2500	0.9386	0.8328	0.00036	0.0216	38.7%
58:42	3000	0.9100	0.8279	0.00093	0.0167	253.7%
37:63	4000	0.9106	0.7821	0.00025	-0.0291	-4.3%
25:75	5000	0.8767	0.7439	0.00009	-0.0673	-65.7%

Figure A.IX.18: Raw data

VGG16						
Real:Rynth Ratio	Total Number of Images	Avg TP (%)	Avg TN (%)	Avg FP (%)	Avg FN (%)	Cost per Test Sample
100:0	2000	42.8	43.5	7.2	6.5	10.3
75:25	2500	45.5	38.8	4.5	11.2	6.8
58:42	3000	44.7	38.4	5.3	11.6	7.5
37:63	4000	47.1	35.8	2.9	14.2	6.8
25:75	5000	47.7	31.6	2.3	18.4	7.1

Figure A.IX.19: Raw data

MobileNetV3-small						
Real:Rynth Ratio	Total Number of Images	Avg Train_Acc	Avg Test_Acc	Test_Acc Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.8642	0.7833	0.01003	-	-
75:25	2500	0.8428	0.7900	0.00190	0.0067	-81.06%
58:42	3000	0.7678	0.7833	0.00263	0.0000	-73.75%
37:63	4000	0.7367	0.7167	0.00203	-0.0667	-79.73%
25:75	5000	0.6991	0.6467	0.00053	-0.1367	-94.68%

Figure A.IX.20: Raw data

MobileNetV3-small						
Real:Rynth Ratio	Total Number of Images	Avg Train_Loss	Avg Test_Loss	Test_Loss Variance	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.2418	0.8312	0.00037	-	-
75:25	2500	0.1765	1.0525	0.08723	0.2213	23569.36%
58:42	3000	0.2675	0.8916	0.01107	0.0603	2904.19%
37:63	4000	0.2774	0.8760	0.01654	0.0447	4387.65%
25:75	5000	0.2806	0.8096	0.00192	-0.0216	421.92%

Figure A.IX.21: Raw data

MobileNetV3-small						
Real:Rynth Ratio	Total Number of Images	Avg Train_AUC	Avg Test_AUC	Variance Test_AUC	Accuracy change compared to RSR 100:0	Variance change compared to RSR 100:0
100:0	2000	0.9363	0.5773	0.00116	-	-
75:25	2500	0.9309	0.5417	0.00305	-0.0356	162.07%
58:42	3000	0.8641	0.5657	0.00181	-0.0117	55.59%
37:63	4000	0.8269	0.5556	0.00009	-0.0217	-92.23%
25:75	5000	0.7701	0.5696	0.00010	-0.0077	-91.18%

Figure A.IX.22: Raw data

MobileNetV3-small						
Real:Rynth Ratio	Total Number of Images	Avg TP (%)	Avg TN (%)	Avg FP (%)	Avg FN (%)	Cost per Sample
100:0	2000	36.1	42.0	13.9	8.0	10.933
75:25	2500	45.6	33.3	4.4	16.7	8.600
58:42	3000	47.0	31.3	3.0	18.7	8.333
37:63	4000	48.5	23.0	1.5	27.0	10.167
25:75	5000	49.5	15.0	0.5	35.0	10.833

Figure A.IX.23: Raw data

6100  
 6101  
 6102  
 6103 Definition of Experiment 5 - Rebalancing Imbalanced Datasets with Synthetic data  
 6104  
 6105

case	Total Number of Images	# of GOOD	# of DEFECT	# of TRAINING	# of VALIDATION	# of TESTING
No Change	1800	600	1400	1600	200	200
Cost-Learning	1800	600	1400	1600	200	200
Undersample	1200	600	600	960	120	120
Oversample	2400	1200	1200	1920	240	240
Synthetic Padding	2400	1200	1200	1920	240	240

6106  
 6107  
 6108  
 6109  
 6110  
 6111  
 6112  
 6113  
 6114  
 6115  
 6116  
 6117  
 6118  
 6119  
 6120  
 6121  
 6122  
 6123  
 6124  
 6125  
 6126  
 6127  
 6128  
 6129  
 6130  
 6131  
 6132  
 6133  
 6134  
 6135  
 6136  
 6137  
 6138  
 6139  
 6140  
 6141  
 6142  
 6143  
 6144  
 6145  
 6146  
 6147  
 6148  
 6149  
 6150  
 6151  
 6152  
 6153  
 6154  
 6155  
 6156  
 6157  
 6158  
 6159  
 6160  
 6161  
 6162  
 6163  
 6164  
 6165  
 6166  
 6167  
 6168  
 6169  
 6170  
 6171  
 6172  
 6173  
 6174  
 6175  
 6176  
 6177  
 6178  
 6179  
 6180  
 6181  
 6182  
 6183  
 6184  
 6185  
 6186  
 6187  
 6188  
 6189  
 6190  
 6191  
 6192  
 6193  
 6194  
 6195  
 6196  
 6197  
 6198  
 6199  
 Figure A.IX.24: Definition of Experiment 5. General description of what images are used.

Figure A.IX.25: Definition of Experiment 5. Further breakdown of images used in each rebalancing method.  
 \*In the oversample method, all 480 available Good images not used in the other sets are duplicated to yield 960 images.

case	Total Number of Images	Train Images: REAL	# of Training Real GOOD	# of Training Real Defect	Train Images: Synthetic	# of Added Synthetic GOOD	# of Added Synthetic DEFECT
No Change	1800	1600	400	1200	0	0	0
Cost-Learning	1800	1600	400	1200	0	0	0
Undersample	1200	960	480	480	0	0	0
Oversample	2400	1920	960*	960	0	0	0
Synthetic Padding	2400	1320	360	960	600	600	0

6200  
6201  
6202  
6203 Raw Data  
6204  
6205  
6206

VGG16					
case	Avg Test_Acc	Avg Train_AUC	Avg Test_AUC	Test_AUC Variance	Effect of Sampling Method compared to the imbalanced dataset
No Change	0.8433	0.9847	0.8519	0.0001	-
Cost-Learning	0.8433	0.9843	0.8257	0.0009	-0.0263
Undersample	0.8633	0.9660	0.8234	0.0002	-0.0285
Oversample	0.8733	0.9721	0.7676	0.0017	-0.0844
Synthetic Padding	0.8633	0.9925	0.8137	0.0050	-0.0382

Figure A.IX.26: Raw data

VGG16				
case	Avg Train_Loss	Avg Test_Loss	Test_Loss Variance	Effect of Sampling Method compared to the imbalanced dataset
No Change	0.4369	0.8433	0.0000	-
Cost-Learning	0.4385	0.8433	0.0082	0.0000
Undersample	0.4701	0.8633	0.0000	0.0200
Oversample	0.4841	0.8733	0.0006	0.0300
Synthetic Padding	0.4151	0.8633	0.0037	0.0200

Figure A.IX.27: Raw data

VGG16				
case	Avg Cost per Test Sample	Avg TP+TN (%)	Avg FP (%)	Avg FN (%)
No Change	0.0780	0.844	0.0707	0.0850
Cost-Learning	0.0717	0.842	0.0573	0.1003
Undersample	0.0697	0.795	0.1247	0.0807
Oversample	0.0640	0.872	0.1010	0.0267
Synthetic Padding	0.0687	0.862	0.0640	0.0737

Figure A.IX.28: Raw data

6200  
6201  
6202  
6203 Raw Data  
6204  
6205  
6206  
6207  
6208  
6209  
6210  
6211  
6212  
6213  
6214  
6215  
6216  
6217  
6218  
6219  
6220  
6221  
6222  
6223  
6224  
6225  
6226  
6227  
6228  
6229  
6230  
6231  
6232  
6233  
6234  
6235  
6236  
6237  
6238  
6239  
6240  
6241  
6242  
6243  
6244  
6245  
6246  
6247  
6248  
6249

6250  
6251  
6252  
6253  
6254  
6255  
6256  
6257  
6258  
6259  
6260  
6261  
6262  
6263  
6264  
6265  
6266  
6267  
6268  
6269  
6270  
6271  
6272  
6273  
6274  
6275  
6276  
6277  
6278  
6279  
6280  
6281  
6282  
6283  
6284  
6285  
6286  
6287  
6288  
6289  
6290  
6291  
6292  
6293  
6294  
6295  
6296  
6297  
6298  
6299

MobileNetV3-small					
case	Avg Test_Acc	Avg Train_AUC	Avg Test_AUC	Test_AUC Variance	Effect of Sampling Method compared to the imbalanced dataset
No Change	0.8100	0.9847	0.8519	0.0001	-
Cost-Learning	0.8100	0.9843	0.8257	0.0009	-0.0263
Undersample	0.7967	0.9660	0.8234	0.0002	-0.0285
Oversample	0.7967	0.9721	0.7676	0.0017	-0.0844
Synthetic Padding	0.8100	0.9925	0.8137	0.0050	-0.0382

Figure A.IX.29: Raw data

MobileNetV3-small				
case	Avg Train_Loss	Avg Test_Loss	Test_Loss Variance	Effect of Sampling Method compared to the imbalanced dataset
No Change	0.4369	0.8433	0.0000	-
Cost-Learning	0.4385	0.8433	0.0082	0.0000
Undersample	0.4701	0.8633	0.0000	0.0200
Oversample	0.4841	0.8733	0.0006	0.0300
Synthetic Padding	0.4151	0.8633	0.0037	0.0200

Figure A.IX.30: Raw data

MobileNetV3-small				
case	Cost per Test Sample	Avg TP+TN (%)	Avg FP (%)	Avg FN (%)
No Change	0.0780	0.844	0.0707	0.0850
Cost-Learning	0.0717	0.842	0.0573	0.1003
Undersample	0.0697	0.795	0.1247	0.0807
oversample	0.0640	0.872	0.1010	0.0267
pad w/ synthetic	0.0687	0.862	0.0640	0.0737

Figure A.IX.31: Raw data

6300  
6301  
6302  
6303  
6304  
6305  
6306  
6307  
6308  
6309  
6310  
6311  
6312  
6313  
6314  
6315  
6316  
6317  
6318  
6319  
6320  
6321  
6322  
6323  
6324  
6325  
6326  
6327  
6328  
6329  
6330  
6331  
6332  
6333  
6334  
6335  
6336  
6337  
6338  
6339  
6340  
6341  
6342  
6343  
6344  
6345  
6346  
6347  
6348  
6349 6350  
6351  
6352  
6353  
6354  
6355  
6356  
6357  
6358  
6359  
6360  
6361  
6362  
6363  
6364  
6365  
6366  
6367  
6368  
6369  
6370  
6371  
6372  
6373  
6374  
6375  
6376  
6377  
6378  
6379  
6380  
6381  
6382  
6383  
6384  
6385  
6386  
6387  
6388  
6389  
6390  
6391  
6392  
6393  
6394  
6395  
6396  
6397  
6398  
6399

6400	6450
6401	6451
6402	Acknowledgments 6452
6403	6453
6404	6454
6405	6455
6406	6456
6407	Thank you to Dr Creed Jones, Dr Ryan Williams, and Dr Jason Xuan for being a part of my committee, especially to Dr 6457
6408	Jones for giving me their ear and guidance throughout this project. 6458
6409	6459
6410	6460
6411	6461
6412	Additional thank yous to ... 6462
6413	6463
6414	Northrop Grumman Corporation for funding my Master's Degree at Virginia Tech University. 6464
6415	6465
6416	To the numerous teachers that I consulted with throughout this project. To Brandon Padayao with Baltimore OpenWorks 6466
6417	for allowing me to sit in on a welding course and learn from you about metal welding. To Riley Van Etten and Deion 6467
6418	Waddell for their expertise in 3D graphic design. To numerous others helping me brainstorm through ideas and barriers. 6468
6419	6469
6420	6470
6421	6471
6422	This project gave me the chance to learn about one of the most powerful graphics engines - one of many that fueled my 6472
6423	imagination as a child. 6473
6424	6474
6425	6475
6426	6476
6427	6477
6428	6478
6429	6479
6430	6480
6431	6481
6432	6482
6433	6483
6434	6484
6435	6485
6436	6486
6437	6487
6438	6488
6439	6489
6440	6490
6441	6491
6442	6492
6443	6493
6444	6494
6445	6495
6446	6496
6447	6497
6448	6498
6449	6499

6500	6550
6501	6551
6502	6552
6503	6553
6504	6554
6505	6555
6506	6556
6507	6557
6508	6558
6509	6559
6510	6560
6511	6561
6512	6562
6513	6563
6514	6564
6515	6565
6516	6566
6517	6567
6518	6568
6519	6569
6520	6570
6521	6571
6522	6572
6523	6573
6524	6574
6525	6575
6526	6576
6527	6577
6528	6578
6529	6579
6530	6580
6531	6581
6532	6582
6533	6583
6534	6584
6535	6585
6536	6586
6537	6587
6538	6588
6539	6589
6540	6590
6541	6591
6542	6592
6543	6593
6544	6594
6545	6595
6546	6596
6547	6597
6548	6598
6549	6599
Page intentionally left blank	

6600		6650
6601		6651
6602		6652
6603		6653
6604		6654
6605	I. Papers and Articles	6655
6606		
6607	Chuanqui Tan, Fuchun Sun, Tao Kong, Wenchang Zhang, Chao Yang, and Chunfang Liu. "A Survey on Deep Transfer Learning". State Key Laboratory of Intelligent Technology and Systems. Tsinghua National Laboratory for Information Science and Technology, Tsinghua University. Distributed by Cornell University. arXiv:1808.0197. Aug 6th, 2018	6656
6608		6657
6609		6658
6610		6659
6611	ECCV-18 Unknown author submission ID 2903. "Learning image classifiers from (limited) real and (abundant) synthetic data." Semantic Scholar. 2018. Corpus ID: 53472409. <a href="https://www.semanticscholar.org/paper/and-(abundant)-synthetic-data/1f759dbd74ad3907457dcbf50d4be54b470e33f4">https://www.semanticscholar.org/paper/and-(abundant)-synthetic-data/1f759dbd74ad3907457dcbf50d4be54b470e33f4</a> . Citation possibly for Nguyen, Tan, Chen, H, et al. "Learning image classifiers from (limited) real and (abundant) synthetic data." Published by Rice University. <a href="https://tannguyen.blogs.rice.edu/current-research/">https://tannguyen.blogs.rice.edu/current-research/</a> . 2018.	6660
6612	Gastelum, Zoe Nellie, and Shead, Timothy. How Low Can You Go? Using Synthetic 3D Imagery to Drastically Reduce Real-World Training Data for Object Detection. United States: N. p., 2020. Web. doi:10.2172/1670874.	6661
6613		6662
6614	Howard, Andrew, et al. "Searching for mobilenetv3." Proceedings of the IEEE/CVF international conference on computer vision. 2019.	6663
6615	Man, Keith, and Javaan Chahl. "A Review of Synthetic Image Data and Its Use in Computer Vision." Journal of imaging vol. 8,11 310. 21 Nov. 2022, doi:10.3390/jimaging8110310	6664
6616		6665
6617	Thornström, J. "Domain Adaptation of Unreal Images for Image Classification (Dissertation)." Linköping University Department of Electrical Engineering, Computer Vision (Sweden). 2019. URN: 165758. <a href="https://liu.diva-portal.org/smash/record.jsf?pid=diva2%3A1431281&amp;dswid=6387">https://liu.diva-portal.org/smash/record.jsf?pid=diva2%3A1431281&amp;dswid=6387</a>	6666
6618		6667
6619	Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." arXiv preprint arXiv:1409.1556 (2014).	6668
6620		6669
6621	S. J. Pan and Q. Yang, "A Survey on Transfer Learning," in IEEE Transactions on Knowledge and Data Engineering, vol. 22, no. 10, pp. 1345-1359, Oct. 2010, doi: 10.1109/TKDE.2009.191. keywords: {Machine learning;Training data;Data mining;Knowledge transfer;Space technology;Knowledge engineering;Machine learning algorithms;Labeling;Learning systems;Testing;Transfer learning;survey;machine learning;data mining.}	6670
6622		6671
6623	Weiss, Gary M, McCarthy, Kate, and Zabar, Bibi. "Cost-Sensitive Learning vs. Sampling: Which is Best for Handling Unbalanced Classes with Unequal Error Costs?". Department of Computer and Information Science, Fordham University. 2007. <a href="https://storm.cis.fordham.edu/~gweiss/papers/dmin07-weiss.pdf">https://storm.cis.fordham.edu/~gweiss/papers/dmin07-weiss.pdf</a>	6672
6624		6673
6625		
6626		
6627		
6628		
6629		
6630		
6631		
6632		
6633		
6634		
6635		
6636		
6637		
6638		
6639		
6640		
6641	II. Additional Resources	6691
6642		6692
6643	3D Ace. "Does Polygon Count matter in 3D Modeling for Game Assets?". Nicosia, Cyprus. Jan 3rd, 2024. <a href="https://3dace.com/blog/polygon-count-in-3d-modeling-for-game-assets/">https://3dace.com/blog/polygon-count-in-3d-modeling-for-game-assets/</a>	6693
6644		6694
6645	DowlingSoka, Ryan. "Python in Unreal Tips". Published on Mar 18th, 2024. <a href="https://ryandowlingsoka.com/unreal/python-in-unreal/">https://ryandowlingsoka.com/unreal/python-in-unreal/</a>	6695
6646		
6647	"Keras Documentation: Mobilenet, MobilenetV2, and MobileNetV3." Keras, Google, keras.io/api/applications/mobilenet/. Accessed 25 Apr. 2024.	6696
6648		6697
6649	OpenWorks, Baltimore City. Mar 16th, 2024. Metal 3 with Brandon Padayao. Articles: "Metal 3: MIG Welding" and "Metal 3: TIG Welding". 1400 Greenmount Ave, Baltimore, MD 21202	6698
	Padayao, Brandon. "Metal 3". Baltimore Openworks. Mar 16th, 2024.	6699
	V.3D Models – Cats	6691
		6692
	Anderdali, Cat. Turbosquid. FBX. Product ID: 1353041."Its a cat)+rig+face". <a href="https://www.turbosquid.com/3d-models/cat-1353041">https://www.turbosquid.com/3d-models/cat-1353041</a> . December 4, 2018	6693
		6694
	Idris_1986. "Cat". Turbosquid. ID: 1099111. "cat with blender create .basic bones and paint fur and bacis hair fur .eye textures etc features .others format ojb 3ds". <a href="https://www.turbosquid.com/3d-models/cat-1099111">https://www.turbosquid.com/3d-models/cat-1099111</a> . December 03, 2016	6695
		6696
	Printable_Models. "Cat V1". Free3D. .obj. ID: 123c6a1c5523. <a href="https://free3d.com/3d-model/cat-v1--579827.html">https://free3d.com/3d-model/cat-v1--579827.html</a> . September 6, 2018. Last updated November 20, 2019	6697
		6698
	Printable_Models. "Pallas Cat V1". Free3D. .obj. ID: 123c73c389b. <a href="https://free3d.com/3d-model/pallas-cat-v1--576987.html">https://free3d.com/3d-model/pallas-cat-v1--576987.html</a> . September 27, 2018. Last updated November 20, 2019	6699

6700		6750
6701		6751
6702		6752
6703	Printable_Models. "Cat V1". Free3D. .obj. ID: 123cb1b1943a.. "Cat v1 printable, low poly model.". <a href="https://free3d.com/3d-model/cat-v1-522281.html">https://free3d.com/3d-model/cat-v1-522281.html</a> . September 12, 2018. Last updated November 20, 2019	6753
6704		6754
6705	Printable_Models. "Cat V1". Free3D. .obj. ID: 123c3567cbad.. "Cat v1 printable, low poly model.". <a href="https://free3d.com/3d-model/cat-v1-220685.html">https://free3d.com/3d-model/cat-v1-220685.html</a> . October 10, 2018. Last updated November 20, 2019	6755
6706		6756
6707		6757
6708	Printable_Models. "Cat V3". Free3D. .obj. ID: 123c86a1fcd3.. "Cat V3 printable, low poly model". <a href="https://free3d.com/3d-model/cat-v3-351220.html">https://free3d.com/3d-model/cat-v3-351220.html</a> . October 10, 2018. Last updated November 20, 2019	6758
6709		6759
6710	Printable_Models. "Black Cat Back Arched V2". Free3D. .obj. ID: 123cdff0222e.. "Black Cat Back Arched V2 printable, low poly model". <a href="https://free3d.com/3d-model/black-cat-back-arched-v2-916864.html">https://free3d.com/3d-model/black-cat-back-arched-v2-916864.html</a> . October 10, 2018. Last updated November 20, 2019	6760
6711		6761
6712	Snippysnippets. "Low Poly Cat". Free3D. .obj. "Slightly low poly cat with whiskers". <a href="https://free3d.com/3d-model/low-poly-cat-46138.html">https://free3d.com/3d-model/low-poly-cat-46138.html</a> . February 25, 2017. Last updated November 20, 2019	6762
6713		6763
6714	Studio King. "Base mash Cat". Turbosquid. Product ID: 2080612. "Base Mesh Cat". <a href="https://www.turbosquid.com/3d-models/base-mash-cat-2080612">https://www.turbosquid.com/3d-models/base-mash-cat-2080612</a> . June 7, 2022	6764
6715		6765
6716		6766
6717		6767
6718		6768
6719	VI. 3D Models – Dogs	6769
6720		6770
6721	3d_logoman. "Golden Retriever". Turbosquid. ID: 1845760. "Low poly model, created for 3D printing. You can also use it as a base model for your projects. This model can be used as a decoration in your interior. We make models for individual orders. See my other models". <a href="https://www.turbosquid.com/3d-models/golden-retriever-1850419">https://www.turbosquid.com/3d-models/golden-retriever-1850419</a> . February 05, 2022	6771
6722		6772
6723		6773
6724		6774
6725	Lima, Elvair. "White Dog". Turbosquid. ID: 1845760. "Low poly. Game ready.Textures included an materials applied. All formats tested and working. Textures PBR 2048x2048. I's not rigged.". <a href="https://www.turbosquid.com/3d-models/white-dog-1845760">https://www.turbosquid.com/3d-models/white-dog-1845760</a> . January 26, 2022	6775
6726		6776
6727		6777
6728		6778
6729	lashkoalex. "English Bulldog Real-Time Free Sample". Turbosquid. ID: 1155164. "FREE Sample of a Realistic English Bulldog Real-Time Game-Ready 3D Model. FULL version of this model is available for purchase under ID 1111972". <a href="https://www.turbosquid.com/3d-models/english-bulldog-real-time-free-sample-1155164">https://www.turbosquid.com/3d-models/english-bulldog-real-time-free-sample-1155164</a> . May 9, 2017	6779
6730		6780
6731		6781
6732	Printable_Models. "Australian Cattle Dog V1". Free3D. .obj. ID: 123c86a1fcd3.. "Australian Cattle Dog v1 printable, low poly model". <a href="https://free3d.com/3d-model/australian-cattle-dog-v1-993323.html">https://free3d.com/3d-model/australian-cattle-dog-v1-993323.html</a> . September 27, 2018. Last updated November 20, 2019	6782
6733		6783
6734		6784
6735		6785
6736	Printable_Models. "Canaan Dog V1". Free3D. .obj. ID: 123c82cb7c11. "Canaan Dog v1 printable, low poly model". <a href="https://free3d.com/3d-model/canaan-dog-v1-72376.html">https://free3d.com/3d-model/canaan-dog-v1-72376.html</a> . September 27, 2018. Last updated November 20, 2019	6786
6737		6787
6738		6788
6739	Printable_Models. "Cardigan Welsh Corgi V1". Free3D. .obj. ID: 123ce3e4b2fe.. "Cardigan Welsh Corgi v1 printable, low poly model". <a href="https://free3d.com/3d-model/cardigan-welsh-corgi-v1-784056.html">https://free3d.com/3d-model/cardigan-welsh-corgi-v1-784056.html</a> . September 27, 2018. Last updated November 20, 2019	6789
6740		6790
6741		6791
6742	Printable_Models. "ConcreteDogStatue V1". Free3D. .obj. ID: 123c476a785d. "ConcreteDogStatue v1 printable, low poly model". <a href="https://free3d.com/3d-model/concretedogstatue-v1-25932.html">https://free3d.com/3d-model/concretedogstatue-v1-25932.html</a> . September 27, 2018. Last updated November 20, 2019	6792
6743		6793
6744		6794
6745	Printable_Models. "Dog V2". Free3D. .obj. ID: 123ca88417bb. "Dog v2 printable, low poly model". <a href="https://free3d.com/3d-model/dog-v2-703191.html">https://free3d.com/3d-model/dog-v2-703191.html</a> . October 10, 2018. Last updated November 20, 2019	6795
6746		6796
6747		6797
6748	Printable_Models. "Farm Dog V2". Free3D. .obj. ID: 123c96195bae. "Farm Dog v2 printable, low poly model". <a href="https://free3d.com/3d-model/farm-dog-v2-499533.html">https://free3d.com/3d-model/farm-dog-v2-499533.html</a> . October 10, 2018. Last updated November 20, 2019	6798
6749		6799
	Rocazella, Sam. "Alef". Free3D. .obj. ID: 9v371wmvoc.. "This model was edited solo by me (edited by: Sam Rocazella) from the original balto model. big thanks to the original creator of the balto model letting this be possible.". <a href="https://free3d.com/3d-model/alef-44920.html">https://free3d.com/3d-model/alef-44920.html</a> . August 8, 2013. Last updated November 20, 2019	